



LUND
UNIVERSITY

350

Image Analysis (FMAN20)

Lecture 4, 2018

MAGNUS OSKARSSON



Image Analysis - Motivation

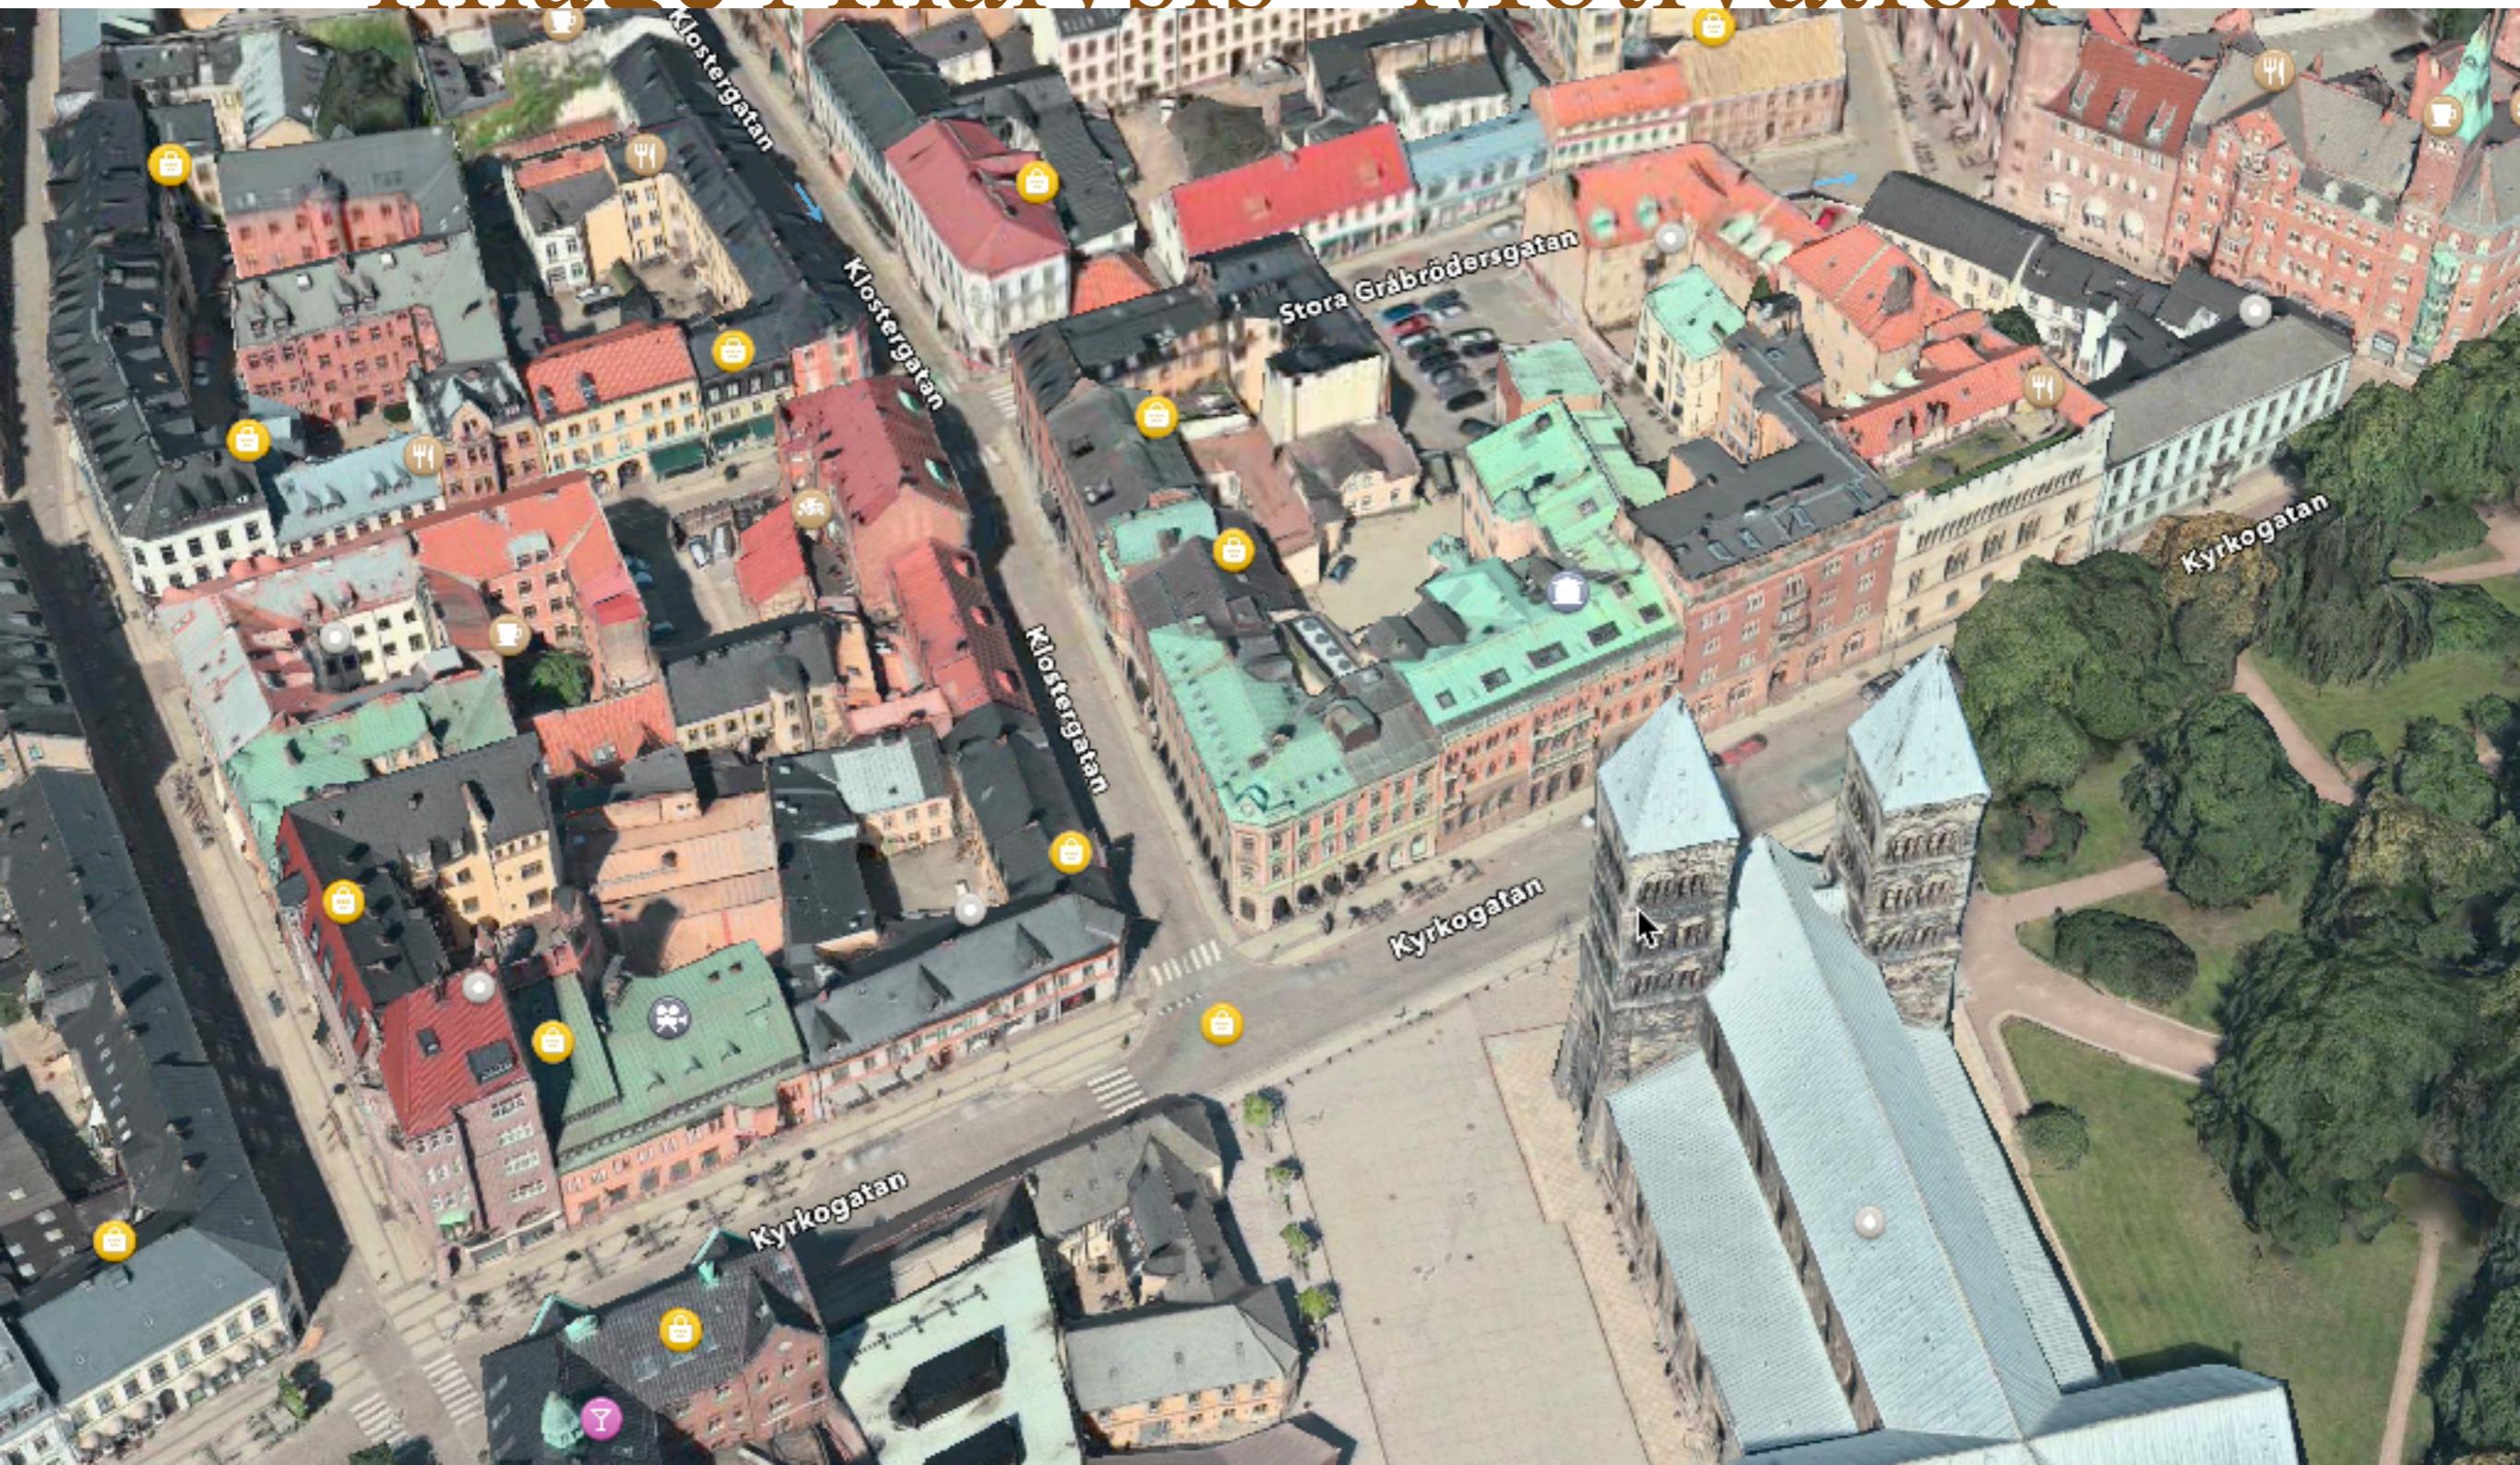
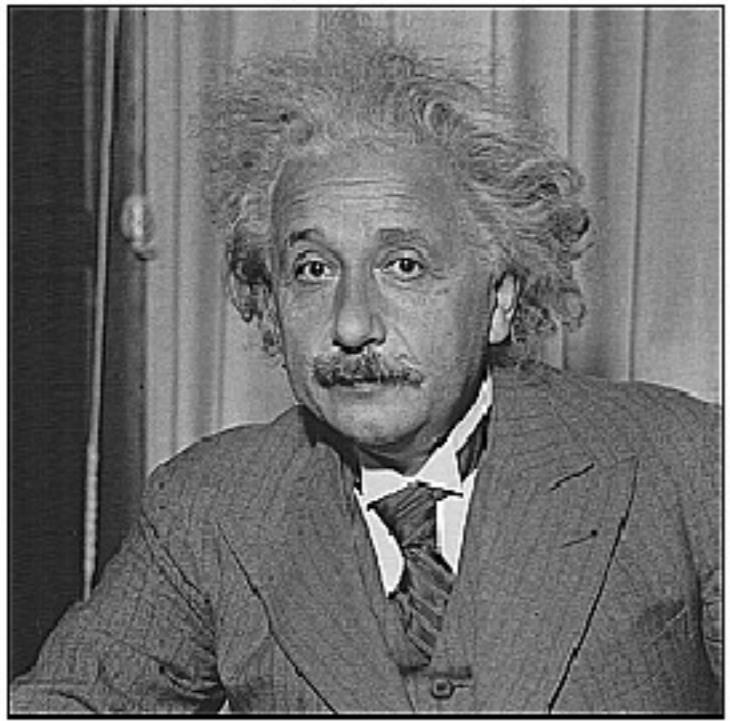
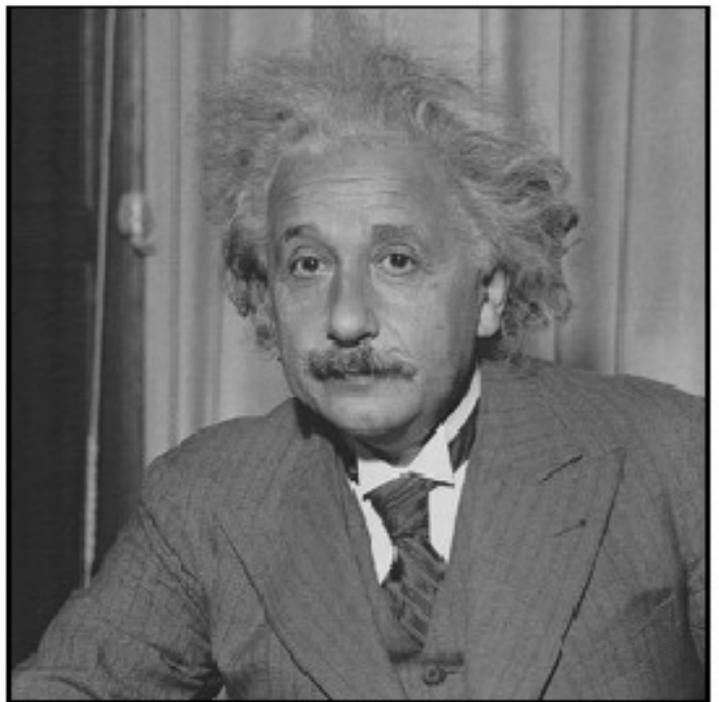


Image Analysis - Motivation



LUND
UNIVERSITY

Today: Image Filters



Smooth/Sharpen Images... Find edges...

Find waldo...

Overview – Convolutions

1. Convolution

1. Uses information in several pixels
2. Can be seen as scalar product for all submatrices
3. Can be seen as linear classifier for submatrices
4. Detects things that looks like itself

2. Convolution theorem

3. Connecting linear algebra, Fourier transform and convolutions

Overview – Convolutions

1. Convolution
 1. Definition, properties
 2. Convolution vs Cross-correlation
 3. Convolution and translation invariant linear systems
 4. Motivation using sliding means (1D and 2D)
 5. Interpretation as ‘sliding’ scalar product.
 6. Median Filter (not a convolution)
 7. Gaussian smoothing
 8. Derivatives + Smoothing
2. Convolution theorem
3. Connecting linear algebra, Fourier transform and convolutions

Convolution Operator

$$g = f * h$$

$$g(i, j) = \sum_u \sum_v f(i - u, j - v)h(u, v)$$



Cross-Correlation

Sliding scalar product

$$g(i, j) = \sum_y \sum_x f(i + y, j + x) h(y, x)$$

Compare with convolution

$$g(i, j) = \sum_u \sum_v f(i - u, j - v) \check{h}(u, v)$$

$$\check{h}(u, v) = h(-u, -v)$$

Why use convolution?

Cross-correlation seems much simpler.

One motivation: Convolution has simpler calculation rules

$$f * h = h * f$$

$$f * (g * h) = (f * g) * h$$

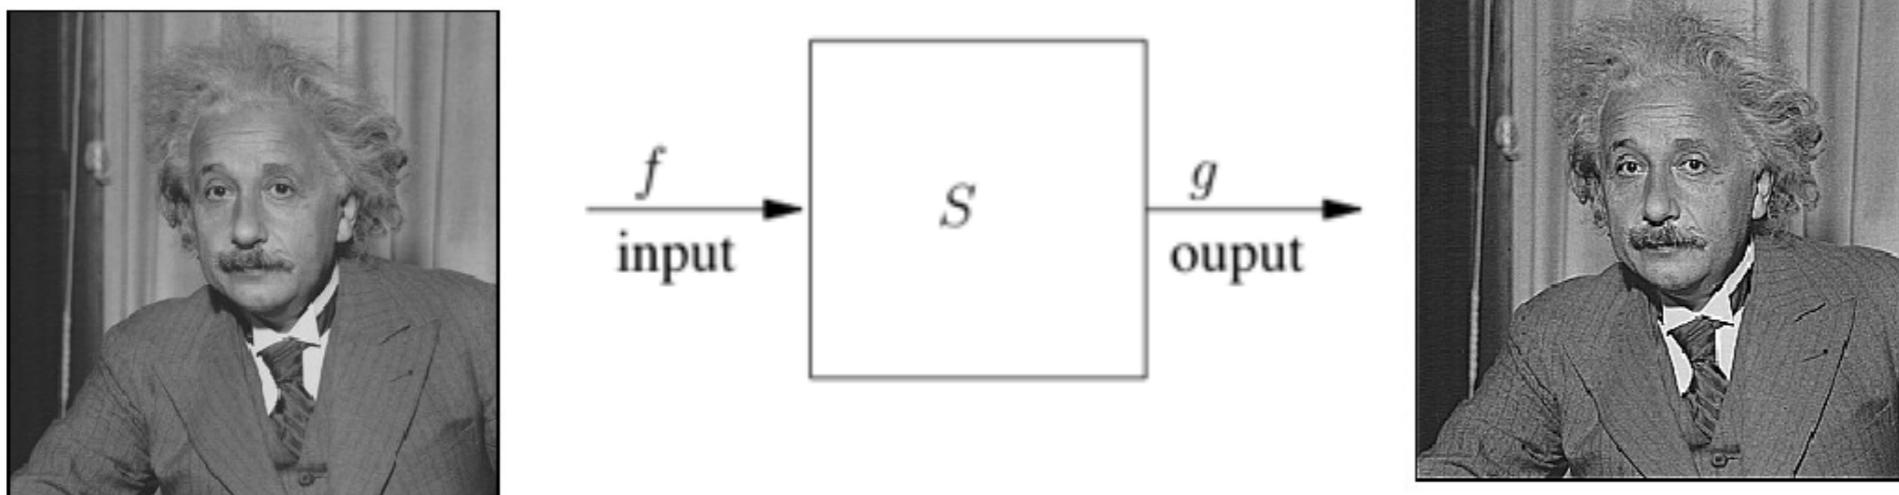
$$f * (g + h) = g * g + f * h$$

$$a(f * g) = (af) * g$$

$$\delta * f = f$$

$$\partial(f * g) = (\partial f) * g$$

Convolutions and linear systems



Any linear and translation invariant system can be represented as a convolution.

Motivation: noise reduction

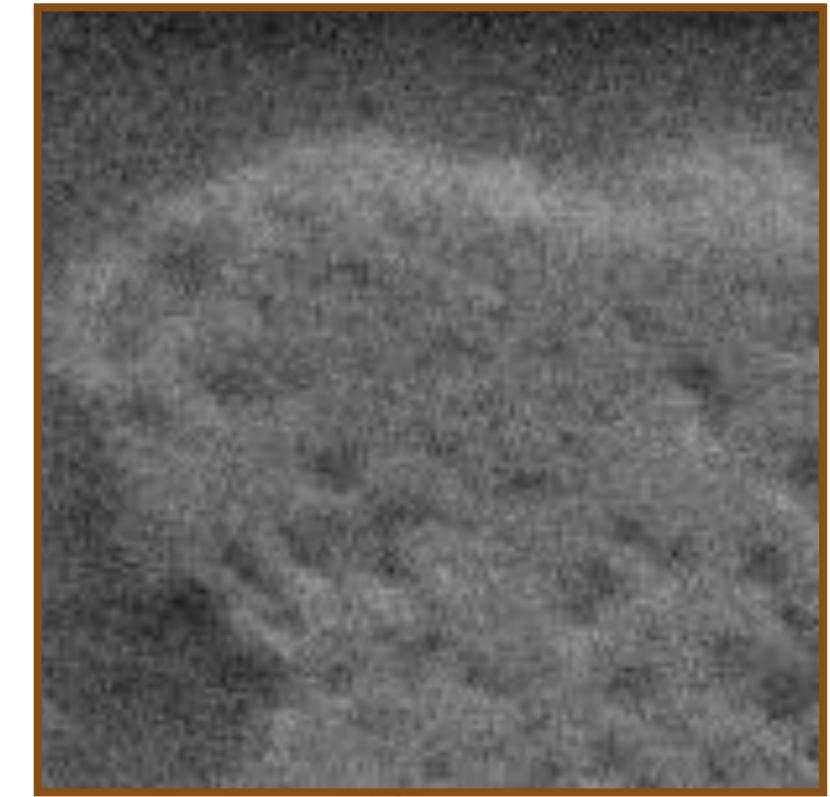
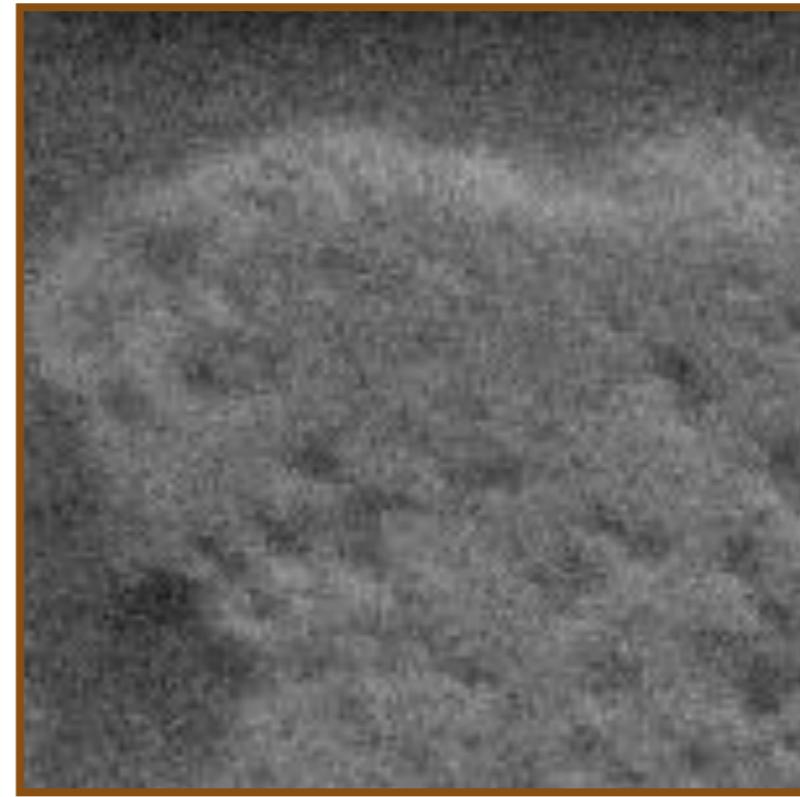
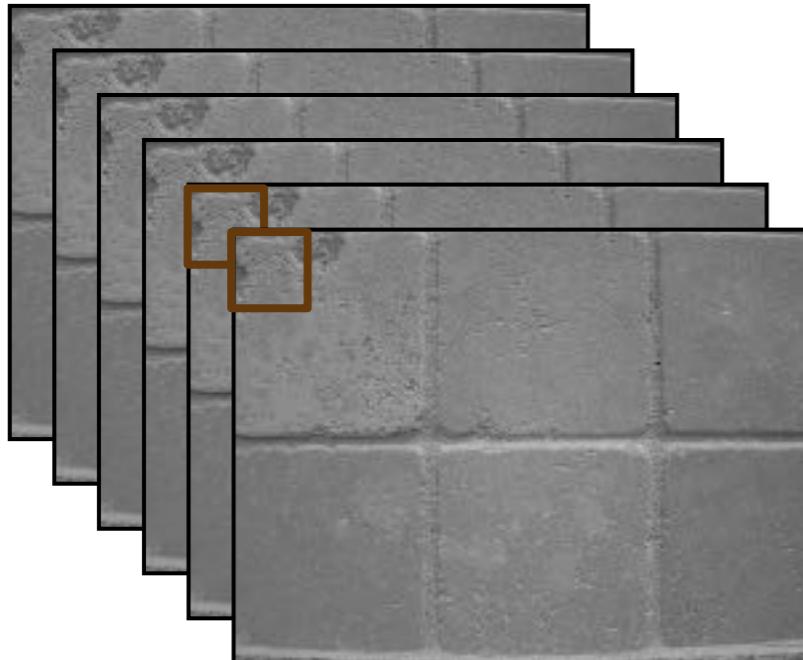
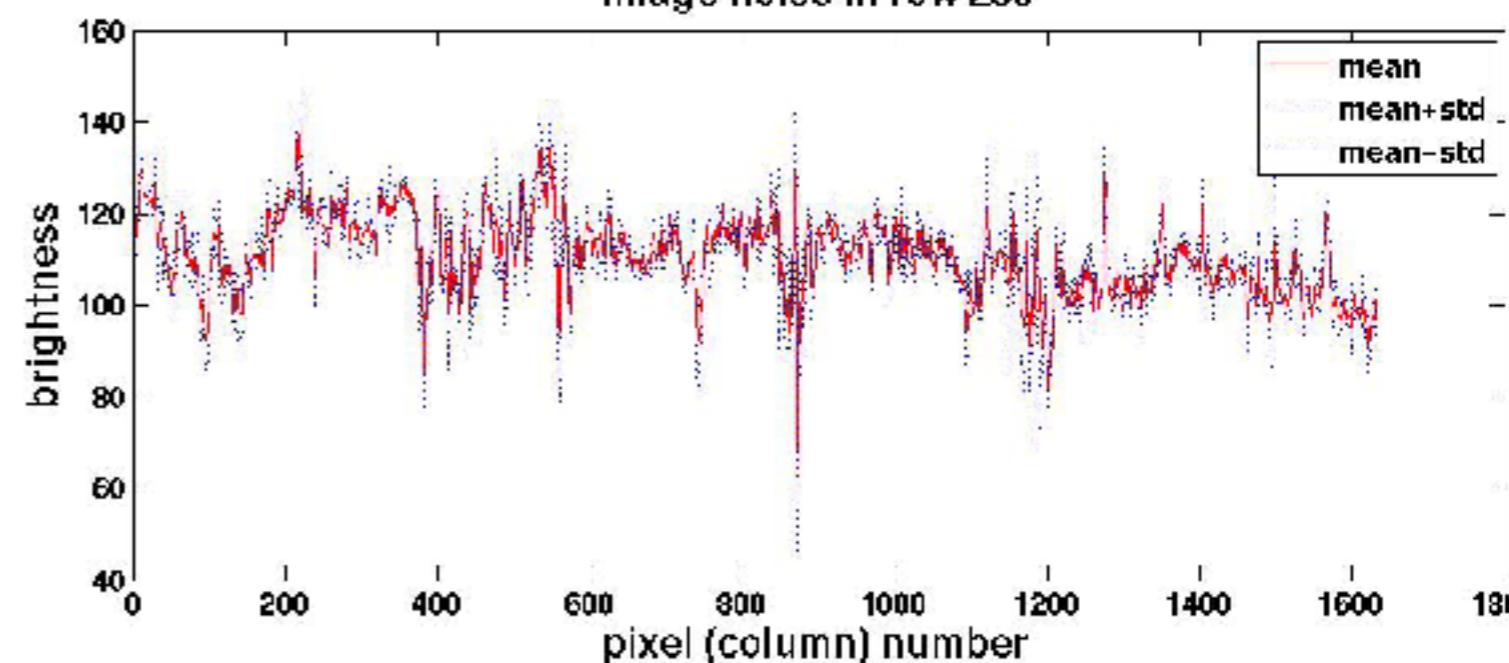
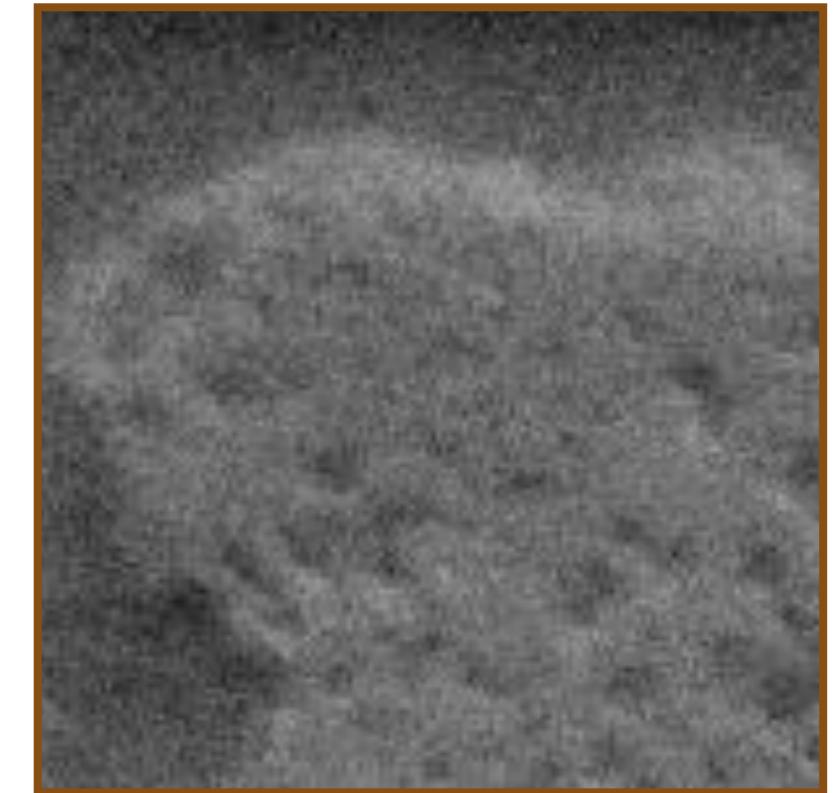
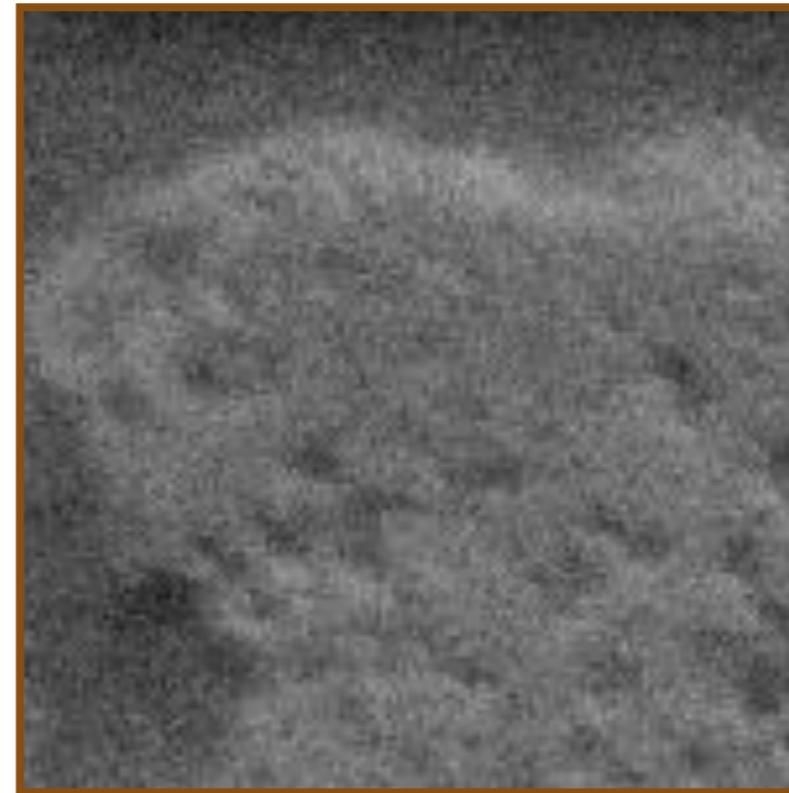
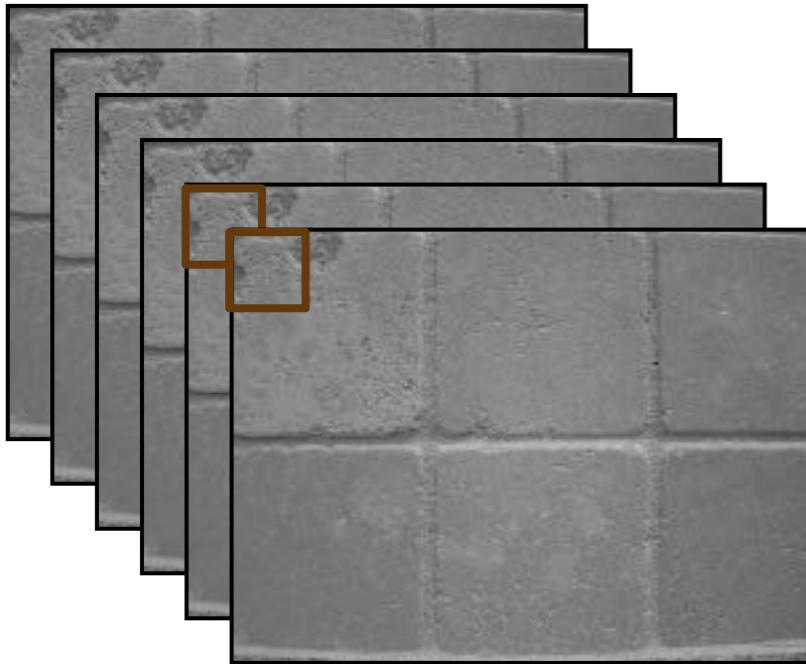


Image noise in row 250

- We can measure **noise** in multiple images of the same static scene.
- How could we reduce the noise, i.e., give an estimate of the true intensities?



Motivation: noise reduction



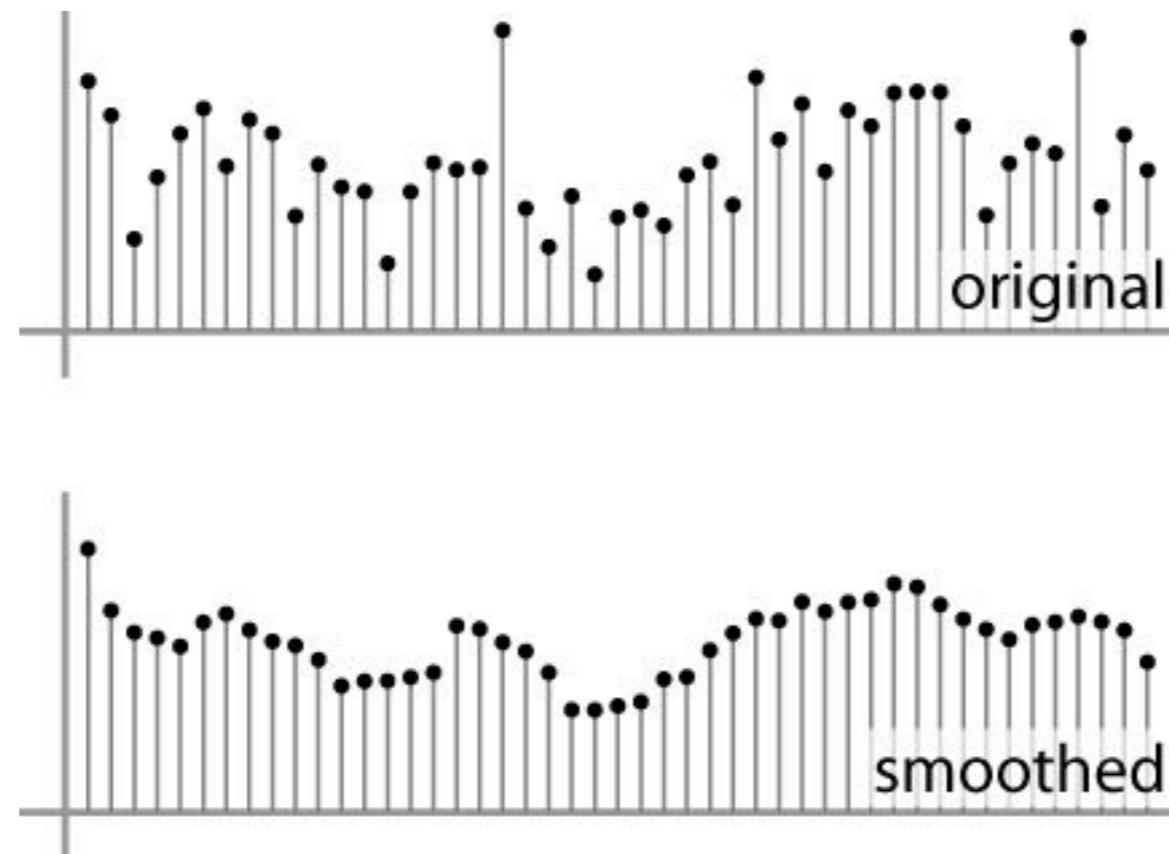
- How could we reduce the noise, i.e., give an estimate of the true intensities?
- **What if there's only one image?**

First attempt at a solution

- Let's replace each pixel with an average of all the values in its neighborhood
- Assumptions:
 - Expect pixels to be like their neighbors
 - Expect noise processes to be independent from pixel to pixel

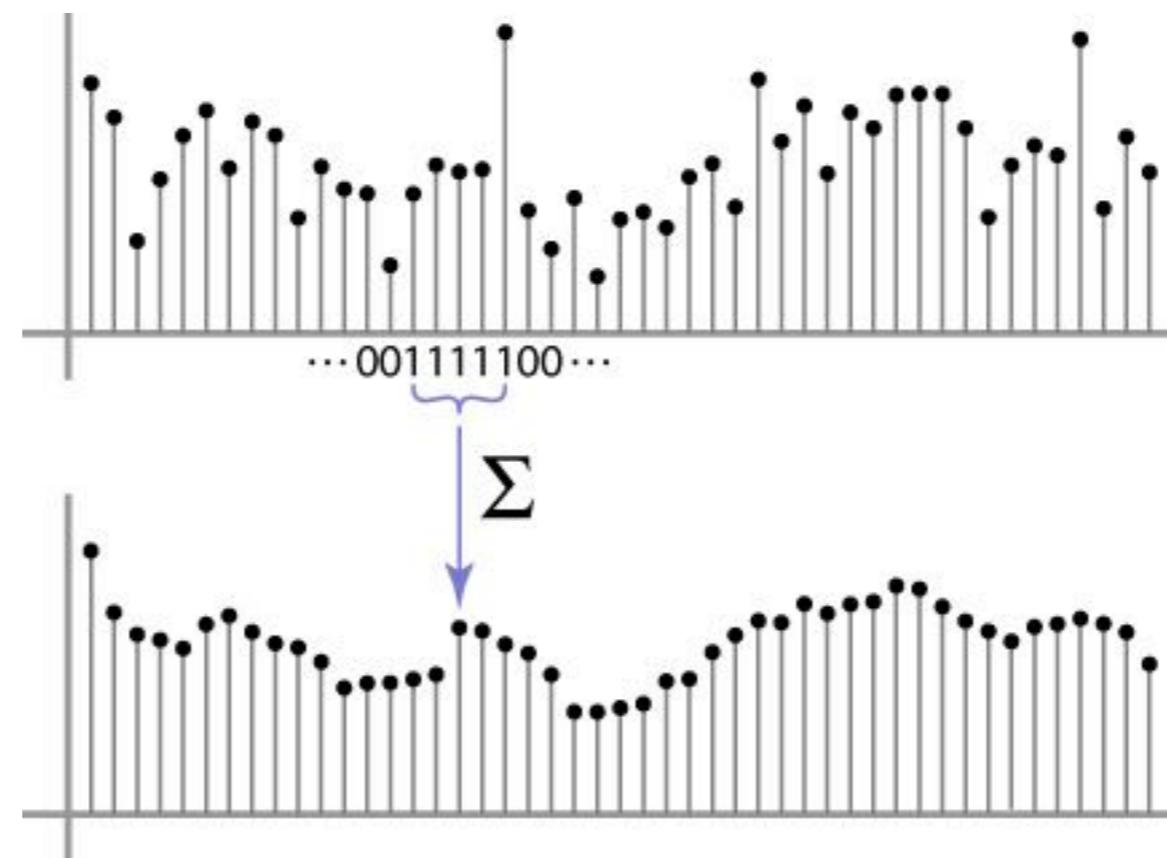
First attempt at a solution

- Let's replace each pixel with an average of all the values in its neighborhood
- Moving average in 1D:



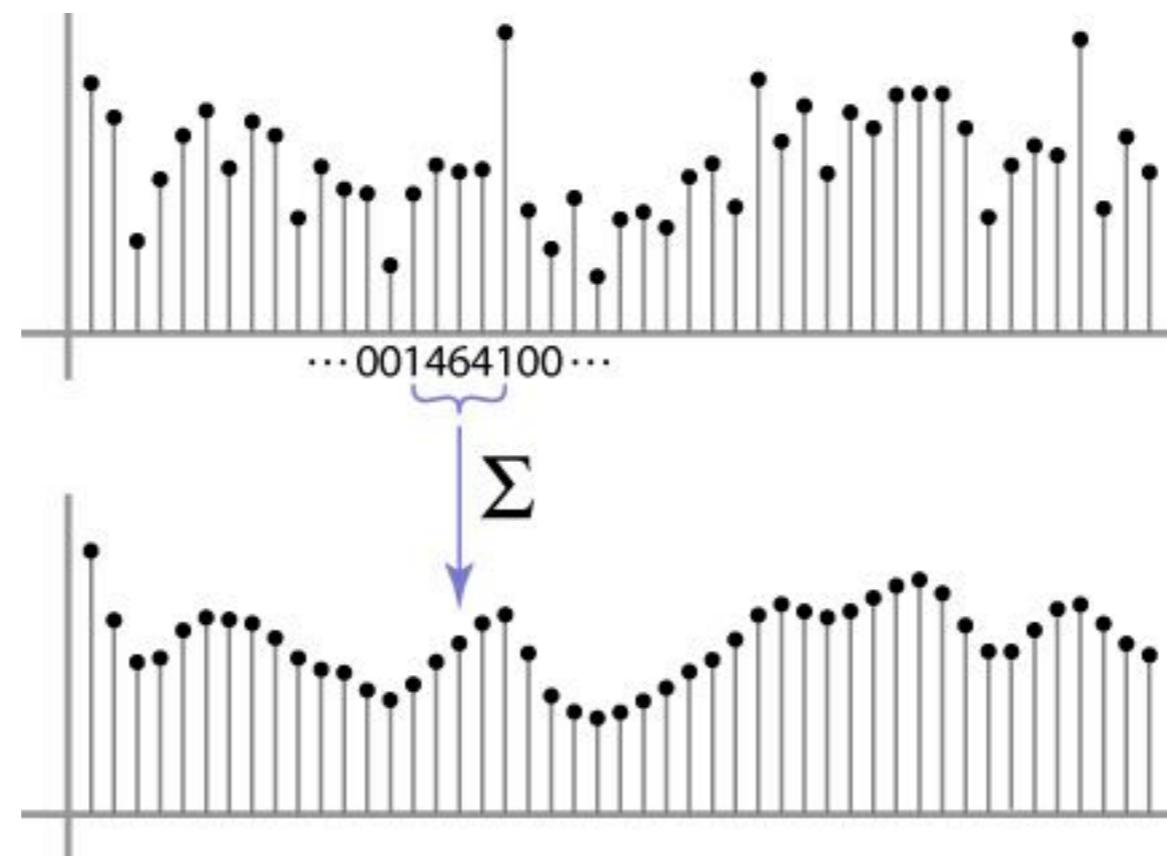
Weighted Moving Average

- Can add weights to our moving average
- *Weights* $[1, 1, 1, 1, 1] / 5$



Weighted Moving Average

- Non-uniform weights $[1, 4, 6, 4, 1] / 16$



Moving Average In 2D

$F[x, y]$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	0	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

$G[x, y]$

			0							

Moving Average In 2D

$F[x, y]$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	0	90	90	90	0	0	0
0	0	0	90	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

$G[x, y]$

	0	10								

Moving Average In 2D

$F[x, y]$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	0	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

$G[x, y]$

	0	10	20							

Moving Average In 2D

$F[x, y]$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	0	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

$G[x, y]$

	0	10	20	30						

Moving Average In 2D

$F[x, y]$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	0	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

$G[x, y]$

	0	10	20	30	30					

Moving Average In 2D

$F[x, y]$

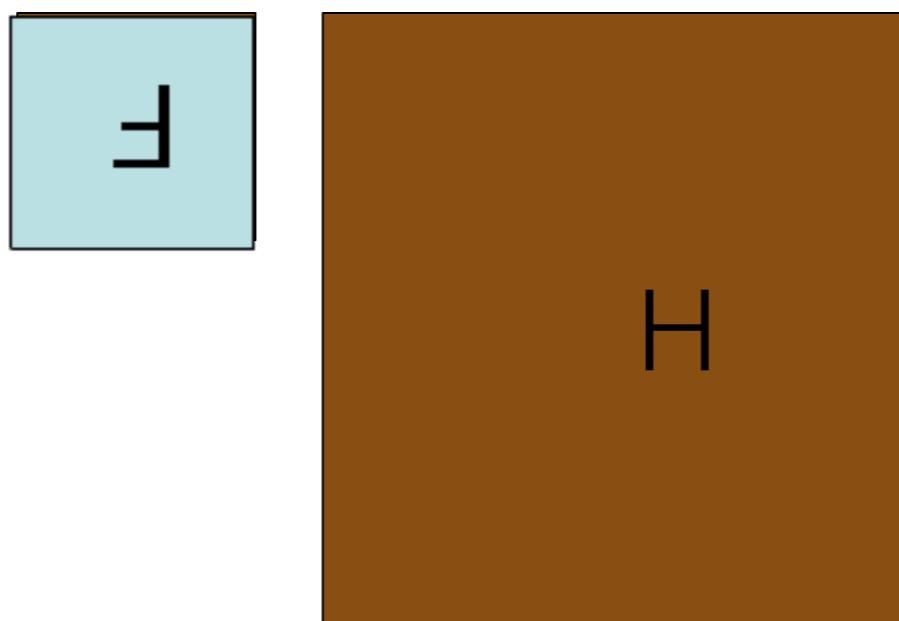
0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$G[x, y]$

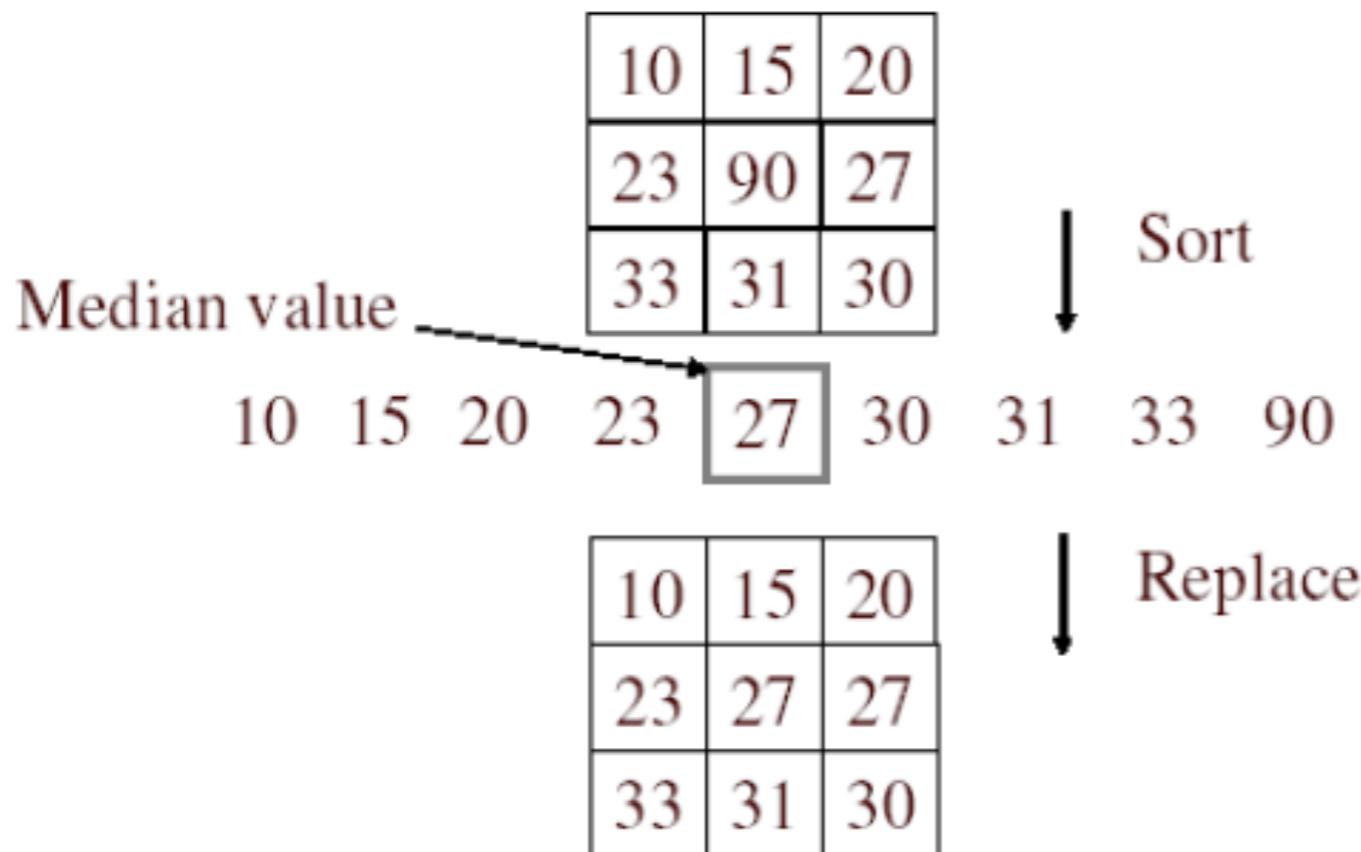
	0	10	20	30	30	30	20	10	
	0	20	40	60	60	60	40	20	
	0	30	60	90	90	90	60	30	
	0	30	50	80	80	90	60	30	
	0	30	50	80	80	90	60	30	
	0	20	30	50	50	60	40	20	
	10	20	30	30	30	30	20	10	
	10	10	10	0	0	0	0	0	

Convolution - repetition

- Convolution:
 - Flip the filter in both dimensions (bottom to top, right to left)
 - Then apply cross-correlation
 - Produces scalar product of flipped filter at every position!



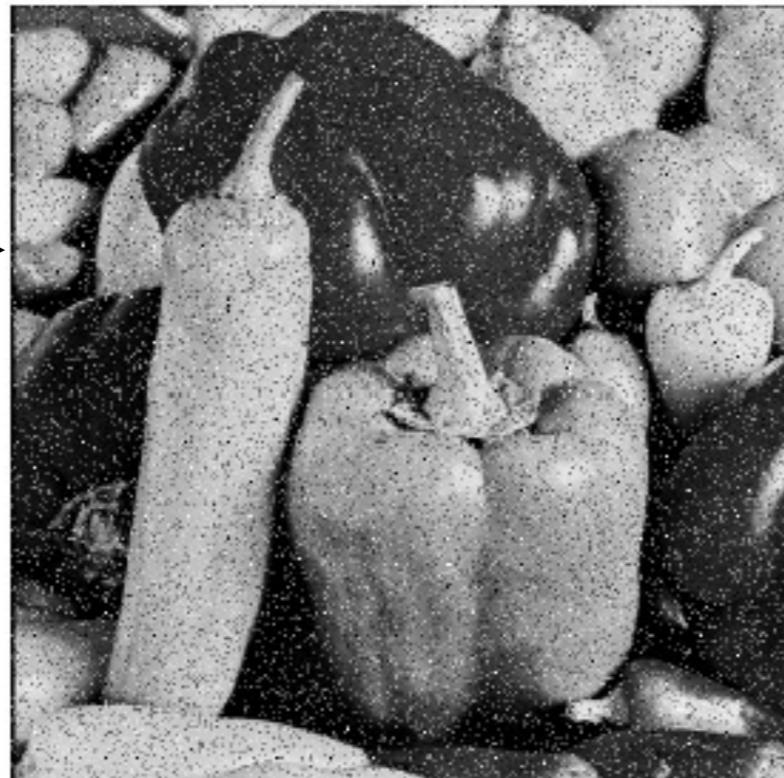
Median filter – an example of a non-linear sliding window smoother (Not a convolution)



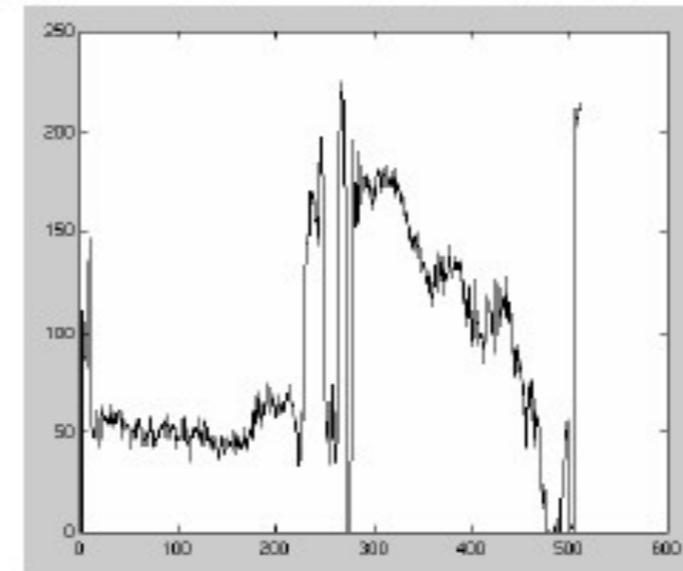
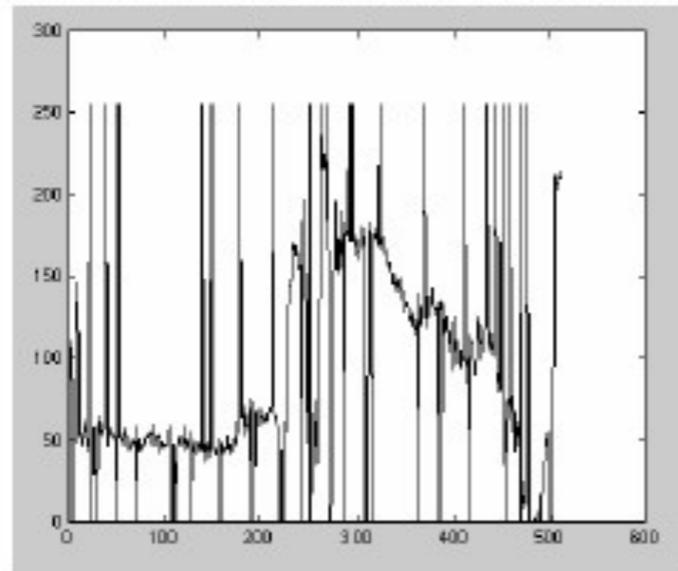
- No new pixel values introduced
- Removes spikes: good for impulse, salt & pepper noise
- Not linear
- Not a convolution

Median filter

Salt and
pepper
noise



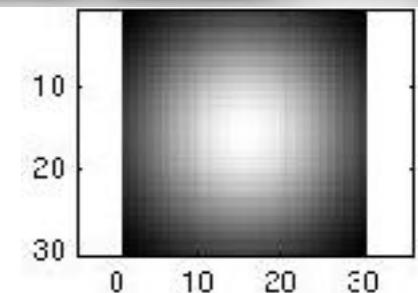
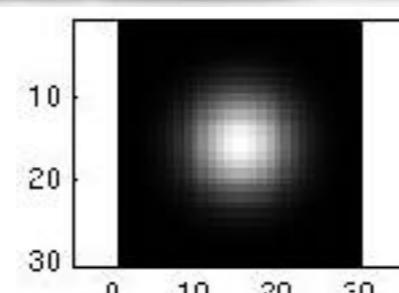
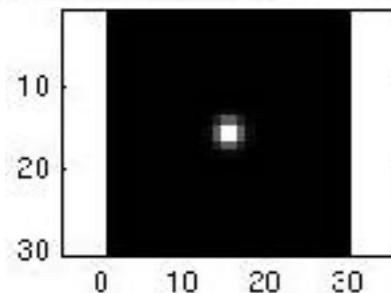
← Median
filtered



Plots of a row of the image

Smoothing with a Gaussian

Parameter σ is the “scale” / “width” / “spread” of the Gaussian kernel, and controls the amount of smoothing.

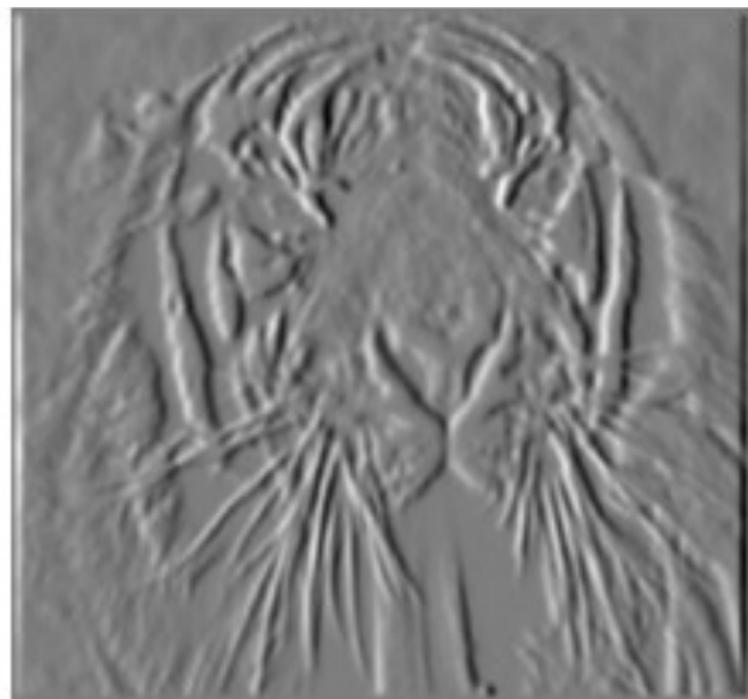


Partial derivatives of an image



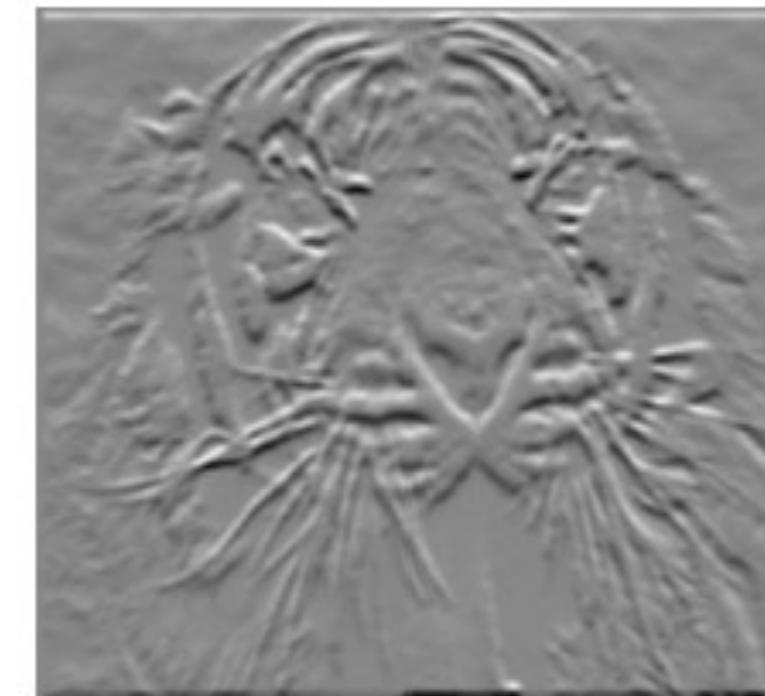
$$\frac{\partial f(x, y)}{\partial x}$$

-1	1
----	---



$$\frac{\partial f(x, y)}{\partial y}$$

-1	?	1
1	Or	-1

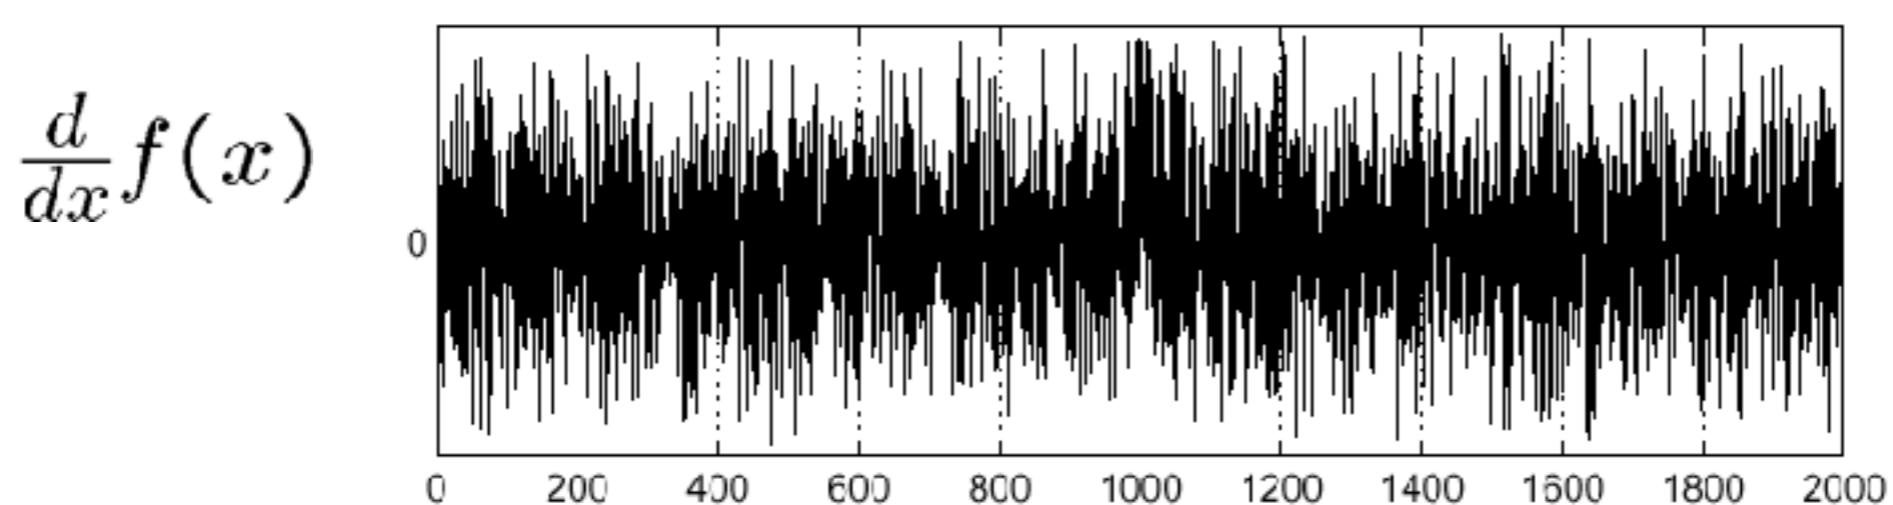
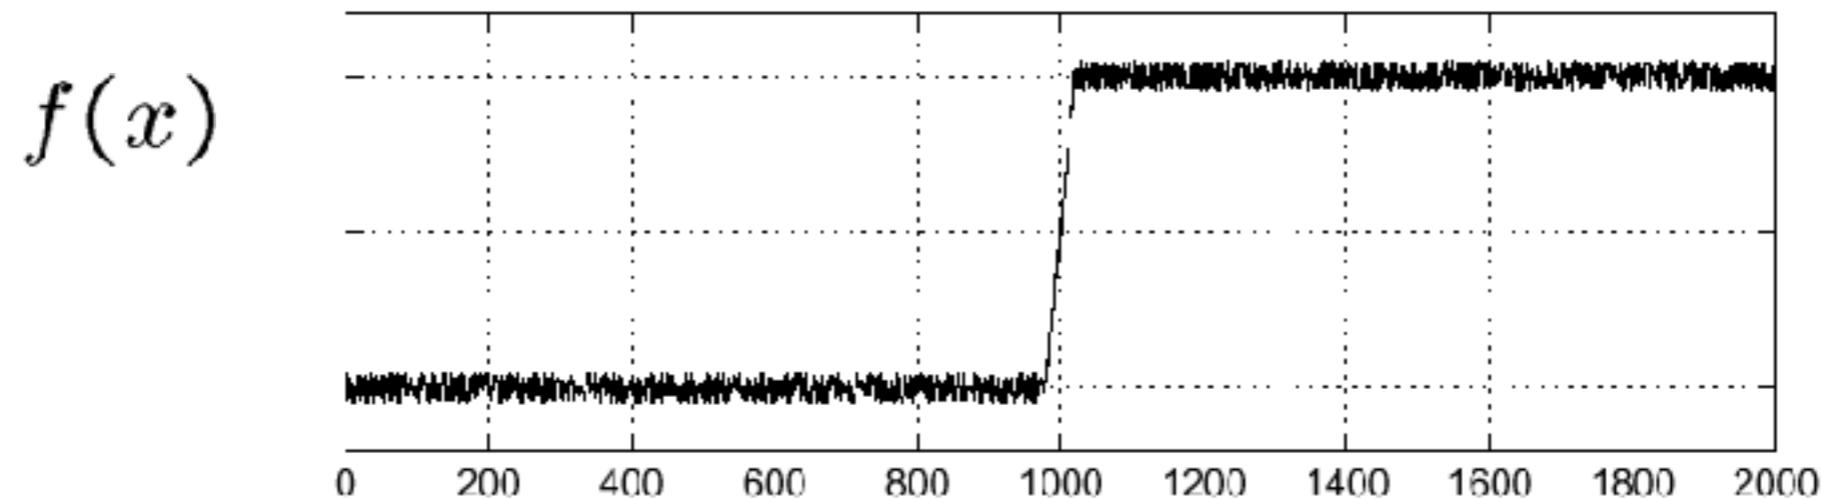


Which shows changes with respect to x?
(showing flipped filters)

Effects of noise

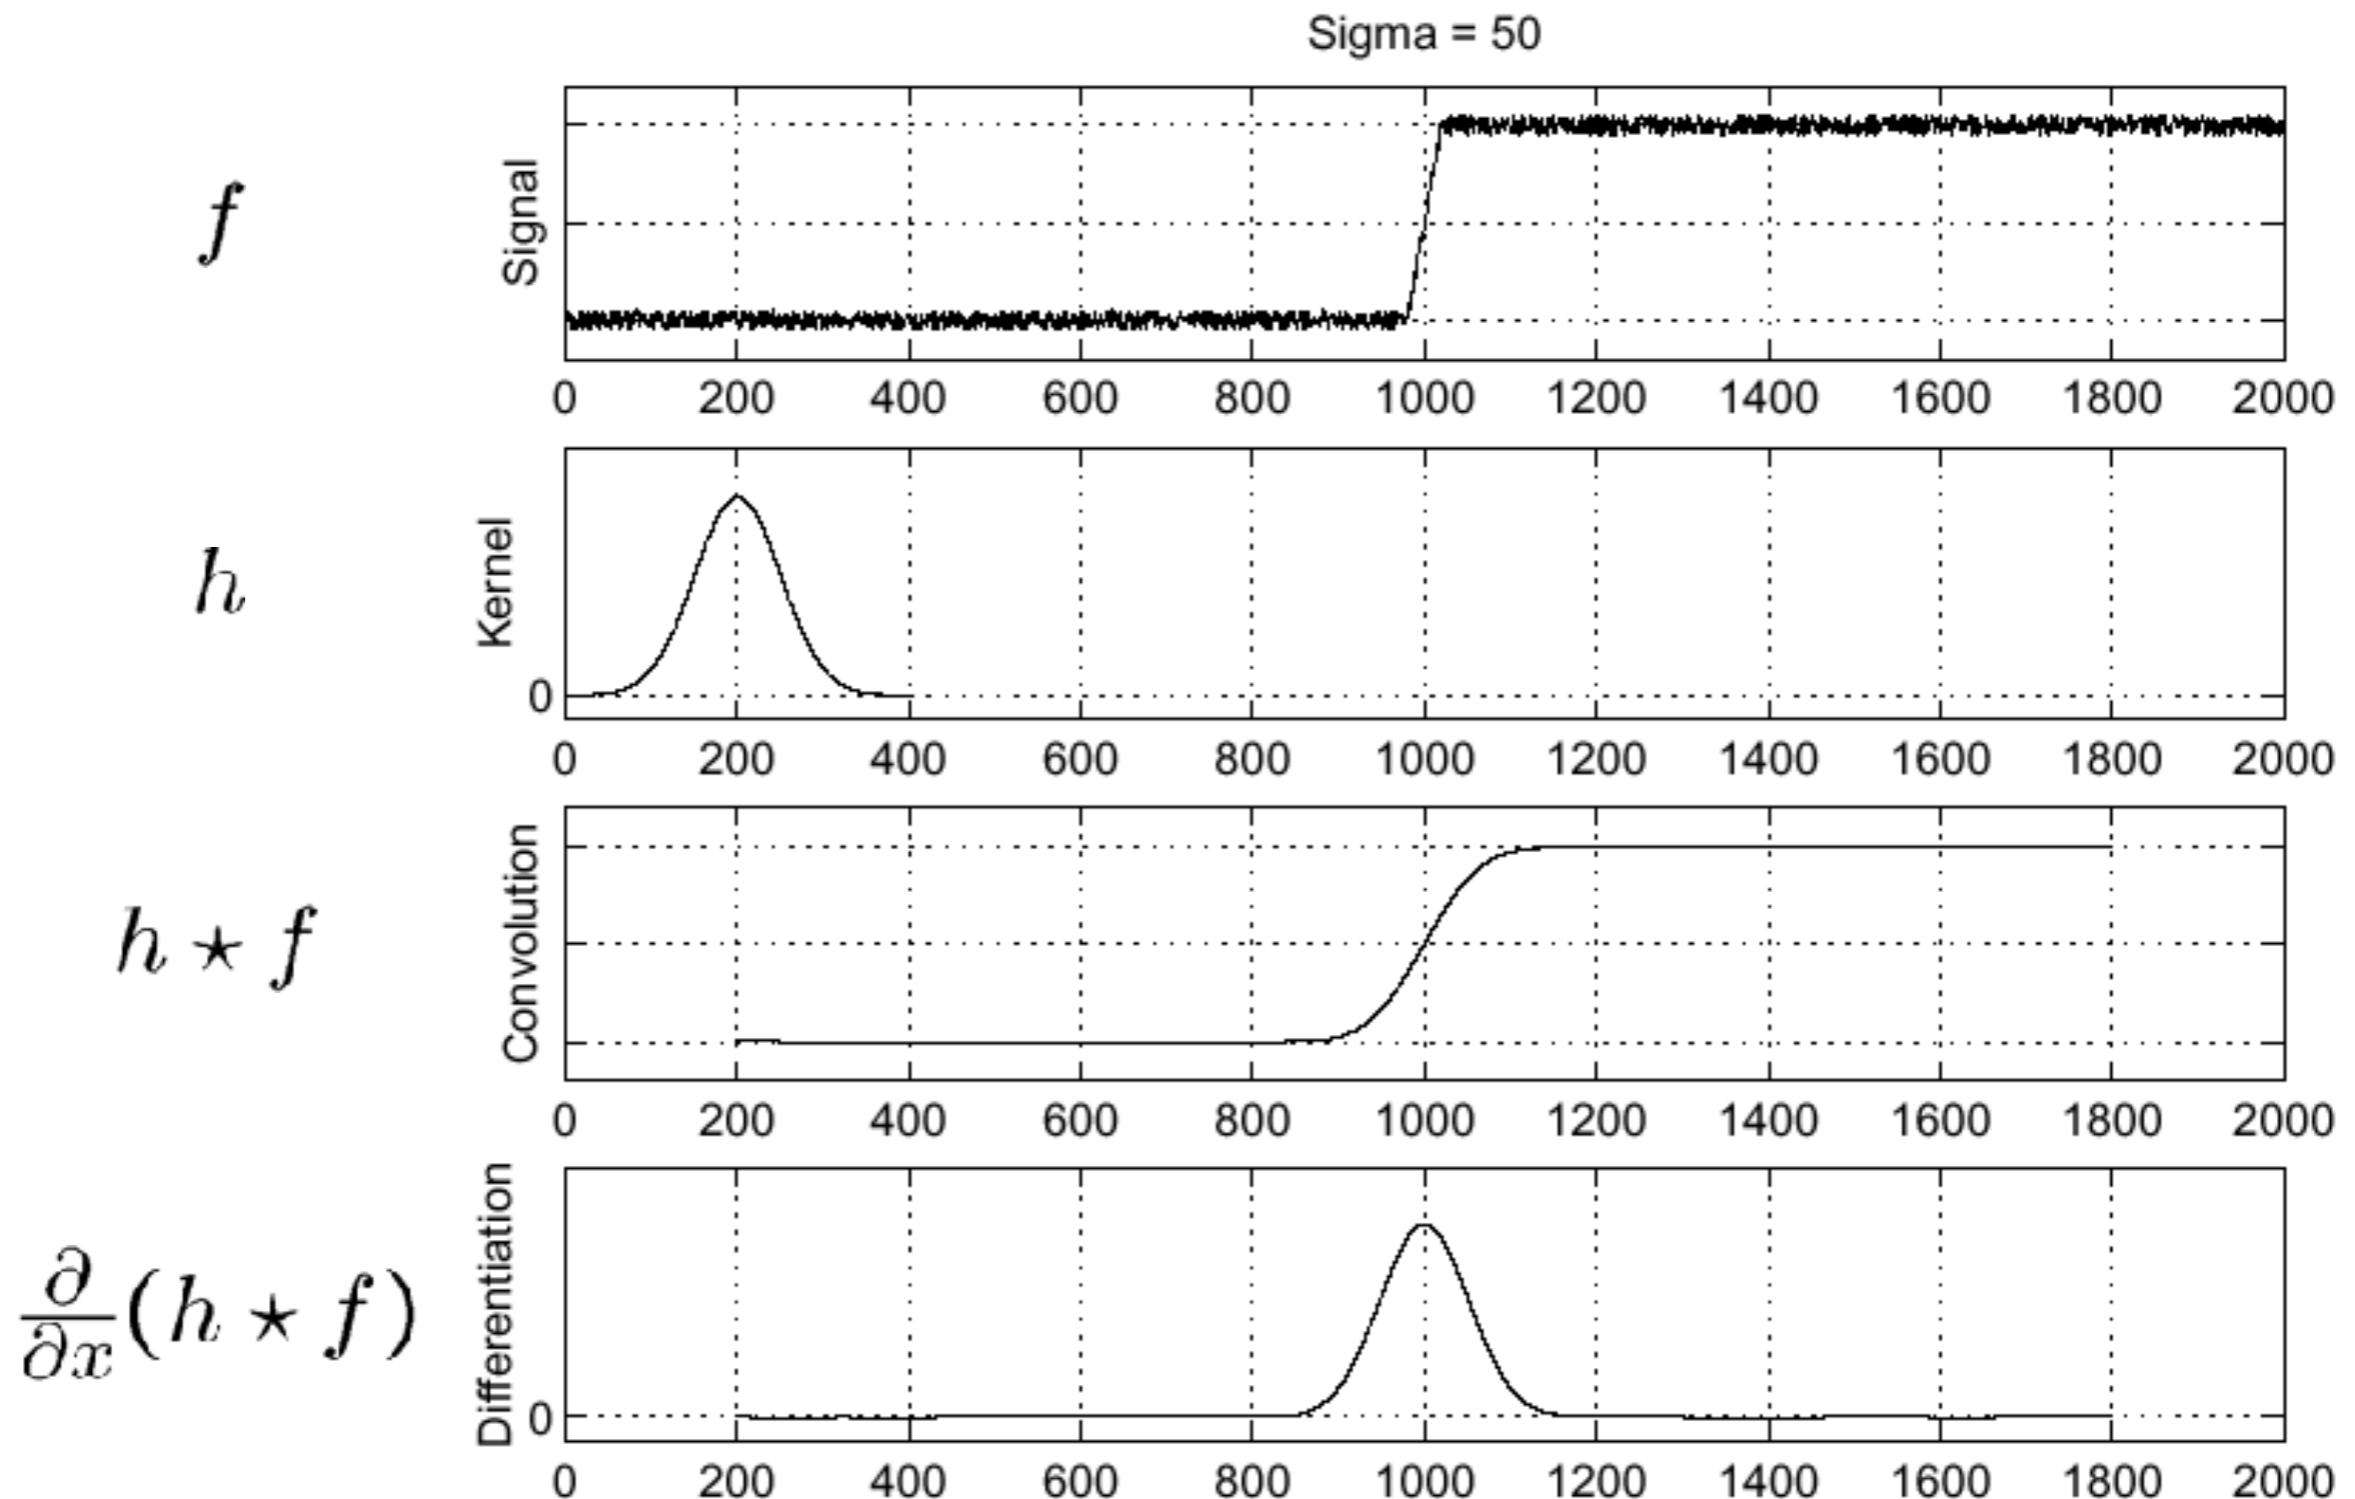
Consider a single row or column of the image

- Plotting intensity as a function of position gives a signal



Where is the edge?

Solution: smooth first



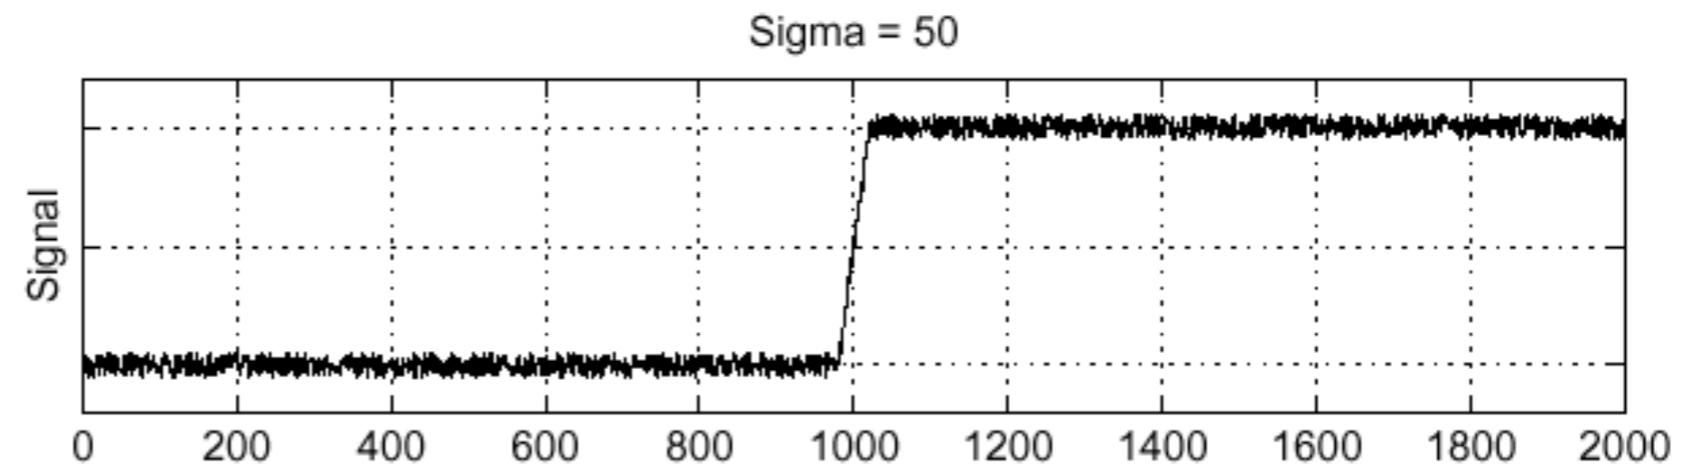
Where is the edge? Look for peaks in $\frac{\partial}{\partial x}(h \star f)$

Derivative property of convolution

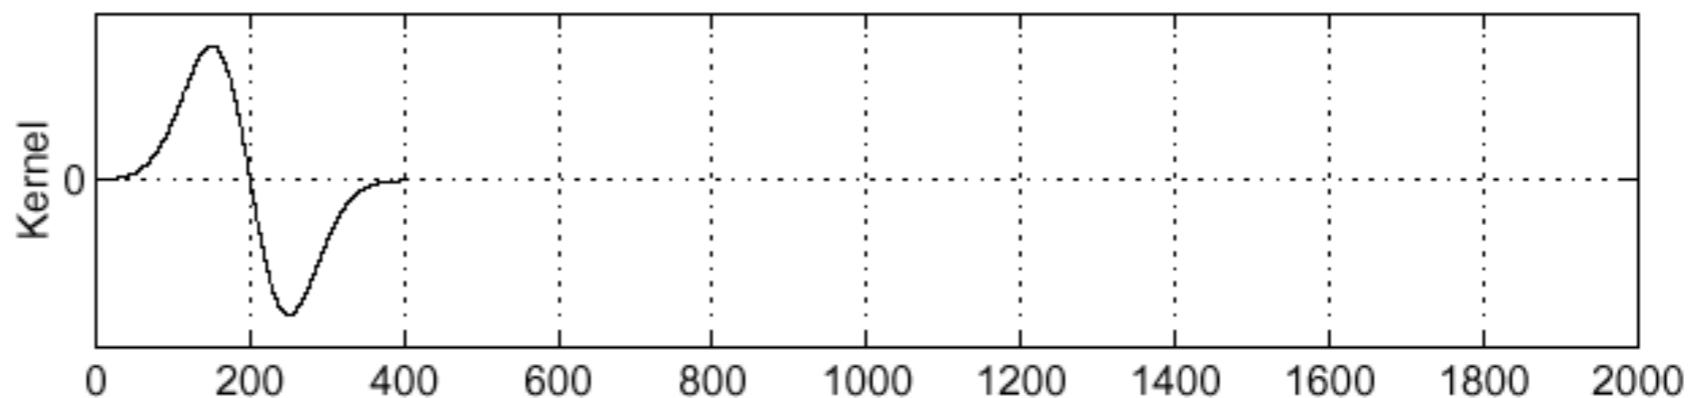
$$\frac{\partial}{\partial x}(h \star f) = (\frac{\partial}{\partial x}h) \star f$$

Differentiation property of convolution.

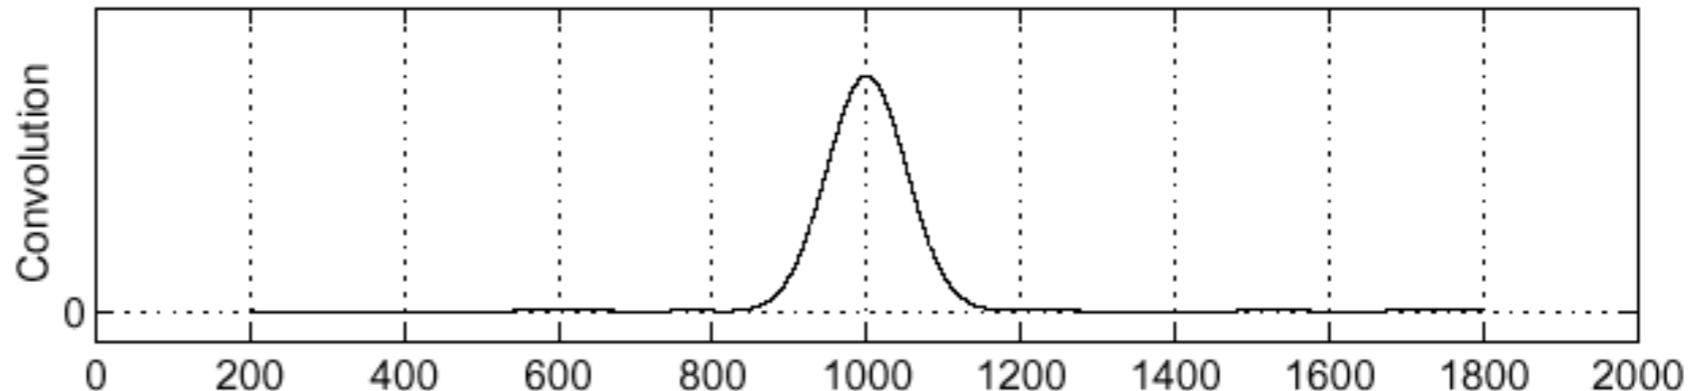
f



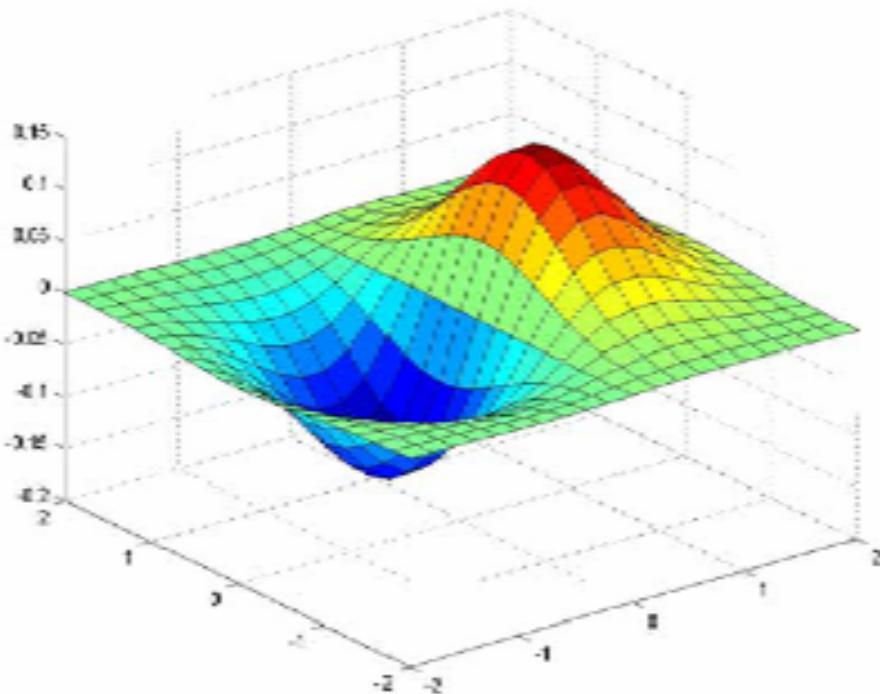
$\frac{\partial}{\partial x}h$



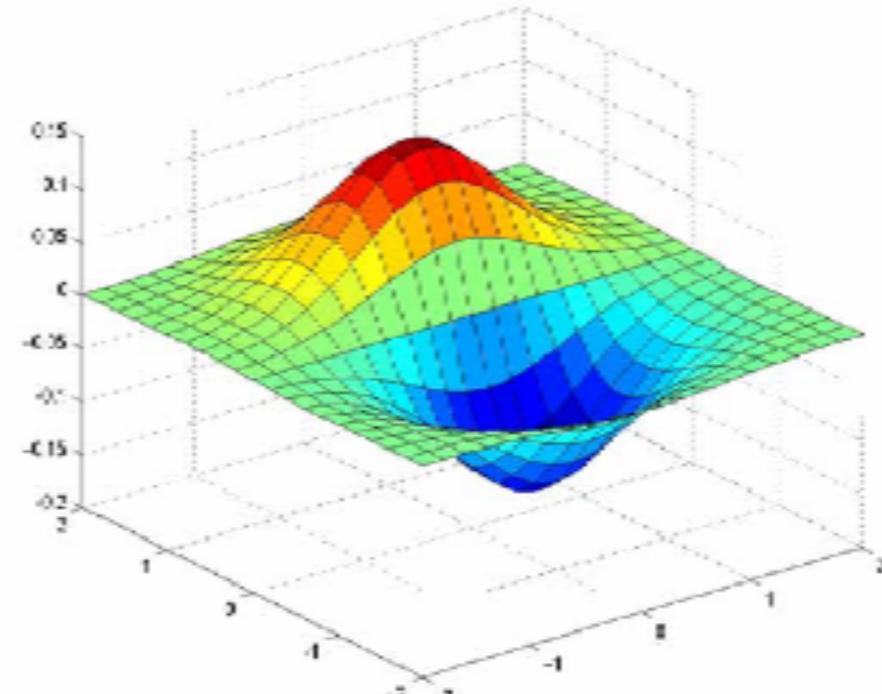
$(\frac{\partial}{\partial x}h) \star f$



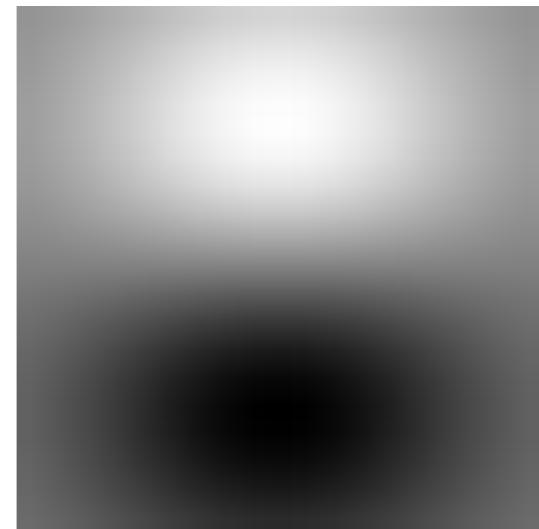
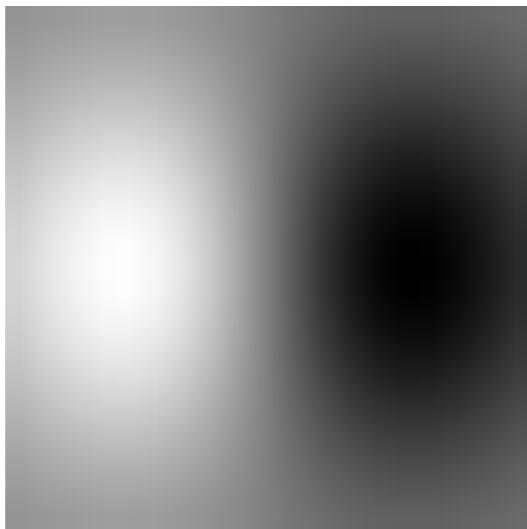
Derivative of Gaussian filters



x-direction



y-direction



Effect of σ on derivatives



$\sigma = 1$ pixel

$\sigma = 3$ pixels

The apparent structures differ depending on Gaussian's scale parameter.

Larger values: larger scale edges detected

Smaller values: finer features detected

So, what scale to choose?

It depends what we're looking for.



FOTOSSEARCH.com

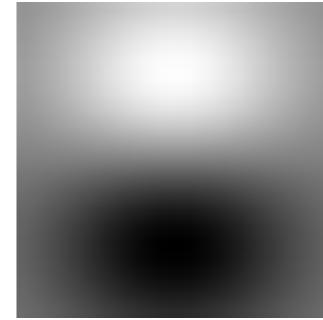
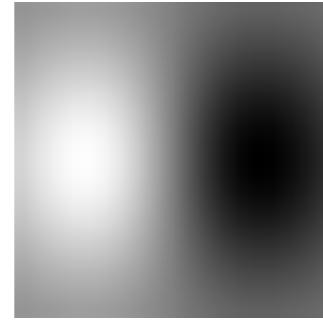


Too fine of a scale...can't see the forest for the trees.
Too coarse of a scale...can't tell the maple grain from the cherry

Template matching

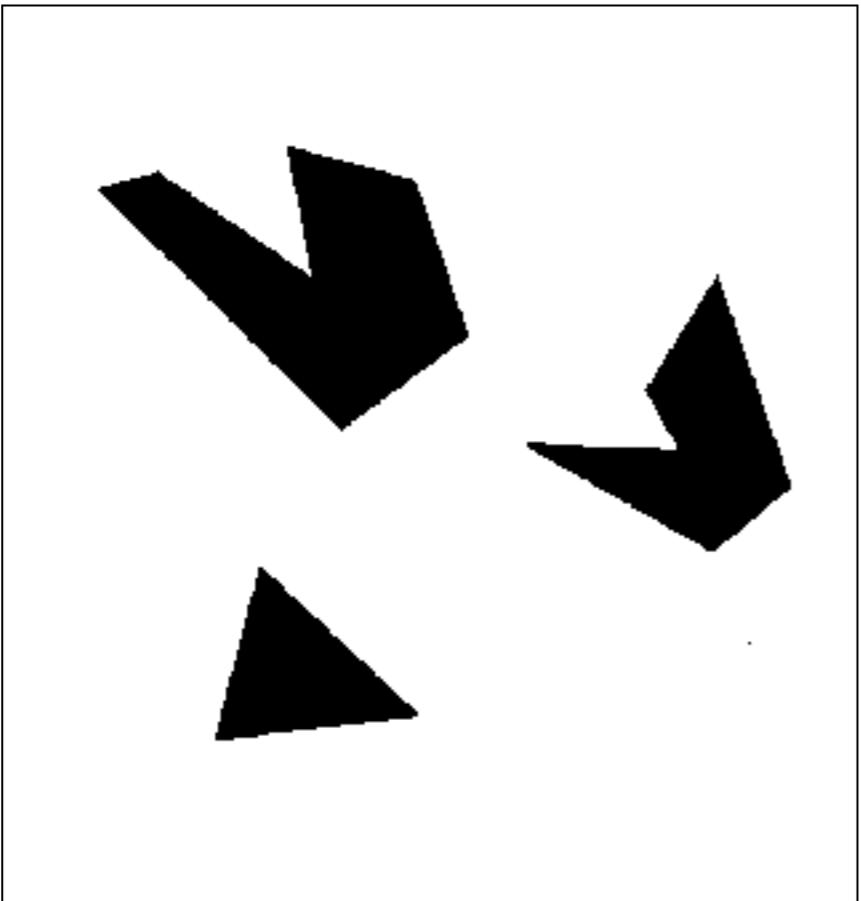
- Filters as **templates**:

Note that filters look like the effects they are intended to find
--- “matched filters”

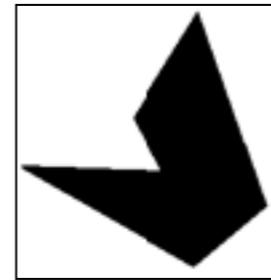


- Use normalized cross-correlation score to find a given pattern (template) in the image.
 - Szeliski Eq. 8.11
 - Normalization needed to control for relative brightnesses.

Template matching



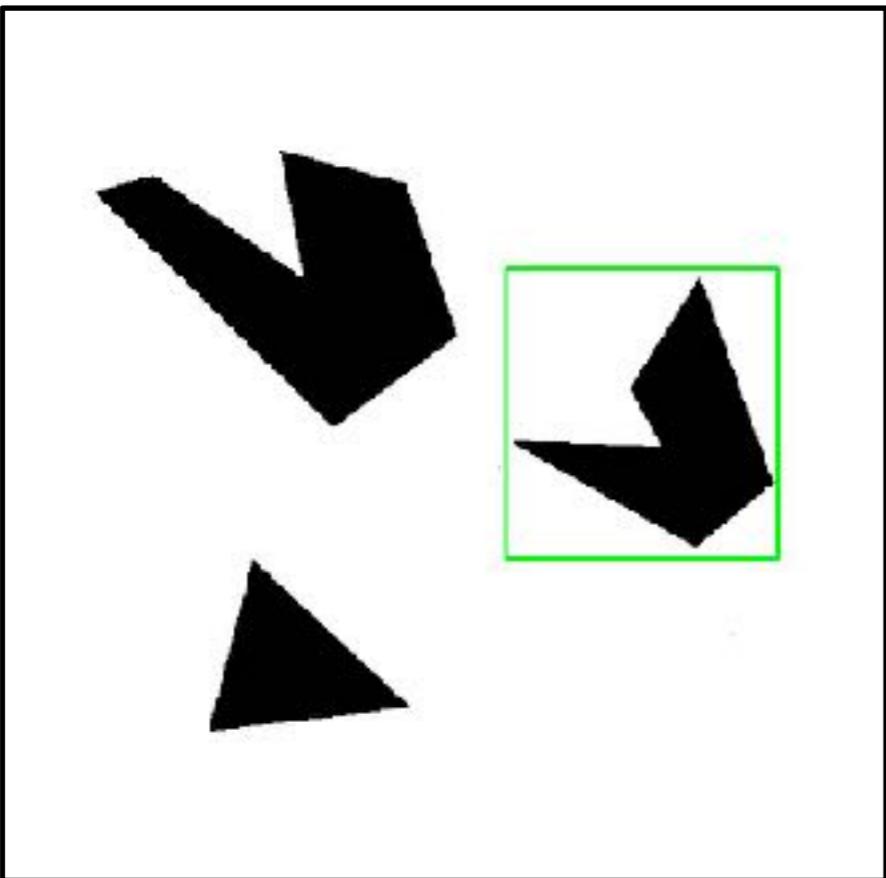
Scene



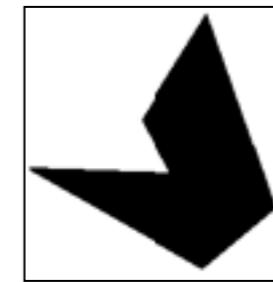
Template (mask)

A toy example

Template matching

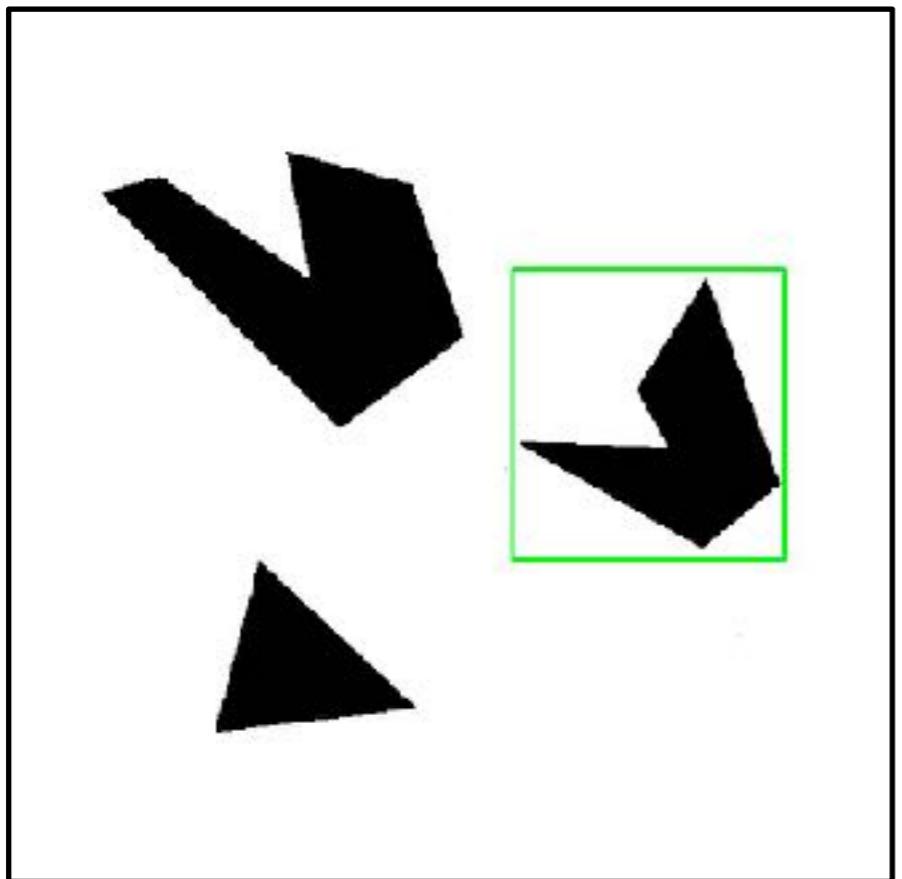


Detected template

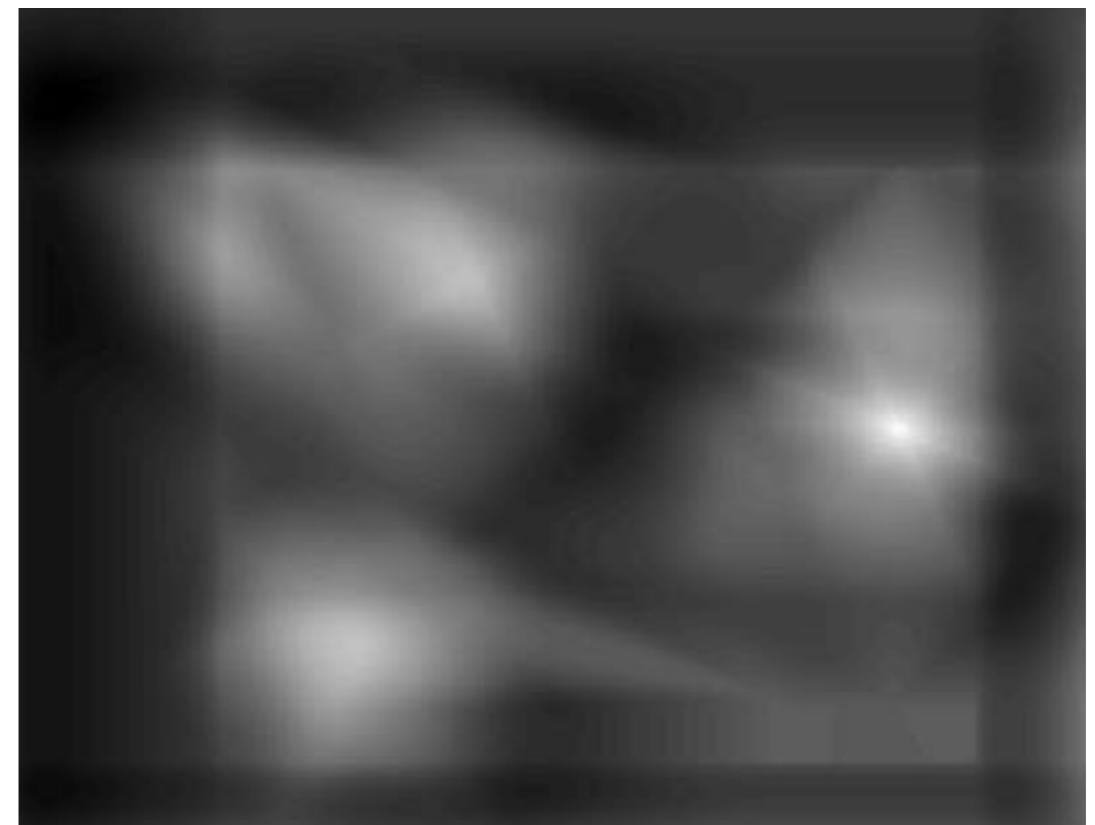


Template

Template matching



Detected template



Correlation map

Template matching

$h =$

0	1	2
1	4	1
0	1	0

Template (mask)

$f =$

5	3	1	5	5
5	2	0	1	2
1	2	1	4	1
1	5	0	1	0
1	4	5	4	2

Scene

f_{cut}

1	5	5
0	1	2
1	4	1

$$r = \|f_{cut} - h\|^2$$

Template matching

$$r = \|f_{cut} - h\|^2$$

$$r = (f_{cut} - h) \cdot (f_{cut} - h) = f_{cut} \cdot f_{cut} - 2f_{cut} \cdot h + h \cdot h$$

$$(f^2) * 1 - 2f * \check{h} + \sum h^2$$

Template matching

$$(f^2) * \mathbf{1} - 2f * \check{h} + \sum h^2$$

```
e = ones(size(h))
hnorm2 = norm(h, 'fro')^2;
hhat = flipud(fliplr(h))
```

```
h =
 0   1   2
 1   4   1
 0   1   0
f =
 5   3   1   5   5
 5   2   0   1   2
 1   2   1   4   1
 1   5   0   1   0
 1   4   5   4   2
```

```
e =
 1   1   1
 1   1   1
 1   1   1
```

```
hhat =
 0   1   0
 1   4   1
 2   1   0
r = conv2(f.^2,e, 'same') - 2*conv2(f,hhat, 'same') + hnorm2
```

```
r =
 31.0000  48.0000  40.0000  26.0000  25.0000
 24.0000  54.0000  55.0000  48.0000  66.0000
 52.0000  51.0000  52.0000 -0.0000  27.0000
 42.0000  40.0000  88.0000  60.0000  54.0000
 29.0000  38.0000  47.0000  22.0000  21.0000
```

Where's Waldo?

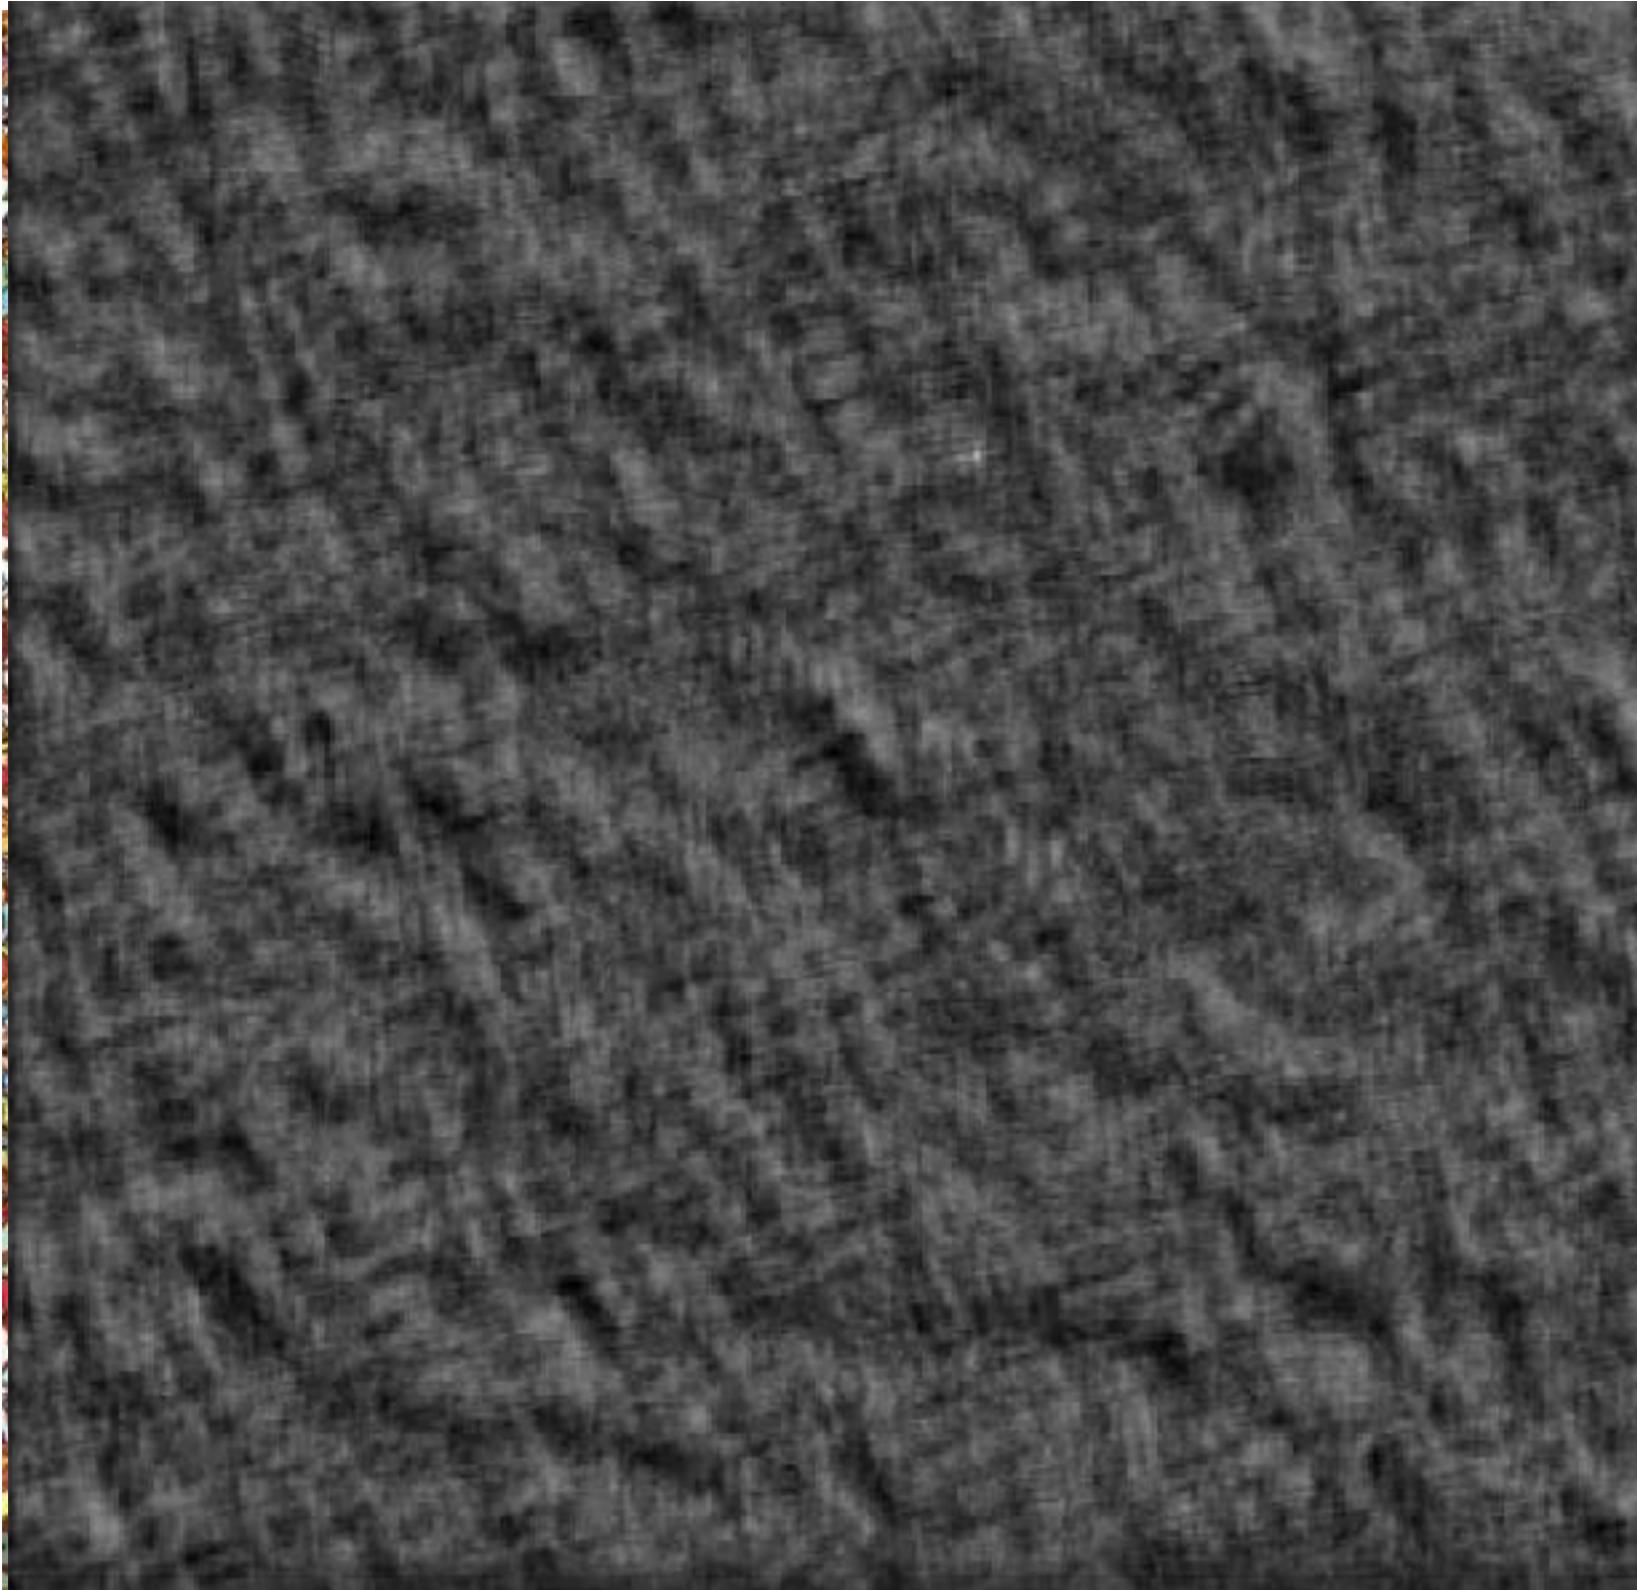


Scene



Template

Where's Waldo?



Scene

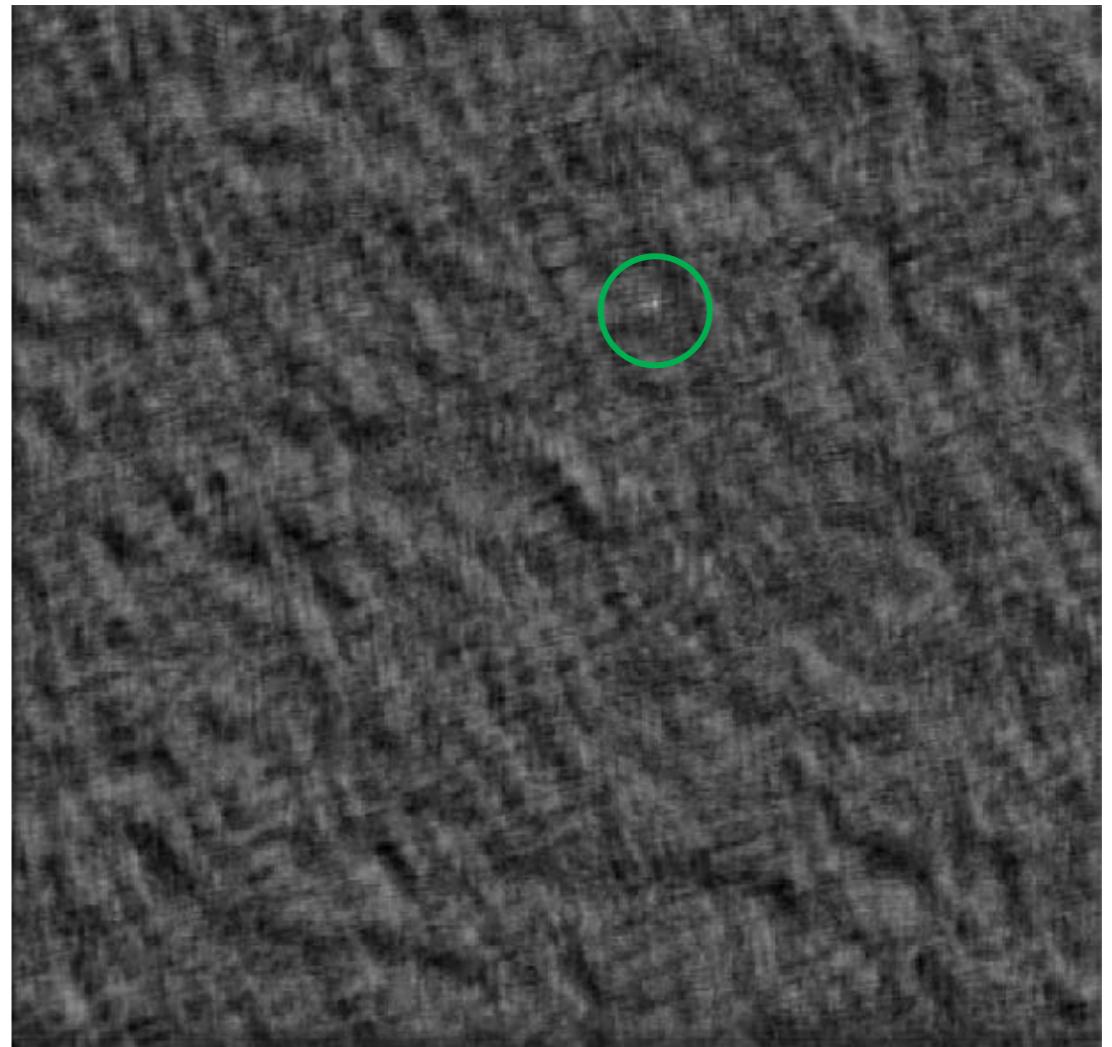


Template

Where's Waldo?



Detected template



Correlation map

Template matching

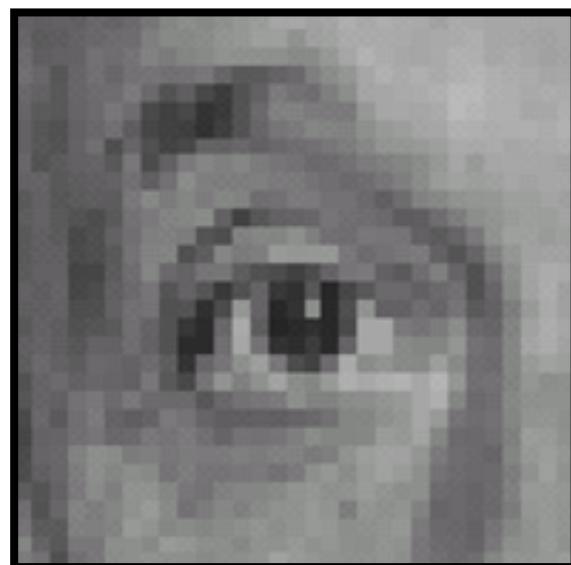


Scene

What if the template is not identical to some
subimage in the scene?

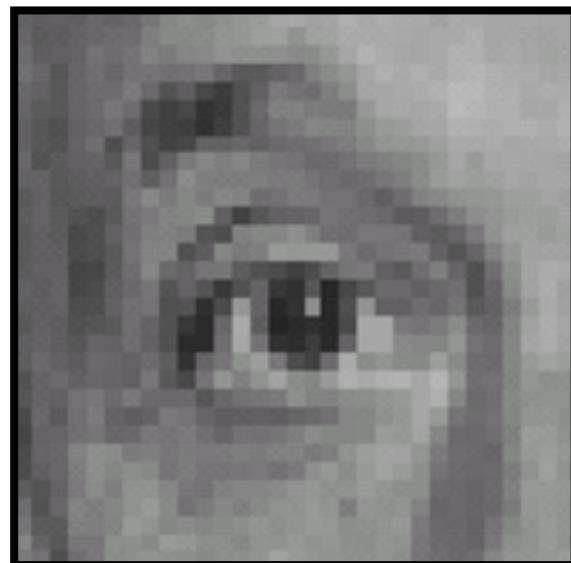
Template

Practice with linear filters



$$\text{Original} \quad * \quad \begin{array}{|c|c|c|} \hline 0 & 0 & 0 \\ \hline 0 & 1 & 0 \\ \hline 0 & 0 & 0 \\ \hline \end{array} = ?$$

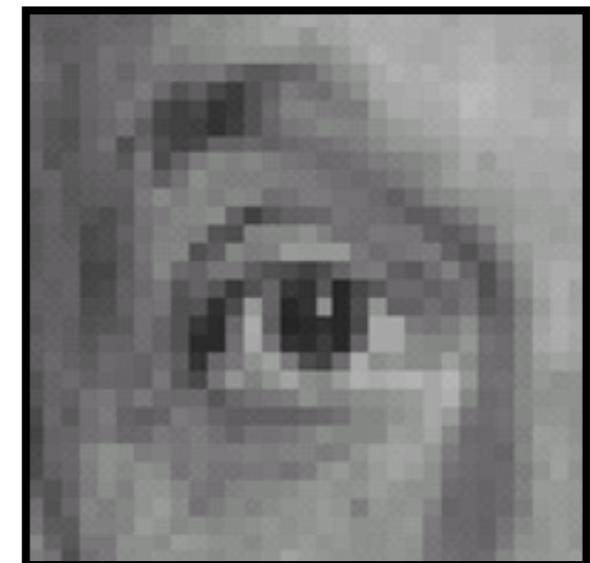
Practice with linear filters



*

0	0	0
0	1	0
0	0	0

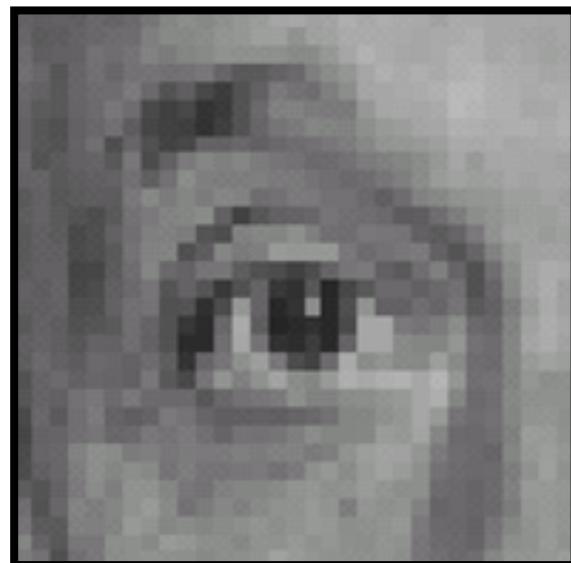
=



Original

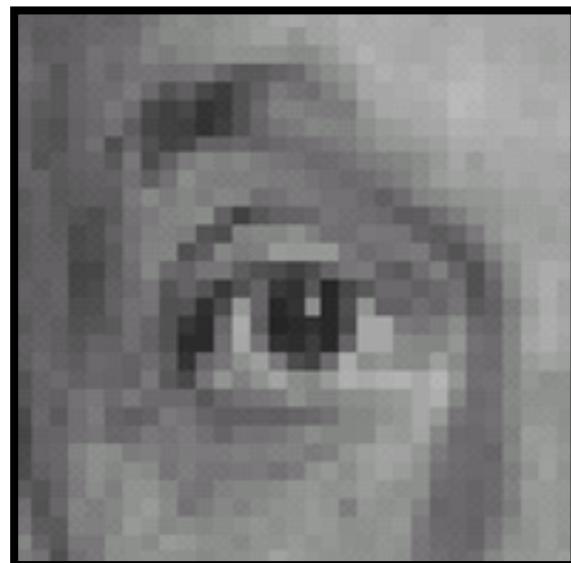
Filtered
(no change)

Practice with linear filters



$$\text{Original} \quad * \quad \begin{array}{|c|c|c|} \hline 0 & 0 & 0 \\ \hline 0 & 0 & 1 \\ \hline 0 & 0 & 0 \\ \hline \end{array} \quad = \quad ?$$

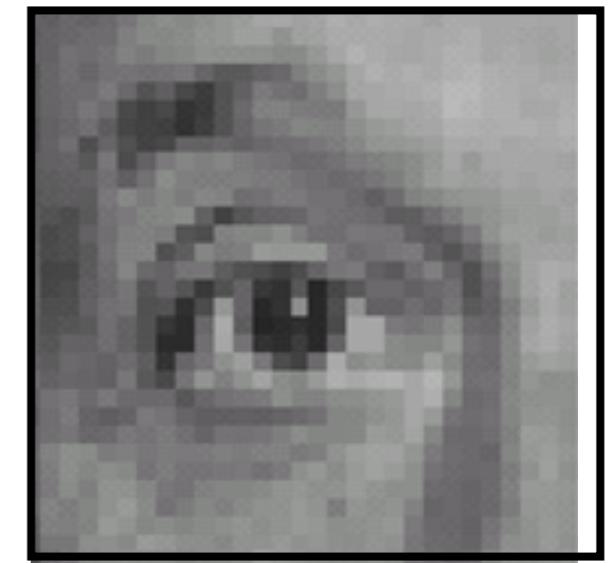
Practice with linear filters



*

0	0	0
0	0	1
0	0	0

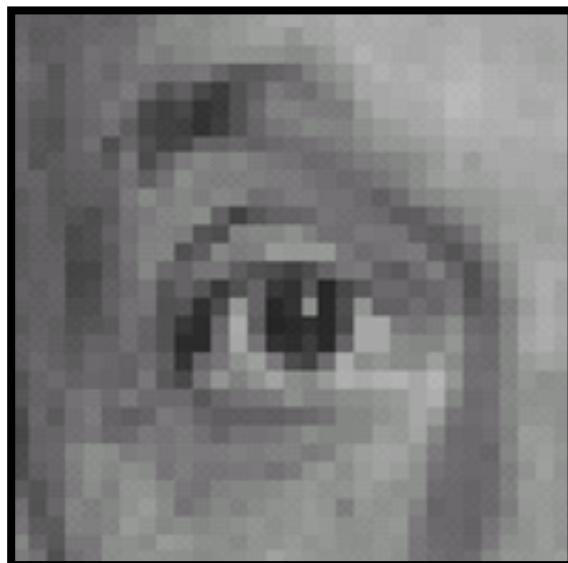
=



Original

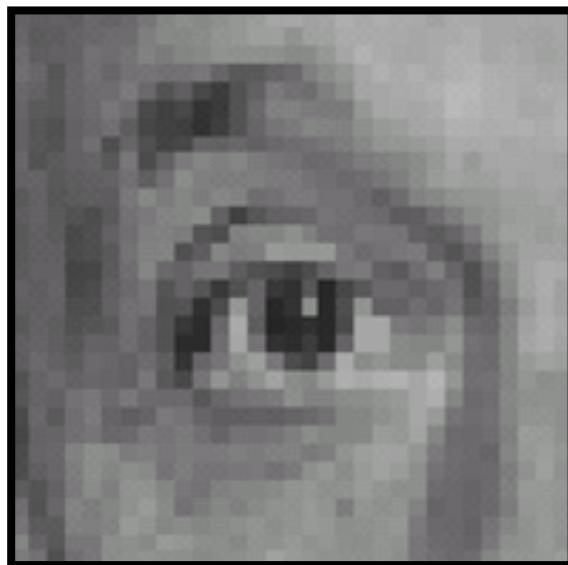
Shifted left
by 1 pixel with
correlation

Practice with linear filters

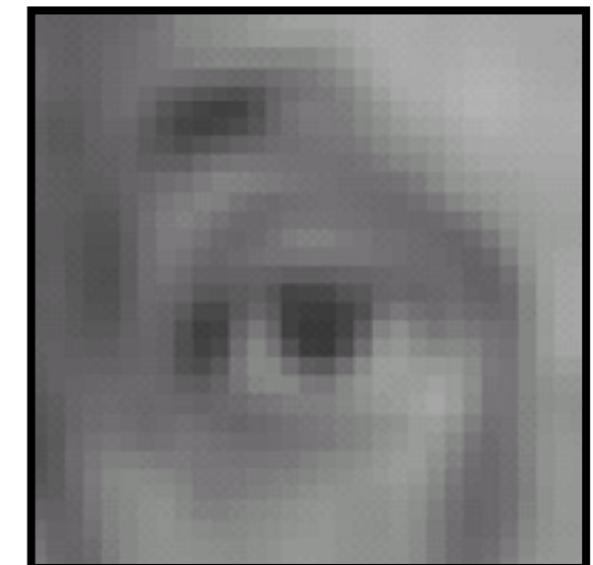


$$\text{Original} \quad * \quad \frac{1}{9} \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array} = ?$$

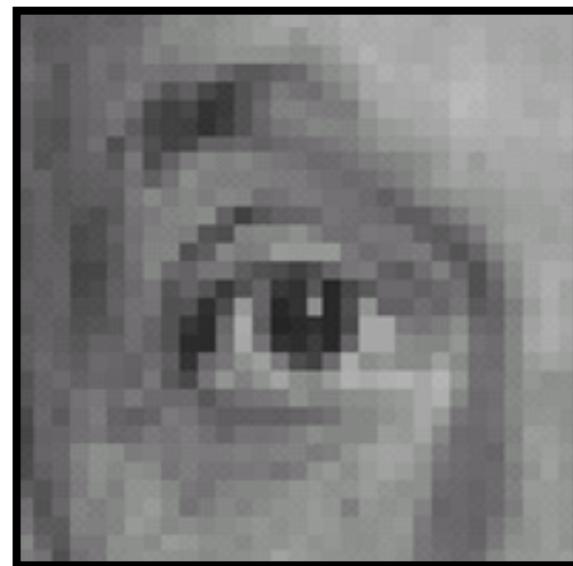
Practice with linear filters



$$\text{Original} \quad * \quad \frac{1}{9} \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array} = \quad \text{Blur (with a box filter)}$$

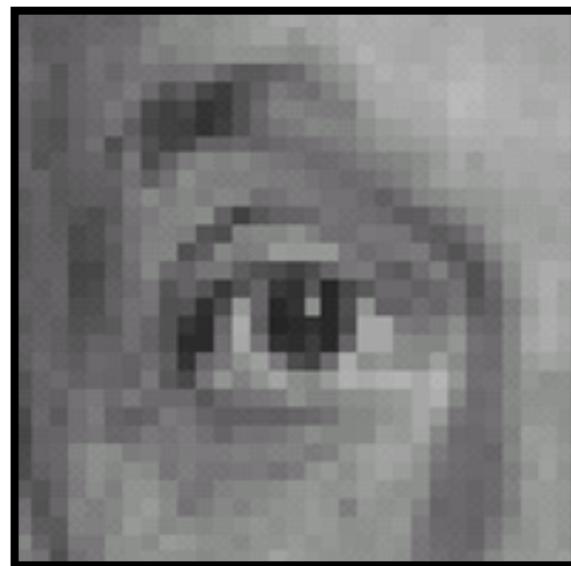


Practice with linear filters



$$\text{Original} \quad * \quad \begin{matrix} 0 & 0 & 0 \\ 0 & 2 & 0 \\ 0 & 0 & 0 \end{matrix} - \frac{1}{9} \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix} = ?$$

Practice with linear filters



*

$$\begin{matrix} 0 & 0 & 0 \\ 0 & 2 & 0 \\ 0 & 0 & 0 \end{matrix}$$

$$-\frac{1}{9}$$

$$\begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$$

=

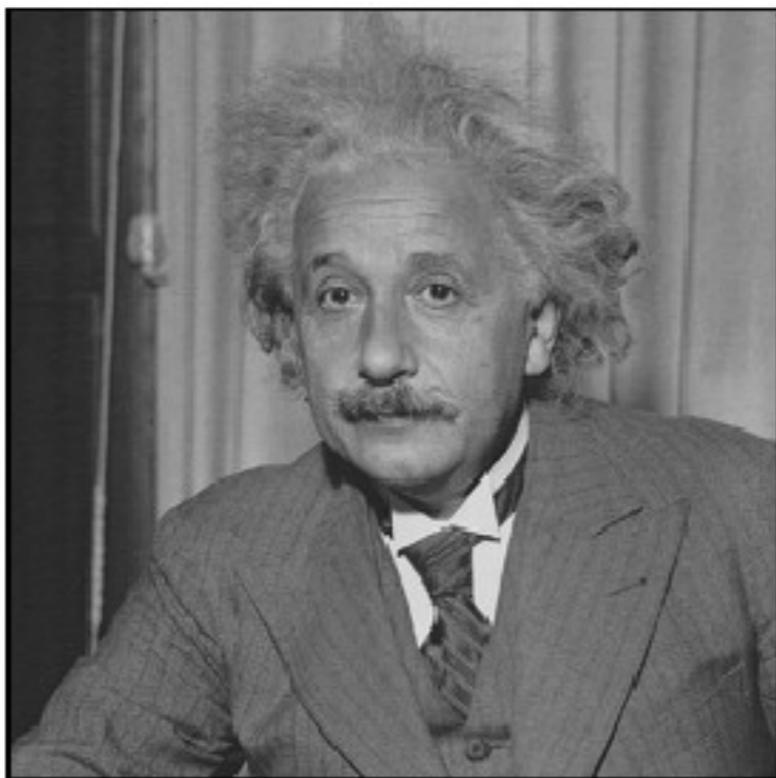


Original

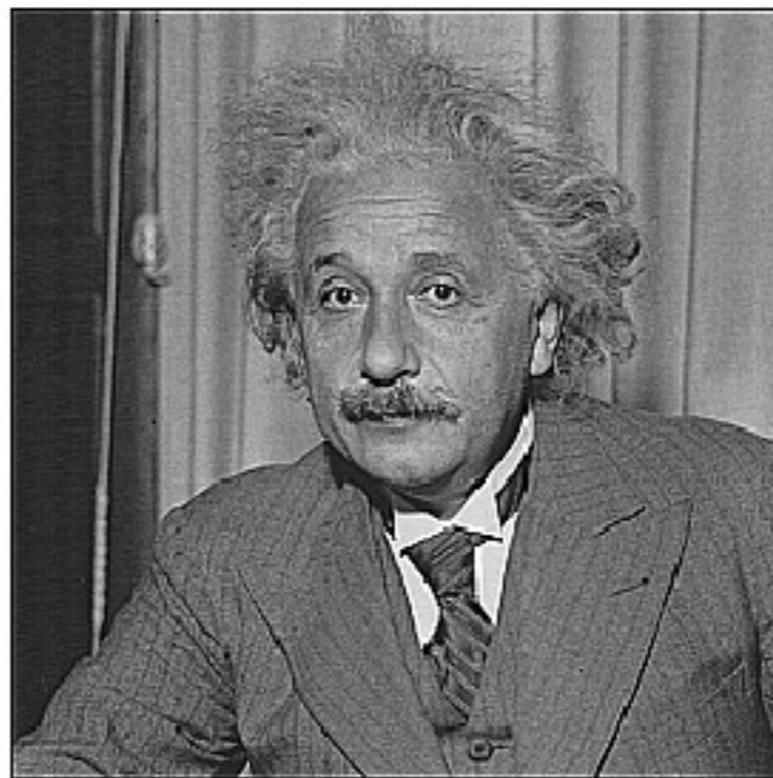
Sharpening filter

- Accentuates differences with local average

Filtering examples: sharpening



before



after

Convolutions and the Fourier Transform

- What to do near the edge of the image?
- Understanding the Fourier Transform
- The convolution Theorem
- Understanding Convolutions using the Fourier Transform

What to do near the edge of the image?

It is important to know where position zero in sequences are.
Compare with decimal-comma in reals.

$$12 \cdot 12 = 144, 1.2 \cdot 12 = 14.4$$

What notations should we use?

Suggestion: underline the element at position zero, e.g.

$$<1,\underline{1}> * <\underline{1},2,1> = <1,3,\underline{3},1>$$

but

$$<1,\underline{1}> * <\underline{1},2,1> = <\underline{1},\underline{3},3,1>$$

and

$$<\underline{1},1> * <\underline{1},2,1> = <\underline{1},3,3,1>$$

What to do near the edge of the image?

In practice we do not have infinite images.

How should we treat the edges of the image? What values should one assume 'outside' the image.

Some common choices are

1. Only calculate the result where we can be certain. The result is a smaller image.
2. Assume that there are zeros outside the image. This often means that we introduce artificial sharp edges at the border.
3. Make a periodic expansion of the image, i.e. assume that the image is periodic. This fits well with the theory for discrete fourier transform.

What to do near the edge of the image?

Assume that one would like to convolute the image

$$f = \begin{bmatrix} 1 & 2 & 3 & 5 \\ 1 & 3 & 2 & 1 \\ 2 & 2 & 2 & 2 \end{bmatrix}$$

with the smoothing filter

$$h = \begin{bmatrix} 1 & 1 \\ 1 & 1 \end{bmatrix}$$

What to do near the edge of the image?

(1) Don't let h extend outside f

$$\begin{bmatrix} 7 & 10 & 11 \\ 8 & 9 & 7 \end{bmatrix}$$

(2) Extend with zeros \Rightarrow equal or larger resulting $h * f$ -image

$$\begin{bmatrix} 1 & 3 & 5 & 8 & 5 \\ 2 & 7 & 10 & 11 & 6 \\ 3 & 8 & 9 & 7 & 3 \\ 2 & 4 & 4 & 4 & 2 \end{bmatrix}$$

What to do near the edge of the image?

- (3) Extend f and h to periodic functions with the same period:
 $f_p, h_p \Rightarrow$ periodic $h_p * f_p$ result with same period

$$\begin{bmatrix} 10 & 7 & 9 & 12 \\ 8 & 7 & 10 & 11 \\ 6 & 8 & 9 & 7 \end{bmatrix}$$

Here we have also made a periodic function of h :

$$h = \begin{bmatrix} 1 & 1 & 0 & 0 \\ 1 & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 \end{bmatrix} .$$

Fourier Transform

$$F(u, v) = \sum_{x=1}^M \sum_{y=1}^N f(x, y) e^{-2\pi i (u-1)(x-1)/M + (v-1)(y-1)/N}$$

- Can be viewed as a change of basis
- Image $f \rightarrow$ Fourier Transform F (and back)
- Has strong connections with convolutions
- Useful for image compression
- Useful for image understanding
- Basically a great tool

Fourier Transform

- Definition, is a change of basis, what does it mean
 - Detour (for increased understanding)
 - Ordinary Fourier Transform (from previous courses)
 - Examples
 - Properties
- Discrete Fourier Transform – 1D

Compare with ordinary Fourier Transform

Definition

Let f be a function from \mathbb{R} to \mathbb{R} . The Fourier transformen of f is defined as

$$(\mathcal{F}f)(u) = F(u) = \int_{-\infty}^{+\infty} e^{-i2\pi xu} f(x) dx .$$

■

Theorem

Under the right assumptions on f , the following inversion formula

$$f(x) = \int_{-\infty}^{+\infty} e^{i2\pi ux} F(u) du$$

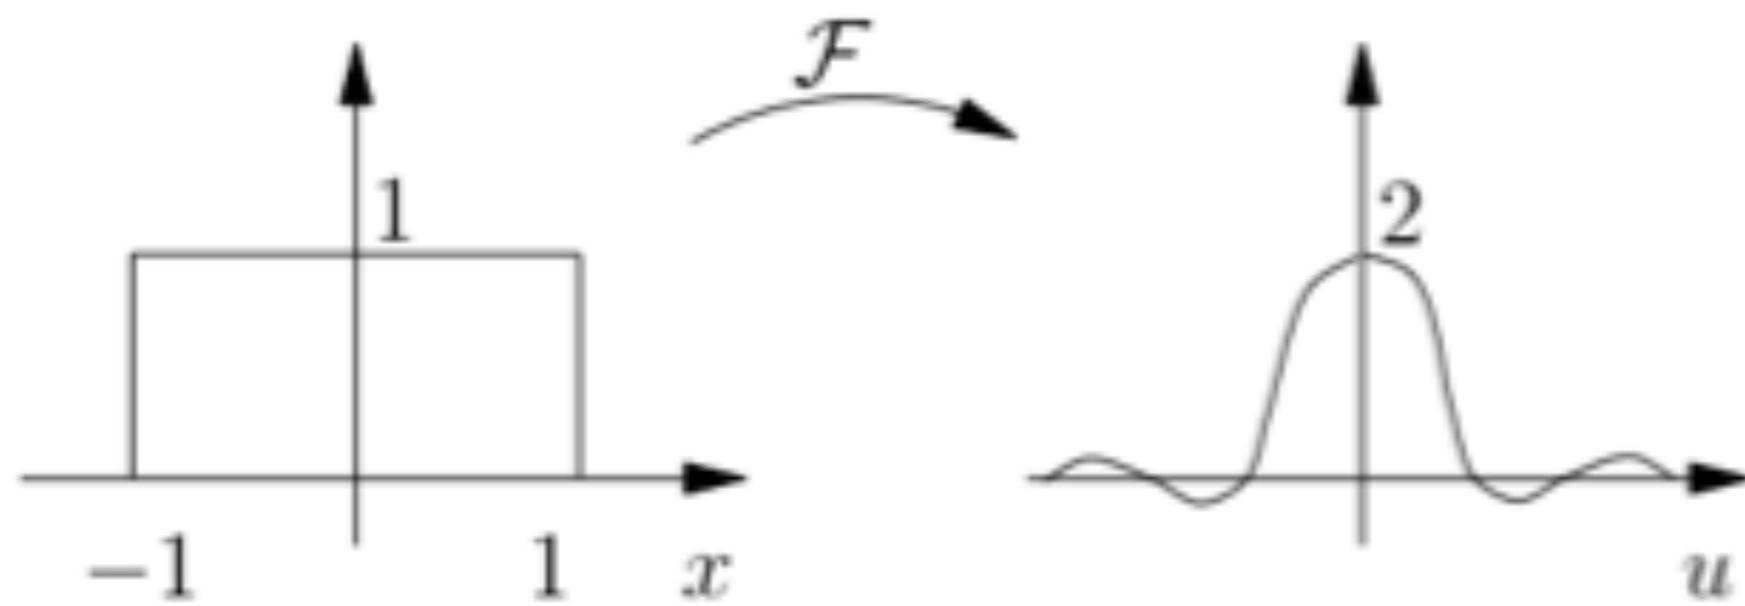
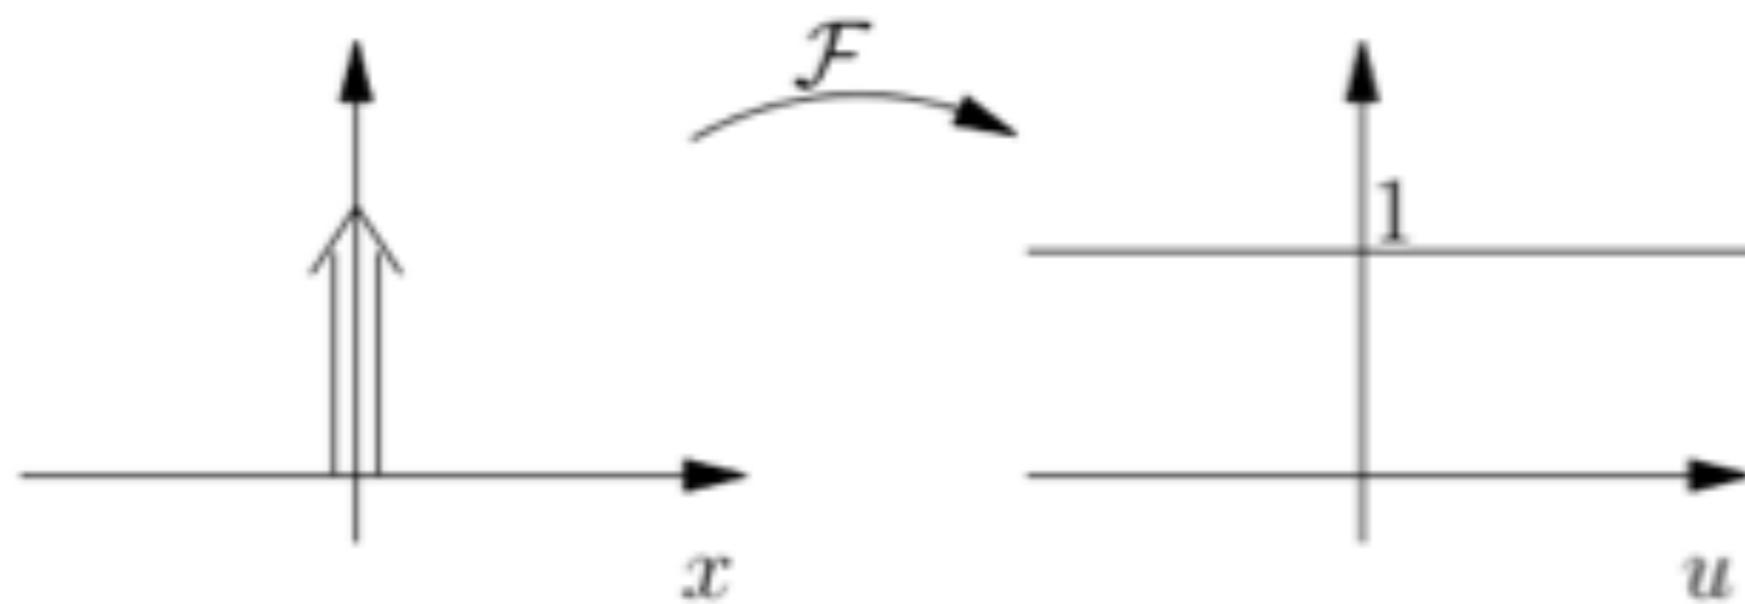
holds.

Examples

$$\delta(x) \mapsto 1(u)$$

$$\text{rect}(x) \mapsto 2 \frac{\sin(2\pi u)}{2\pi u} = 2 \text{sinc}(2\pi u)$$

Examples



Examples

$$c_1 f_1(x) + c_2 f_2(x) \mapsto c_1 F_1(u) + c_2 F_2(u) \text{ (linearity)}$$

$$f(\lambda x) \mapsto \frac{1}{|\lambda|} F\left(\frac{u}{\lambda}\right) \quad \text{(scaling)}$$

$$f(x - a) \mapsto e^{-i2\pi u a} F(u) \quad \text{(translation)}$$

$$e^{-i2\pi x a} f(x) \mapsto F(u + a) \quad \text{(modulation)}$$

$$\overline{f(x)} \mapsto \overline{F(-u)} \quad \text{(conjugation)}$$

$$\frac{df}{dx} \mapsto 2\pi i u F(u) \quad \text{(differentiation I)}$$

$$-2\pi i x f(x) \mapsto \frac{dF}{du} \quad \text{(differentiation II)}$$

Example: $\delta(x - 1) \mapsto e^{-i2\pi u}$

Discrete Fourier Transform - 1D

$$f = \begin{bmatrix} f(1) \\ \vdots \\ f(N) \end{bmatrix}$$

$$F(u) = \sum_{x=1}^N f(x) e^{-2\pi i (u-1)(x-1)/N}$$

$$F(u) = \sum_{x=1}^N f(x) \omega_N^{(u-1)(x-1)}, \quad \omega_N = e^{-2\pi i / N}$$

Discrete Fourier Transform - 1D

Definition

The Fourier Matrix \mathcal{F}_N is given by

$$\mathcal{F}_N = \begin{pmatrix} 1 & 1 & 1 & \dots & 1 \\ 1 & \omega_N & \omega_N^2 & \dots & \omega_N^{N-1} \\ 1 & \omega_N^2 & \omega_N^4 & \dots & \omega_N^{2(N-1)} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 1 & \omega_N^{N-1} & \omega_N^{2(N-1)} & \dots & \omega_N^{(N-1)(N-1)} \end{pmatrix}.$$

$$f \rightarrow F = \mathcal{F}_N f$$

Discrete Fourier Transform - 1D

Theorem 3.3.1. *For the Fourier matrix the following holds,*

$$\mathcal{F}\overline{\mathcal{F}} = NI .$$

From this we obtain $\mathcal{F}^{-1} = \frac{1}{N}\overline{\mathcal{F}}$. The inverse Fourier transform is thus

$$f = \overline{\mathcal{F}}F \iff f(x) = \frac{1}{N} \sum_{u=1}^N F(u) \omega_N^{(x-1)(u-1)}, \quad x = 1, \dots, N .$$

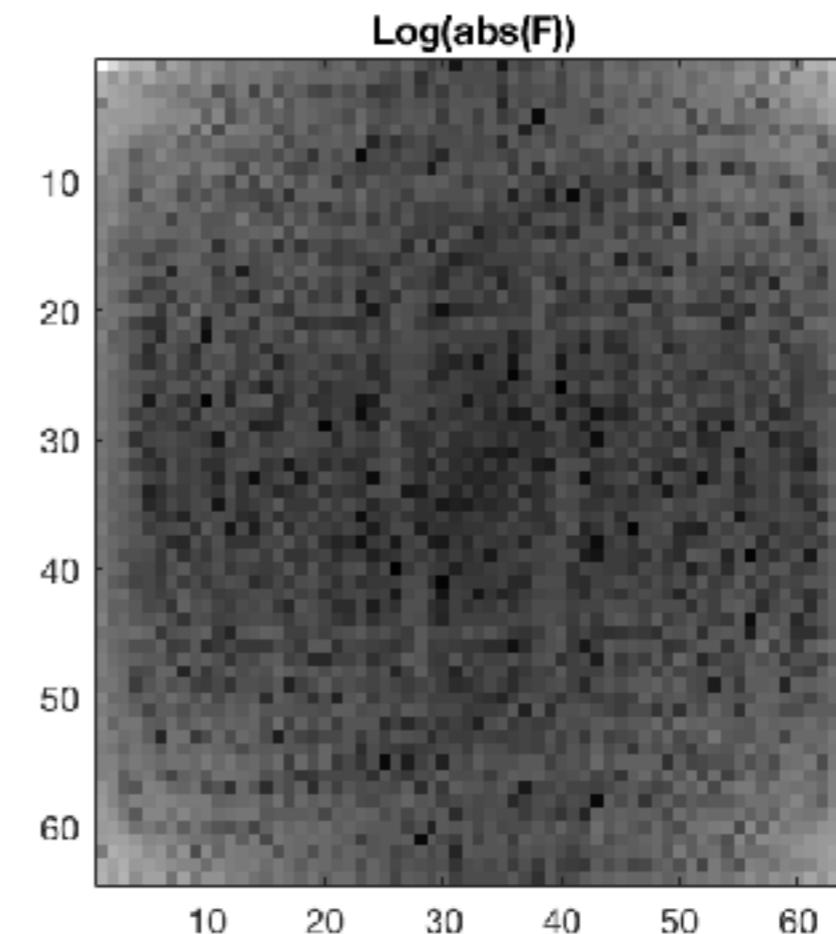
Discrete Fourier Transform - 1D

- Important: DFT assumes that signals are periodic!
- Think of the signal as wrapped periodically
- Fourier transform is complex.
- Plot absolute value and phase
- Low frequencies in the edges/corners.
- Ordinary images typically have large values for low frequencies.

Discrete Fourier Transform - 2D

$$F(u, v) = \sum_{x=1}^M \sum_{y=1}^N f(x, y) e^{-i2\pi((u-1)(x-1)/M + (v-1)(y-1)/N)}$$

$$f(x, y) = \frac{1}{MN} \sum_{u=1}^M \sum_{v=1}^N F(u, v) e^{i2\pi((u-1)(x-1)/M + (v-1)(y-1)/N)}$$



Discrete Fourier Transform - 2D

Let the matrix F represent the Fourier transform of the image $f(x, y)$:

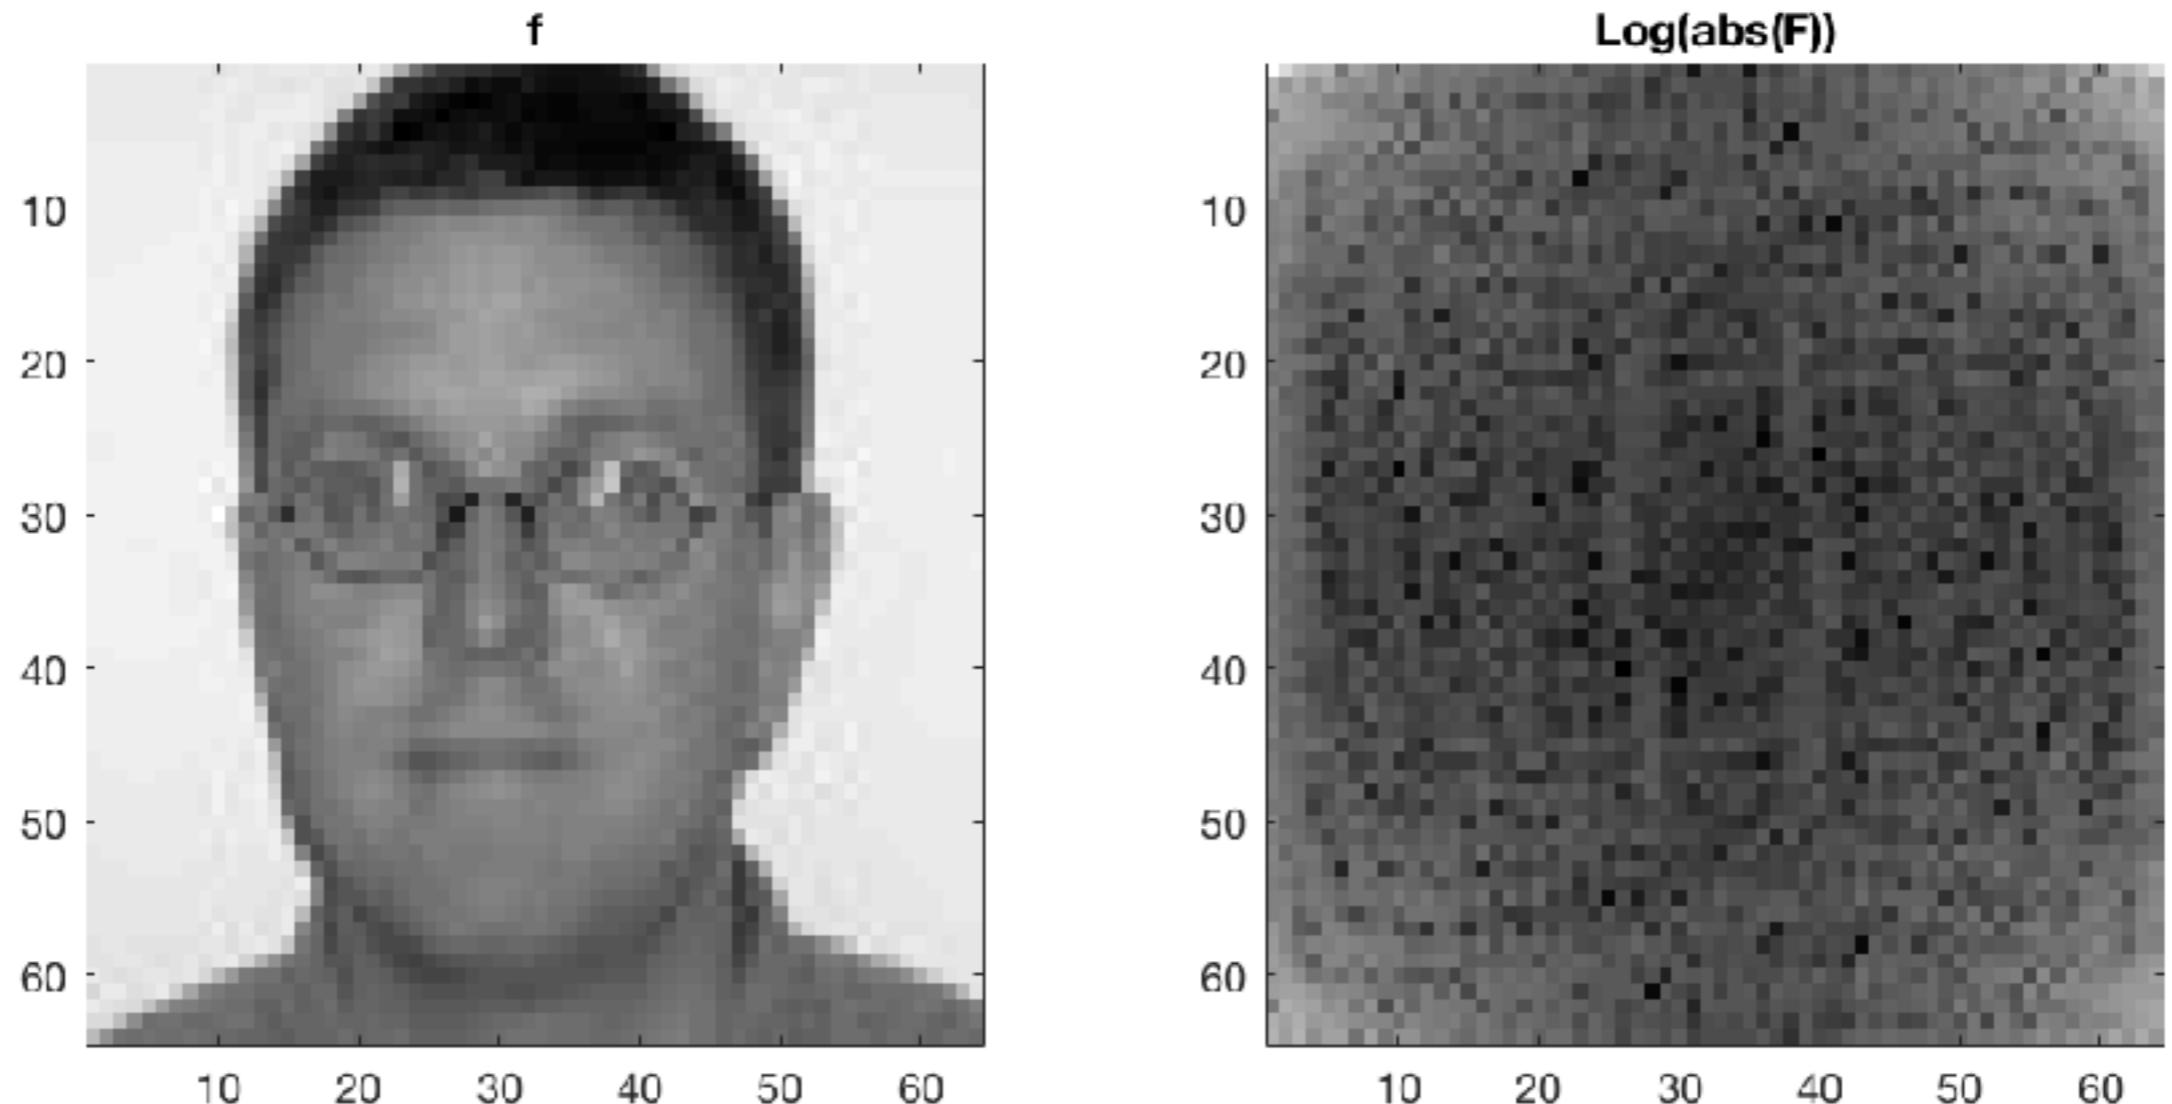
$$F = \mathcal{F}_M f \mathcal{F}_N$$

or

$$F = \mathcal{F}_M (\mathcal{F}_N f^T)^T .$$

i.e. the DFT in two dimensions can be calculated by repeated use of the one-dimensional DFT, first for the rows, then for the columns.

Discrete Fourier Transform - 2D Example

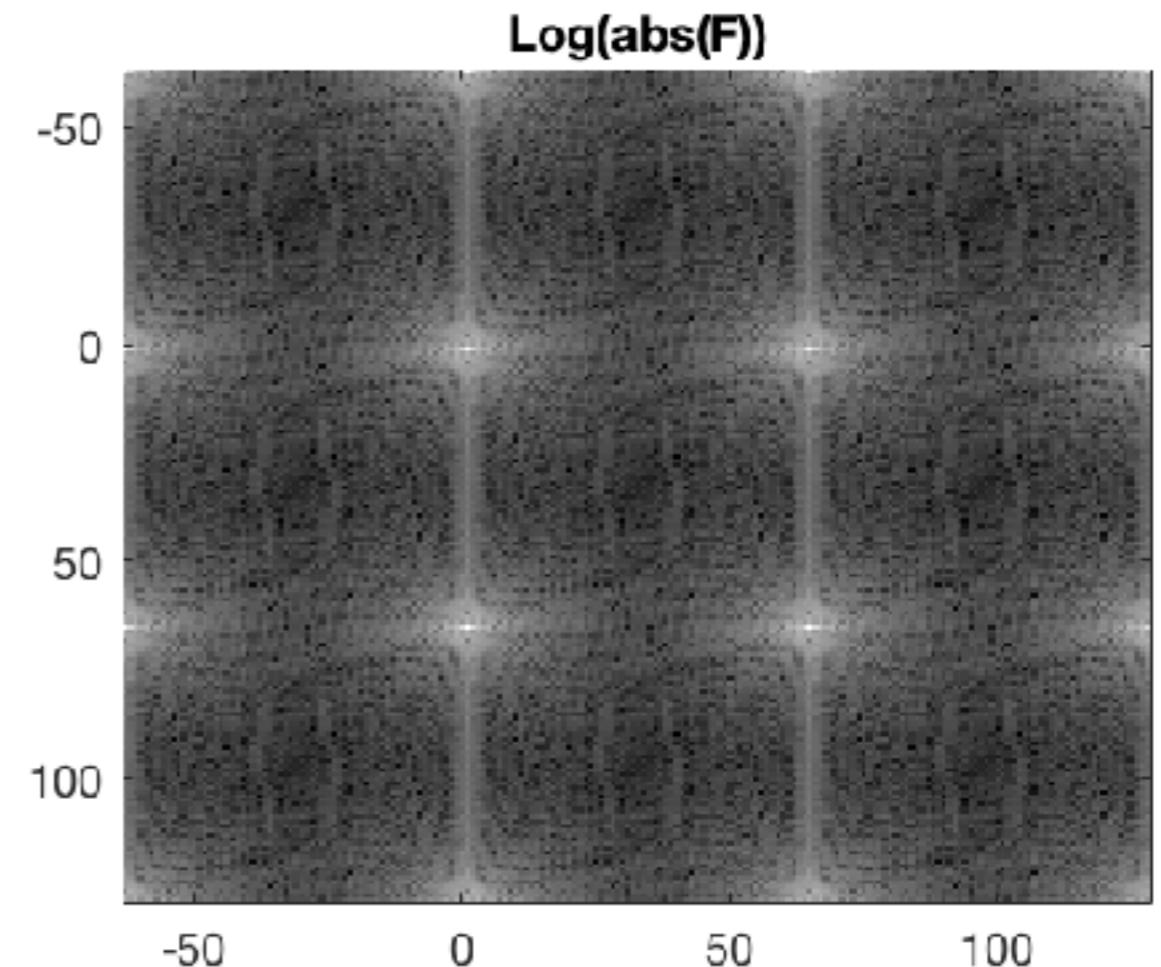
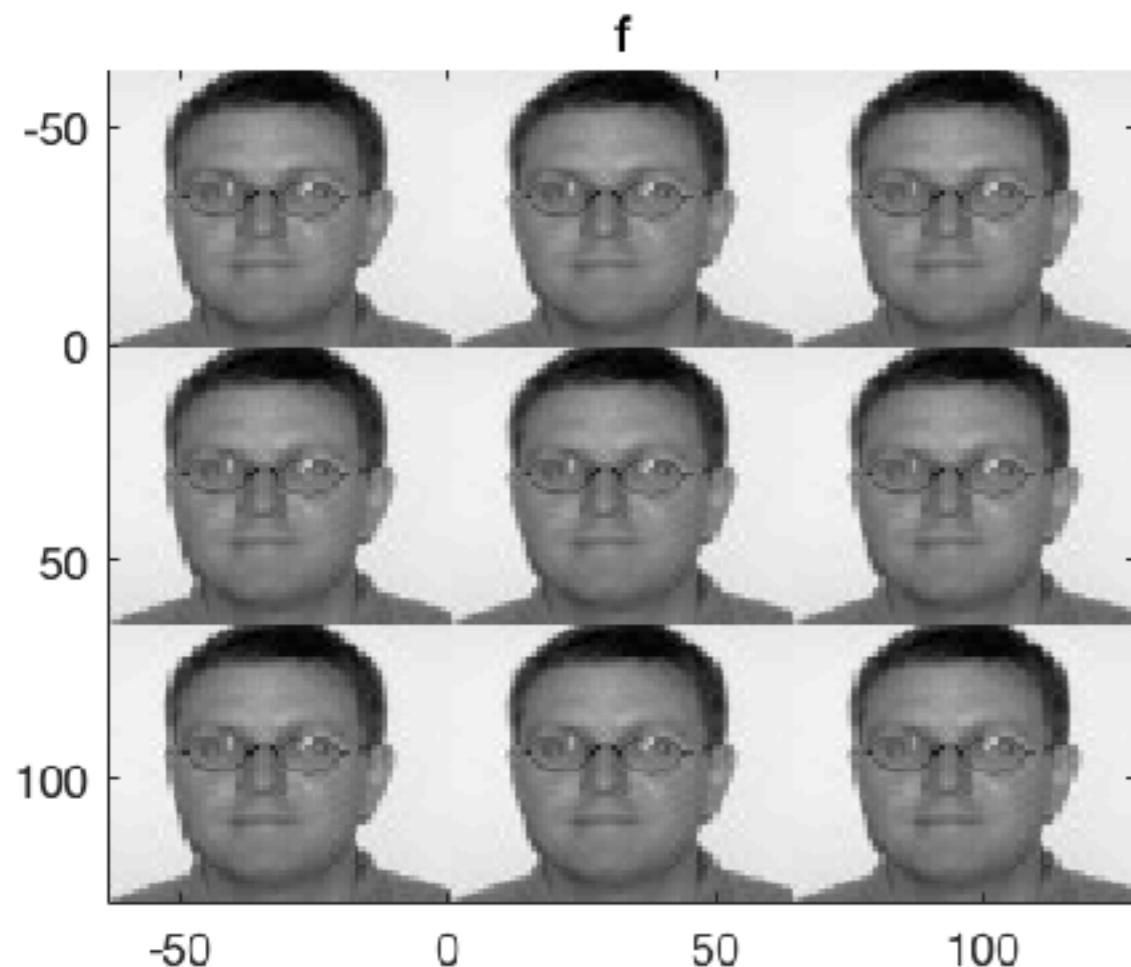


Low frequencies in the edges/corners

Fourier transform is complex. Plot absolute value and phase

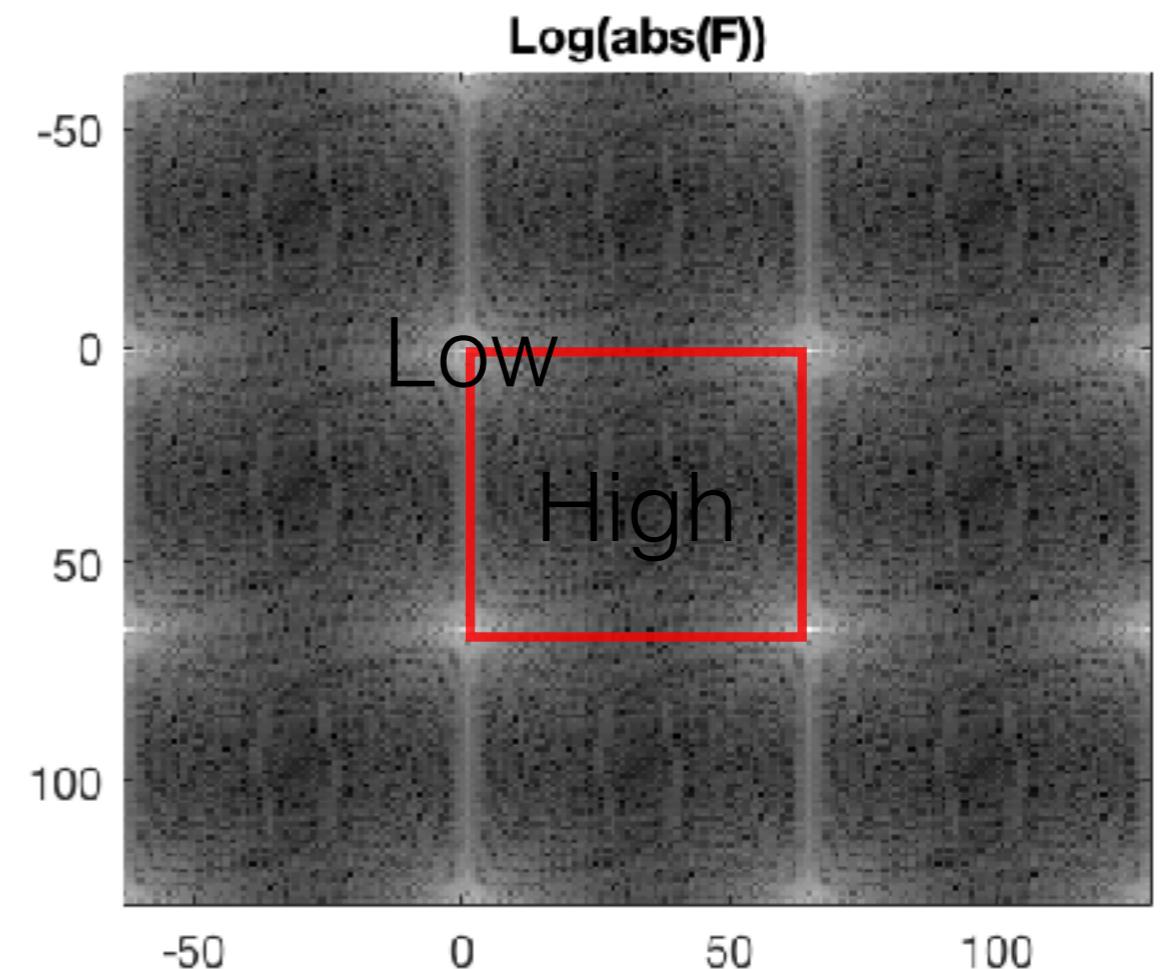
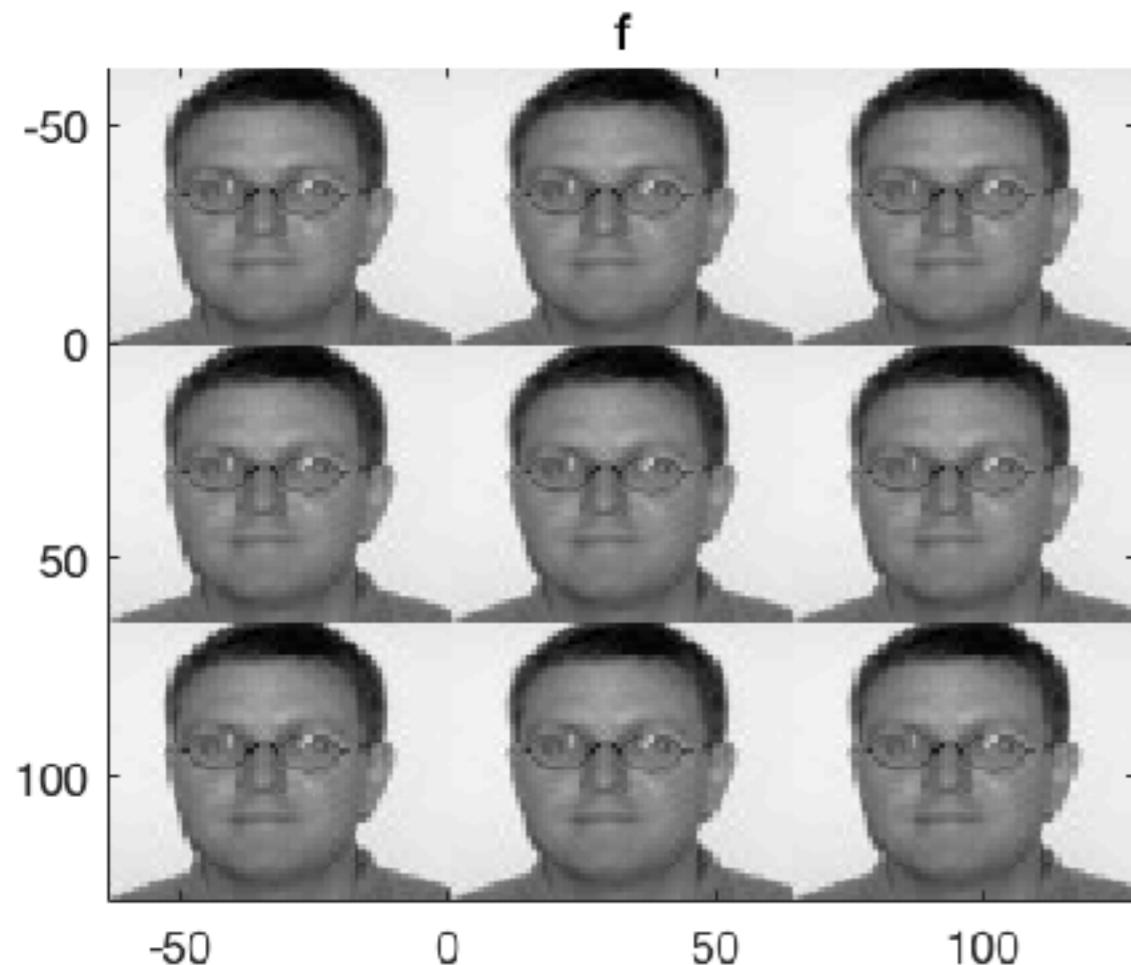
Discrete Fourier Transform - 2D

Example – Periodic expansion



Discrete Fourier Transform - 2D

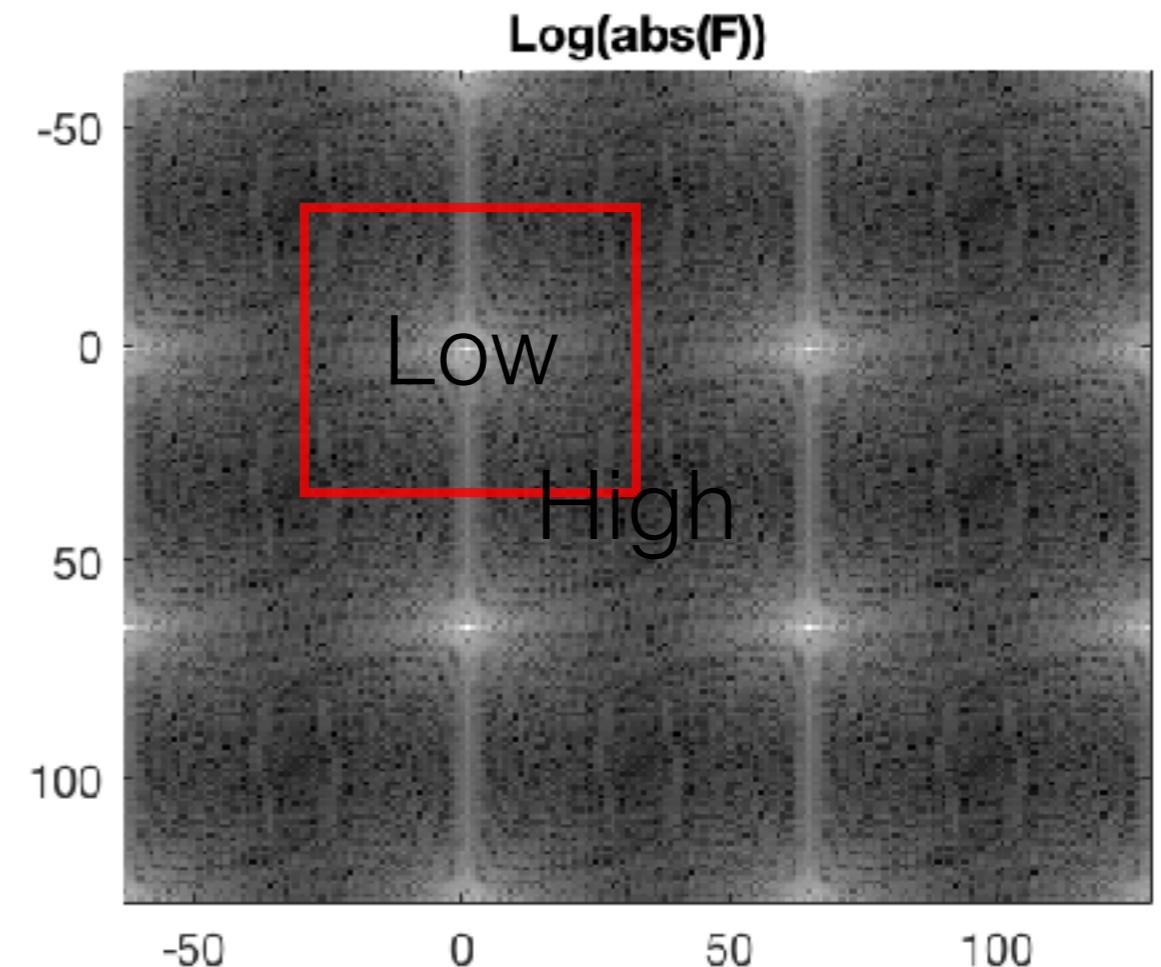
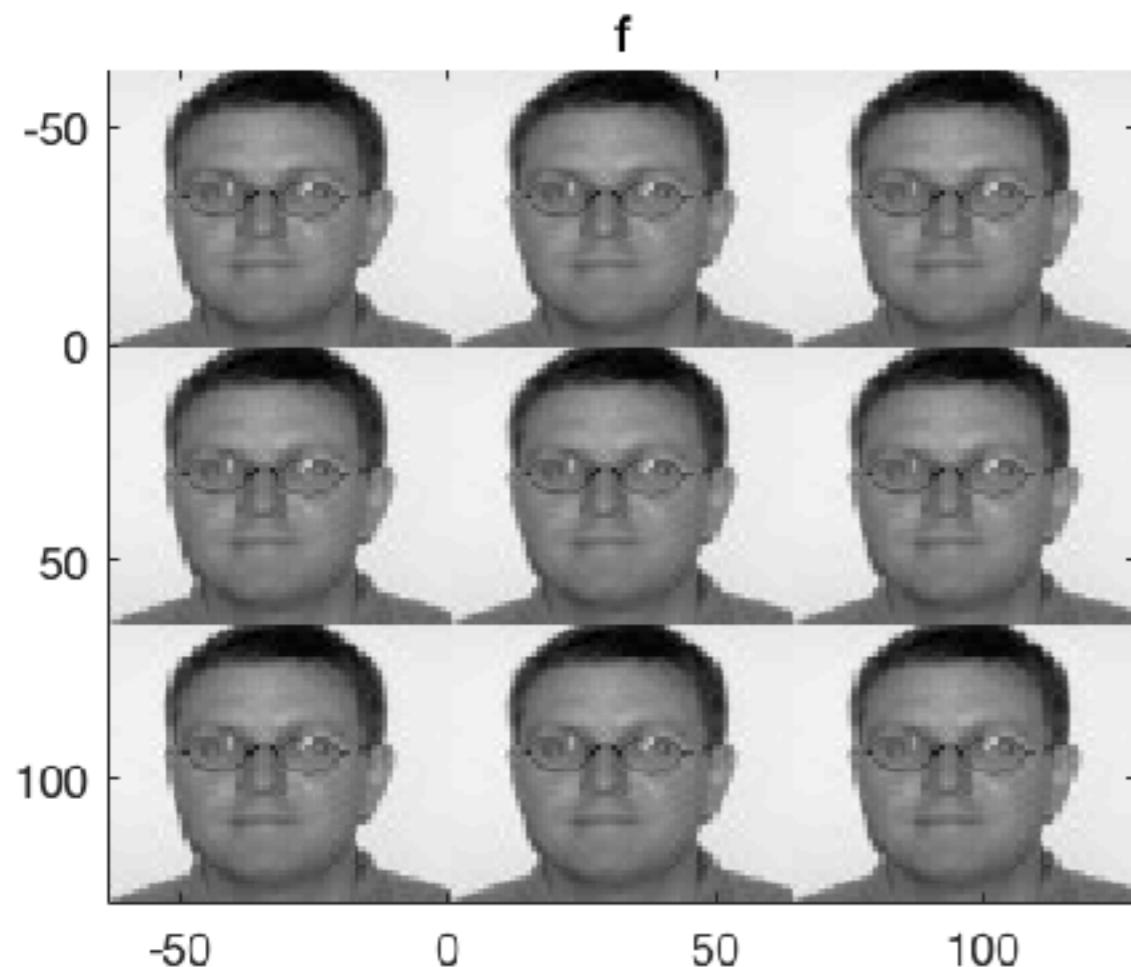
Example – Periodic expansion



(1-M x 1-N)

Discrete Fourier Transform - 2D

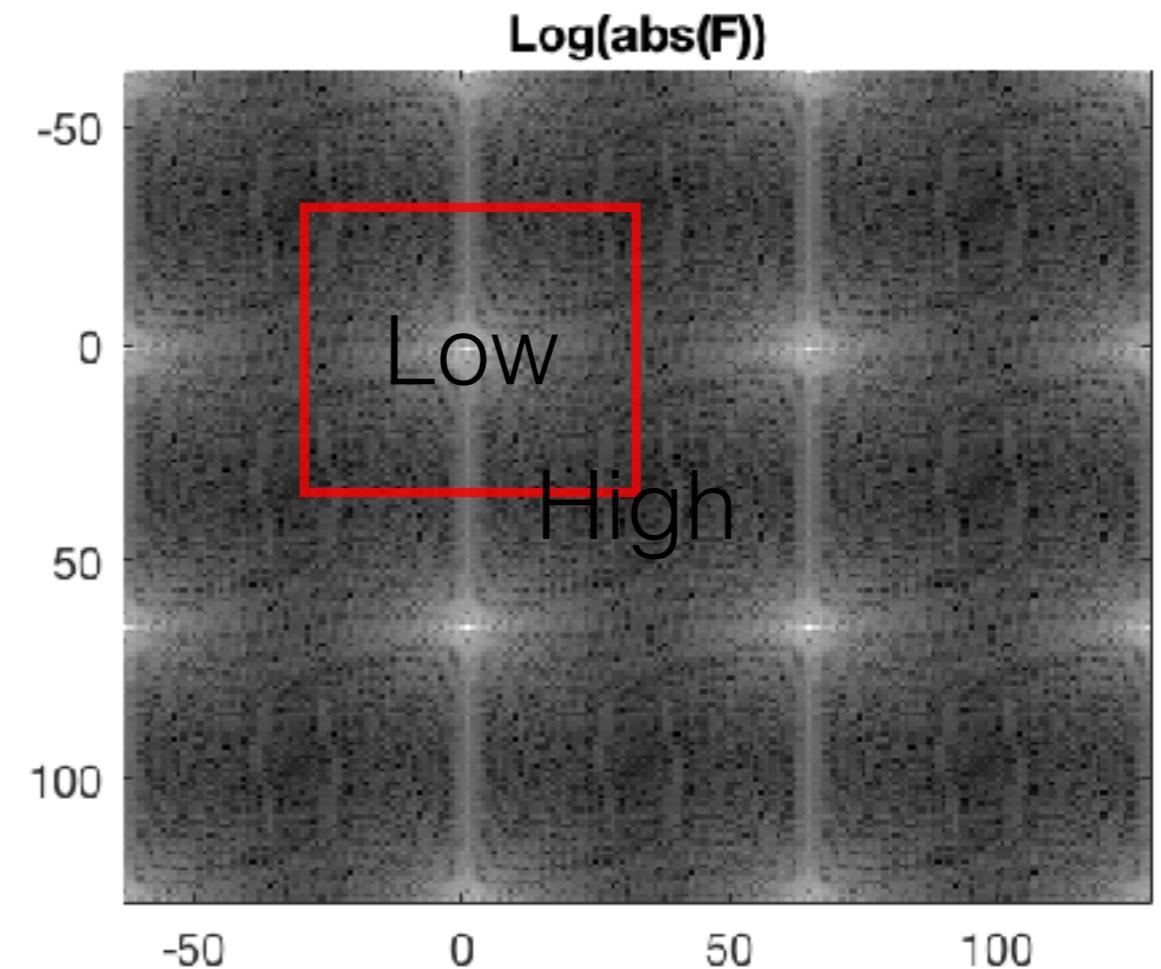
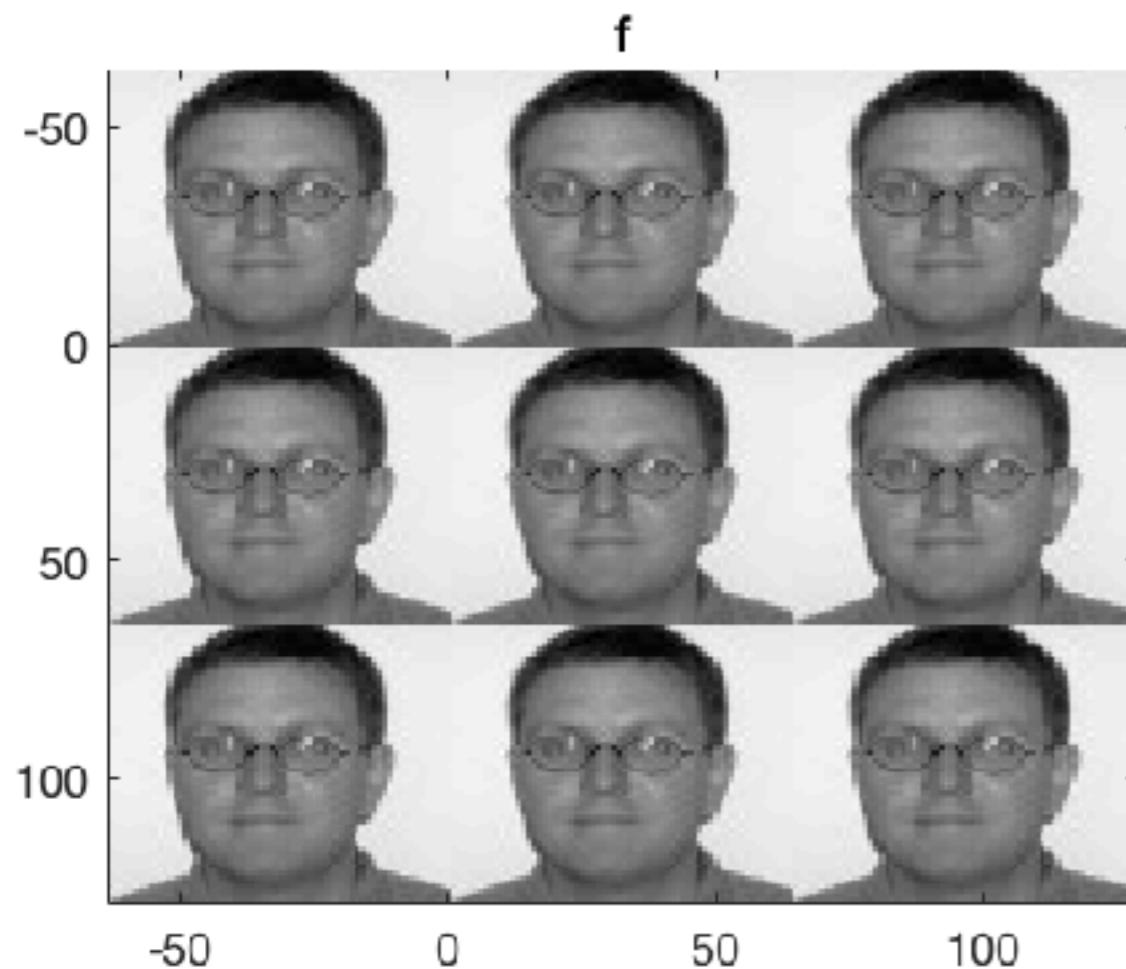
Example – Periodic expansion



$$(-M/2 - M/2 \times -N/2 - N/2)$$

Discrete Fourier Transform - 2D

Example – Periodic expansion



fftshift

$$(-M/2 - M/2 \times -N/2 - N/2)$$

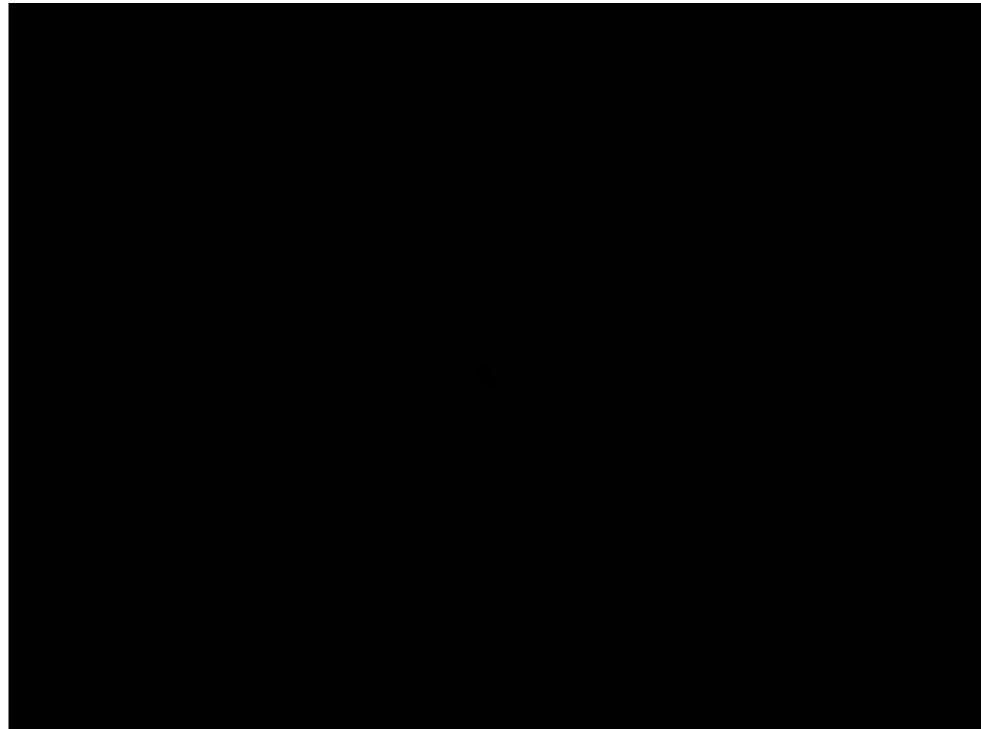
- ▶ Usually, the gray-levels of the Fourier Transform images are scaled using $c \log(1 + |F(u, v)|)$.
- ▶ The middle of the Fourier image (after fftshift) corresponds to low frequencies.
- ▶ Outside the middle high components in F corresponds to higher frequencies and the direction corresponds to "edges" in the images with opposite orientation.

Fourier transform



- Image

Fourier transform

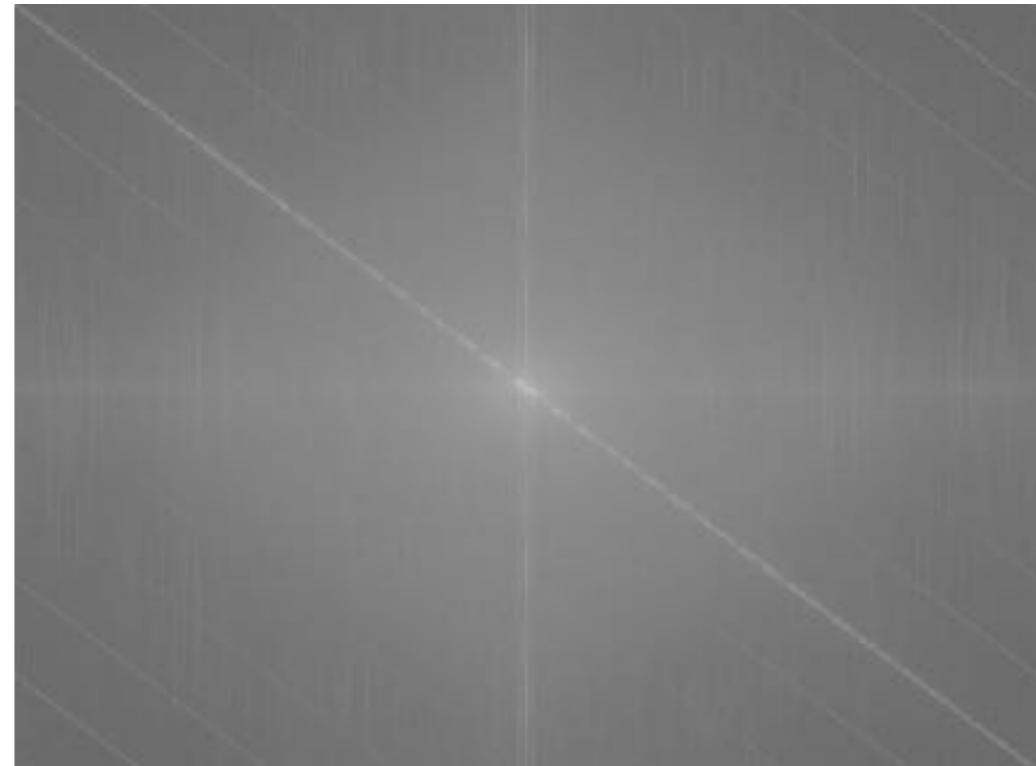


- Image
- $\text{abs}(\text{fft2}(I))$

Fourier transform

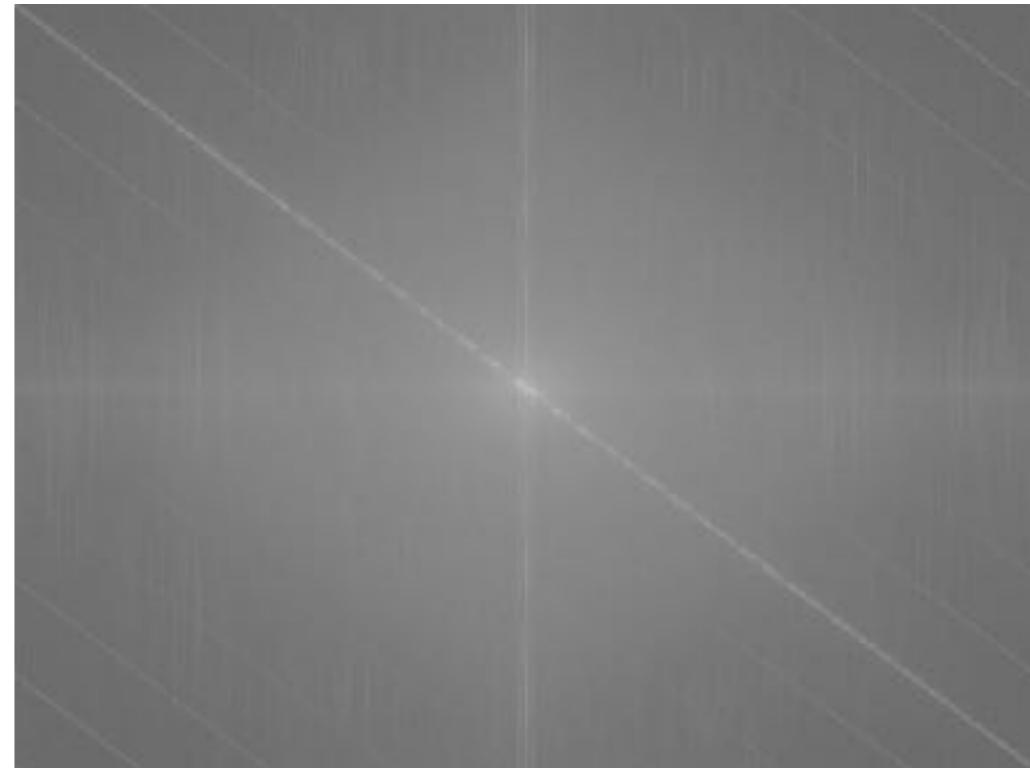


- Image



- $\log(\text{abs}(\text{fft2}(I)))$

Edge effects



- Image
- $\log(\text{abs}(\text{fft2}(I)))$

Fourier transform



- Image

Fourier transform



- Image



- Fourier transform

Fourier transform

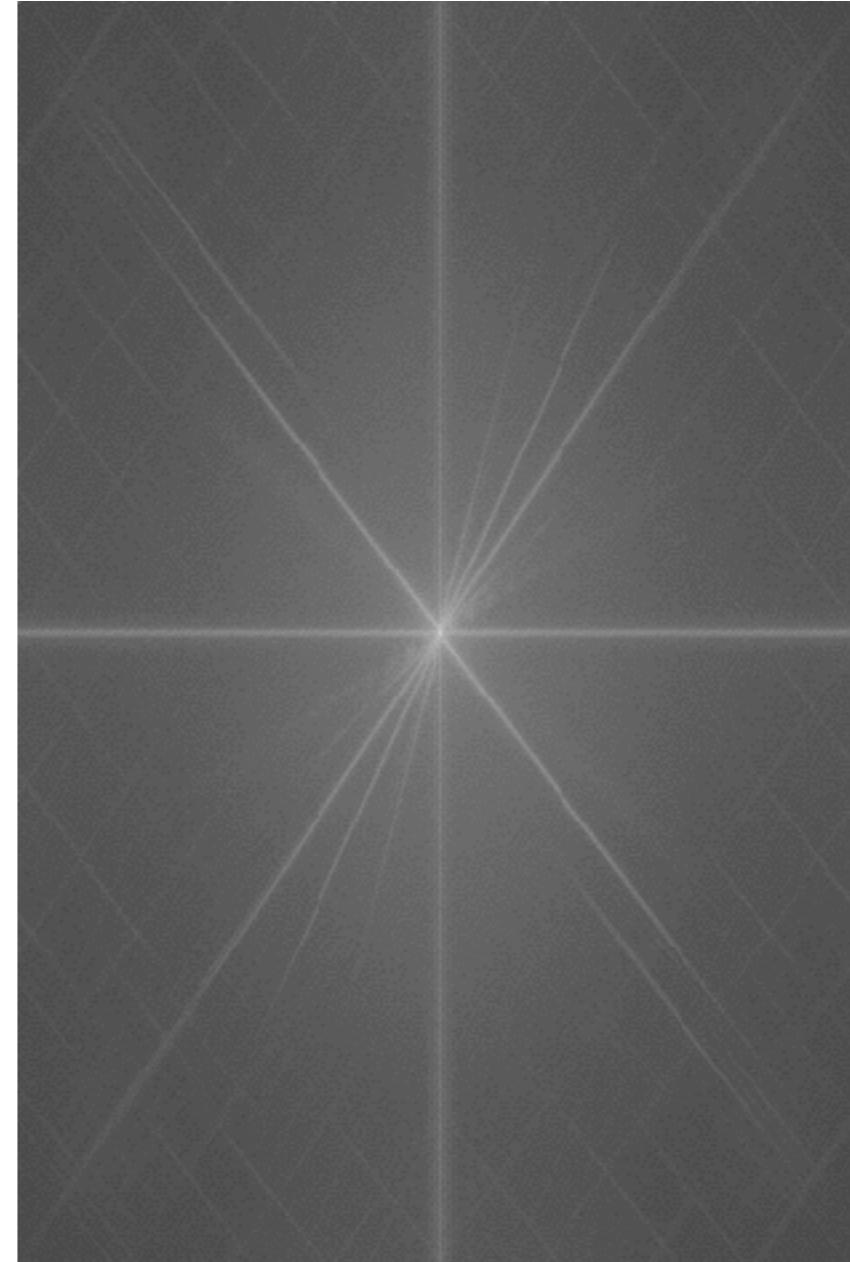


- Image

Fourier transform



- Image



- Fourier transform

Using FFT for convolutions

1. $f \rightarrow FFT \rightarrow F$
2. $h \rightarrow FFT \rightarrow H$
3. $H, F \rightarrow \times \rightarrow H \cdot F$
4. $H \cdot F \rightarrow IFFT \rightarrow h * f$

The computational complexity of using FFT for a convolution is:

$$2 \frac{N \log N}{2} + N + \frac{N \log N}{2} \sim \frac{3}{2} N \log N$$

Calculation based on the definition gives complexity N^2

Frequency function

$$g(x) = h * f = \int h(x - y)f(y)dy$$

$$\mathcal{F}g = G, \mathcal{F}h = H, \mathcal{F}f = F$$

$$G(u, v) = H(u, v)F(u, v) .$$

Definition

$H = \mathcal{F}(h)$ is called the **frequency function** of h .

Filter for image enhancement

signal plane	frequency plane
smoothing	low pass
sharpening	high pass

For discrete functions: $DFT(h * f)(u, v) = H(u, v)F(u, v)$.

Let the output, g be given by the convolution

$$g(x) = S(f)(x) = \int h(x - y)f(y)dy ,$$

where f represents the input and h the impulse response

If $g(x)$ only depends on f :s values in a surrounding (=a small window) of x then S is called a **window operator**.

The window is given by $\{ x \mid h(x) \neq 0 \}$.

Assume that $f(x, y)$ represents a continuous image. Let

$$h(x, y) = \text{rect}(x) \text{rect}(y) .$$

Then

$$S(f) = h * f = \int_{K(x,y)} f(s, t) ds dt ,$$

where the region of integration $K(x, y)$ is a unit square with centre at (x, y) .

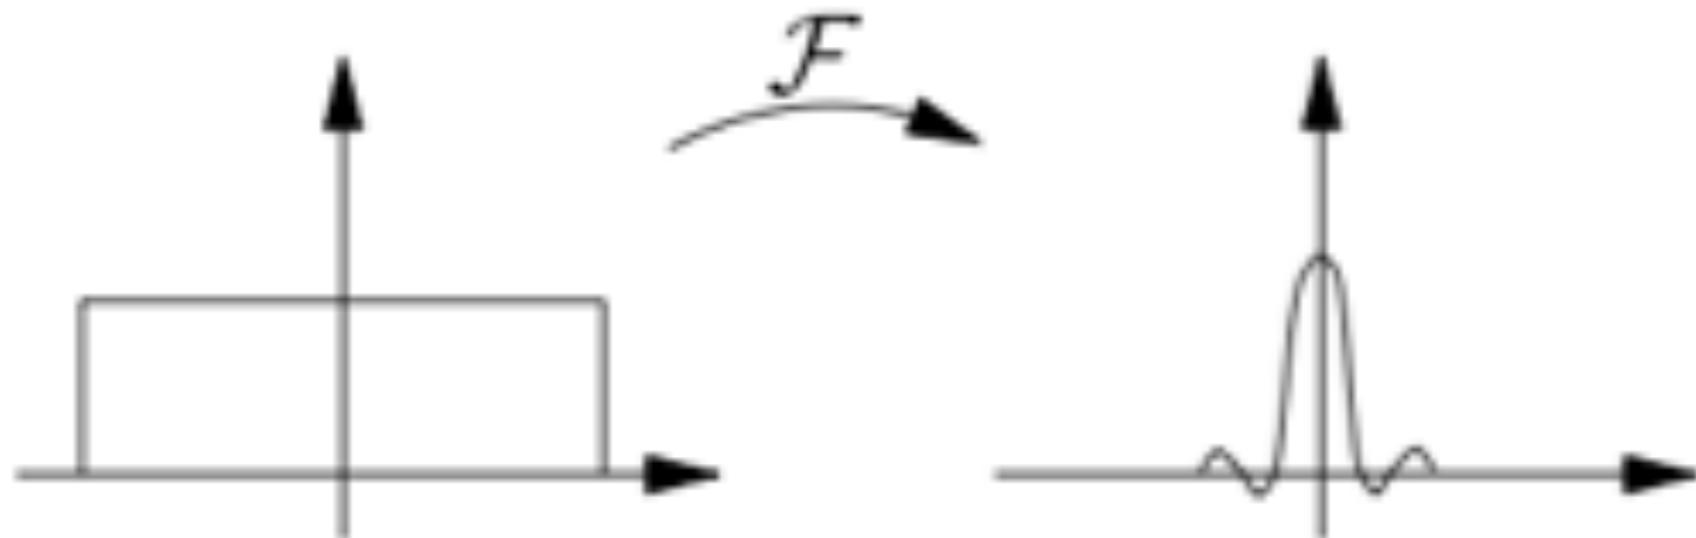
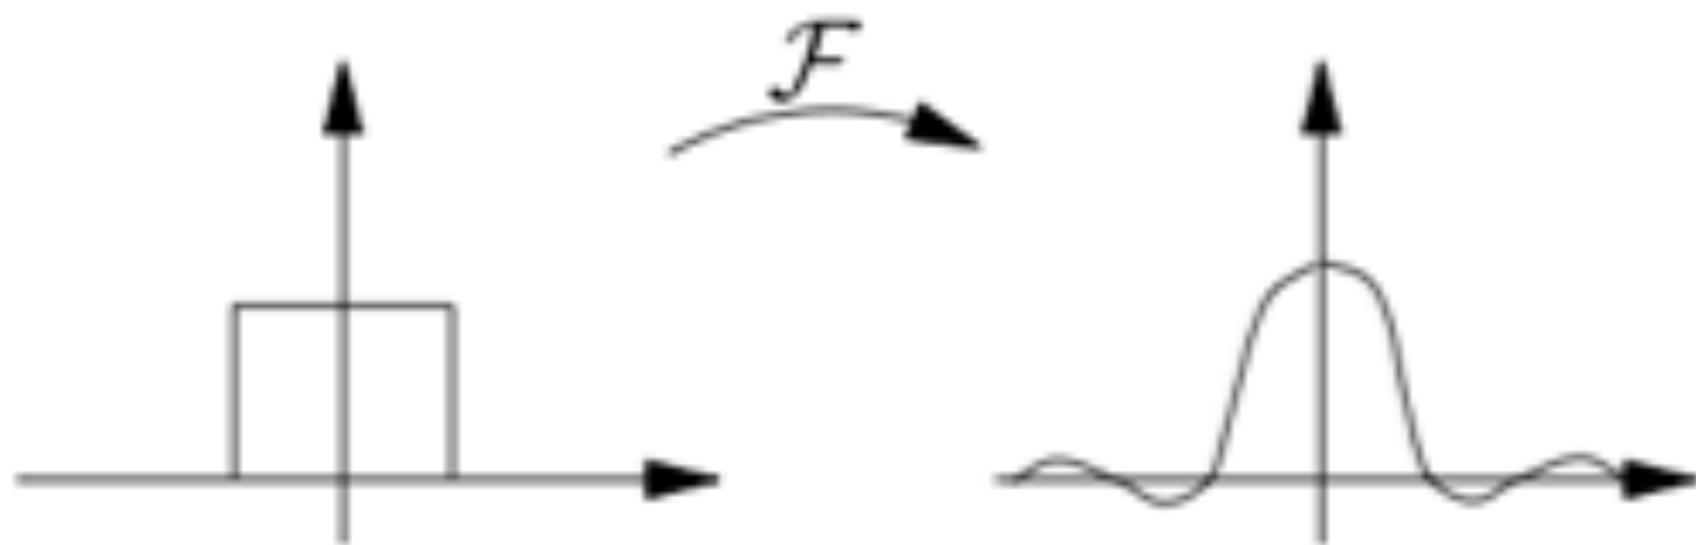
S is called a **mean value operator**.

The fourier transform gives

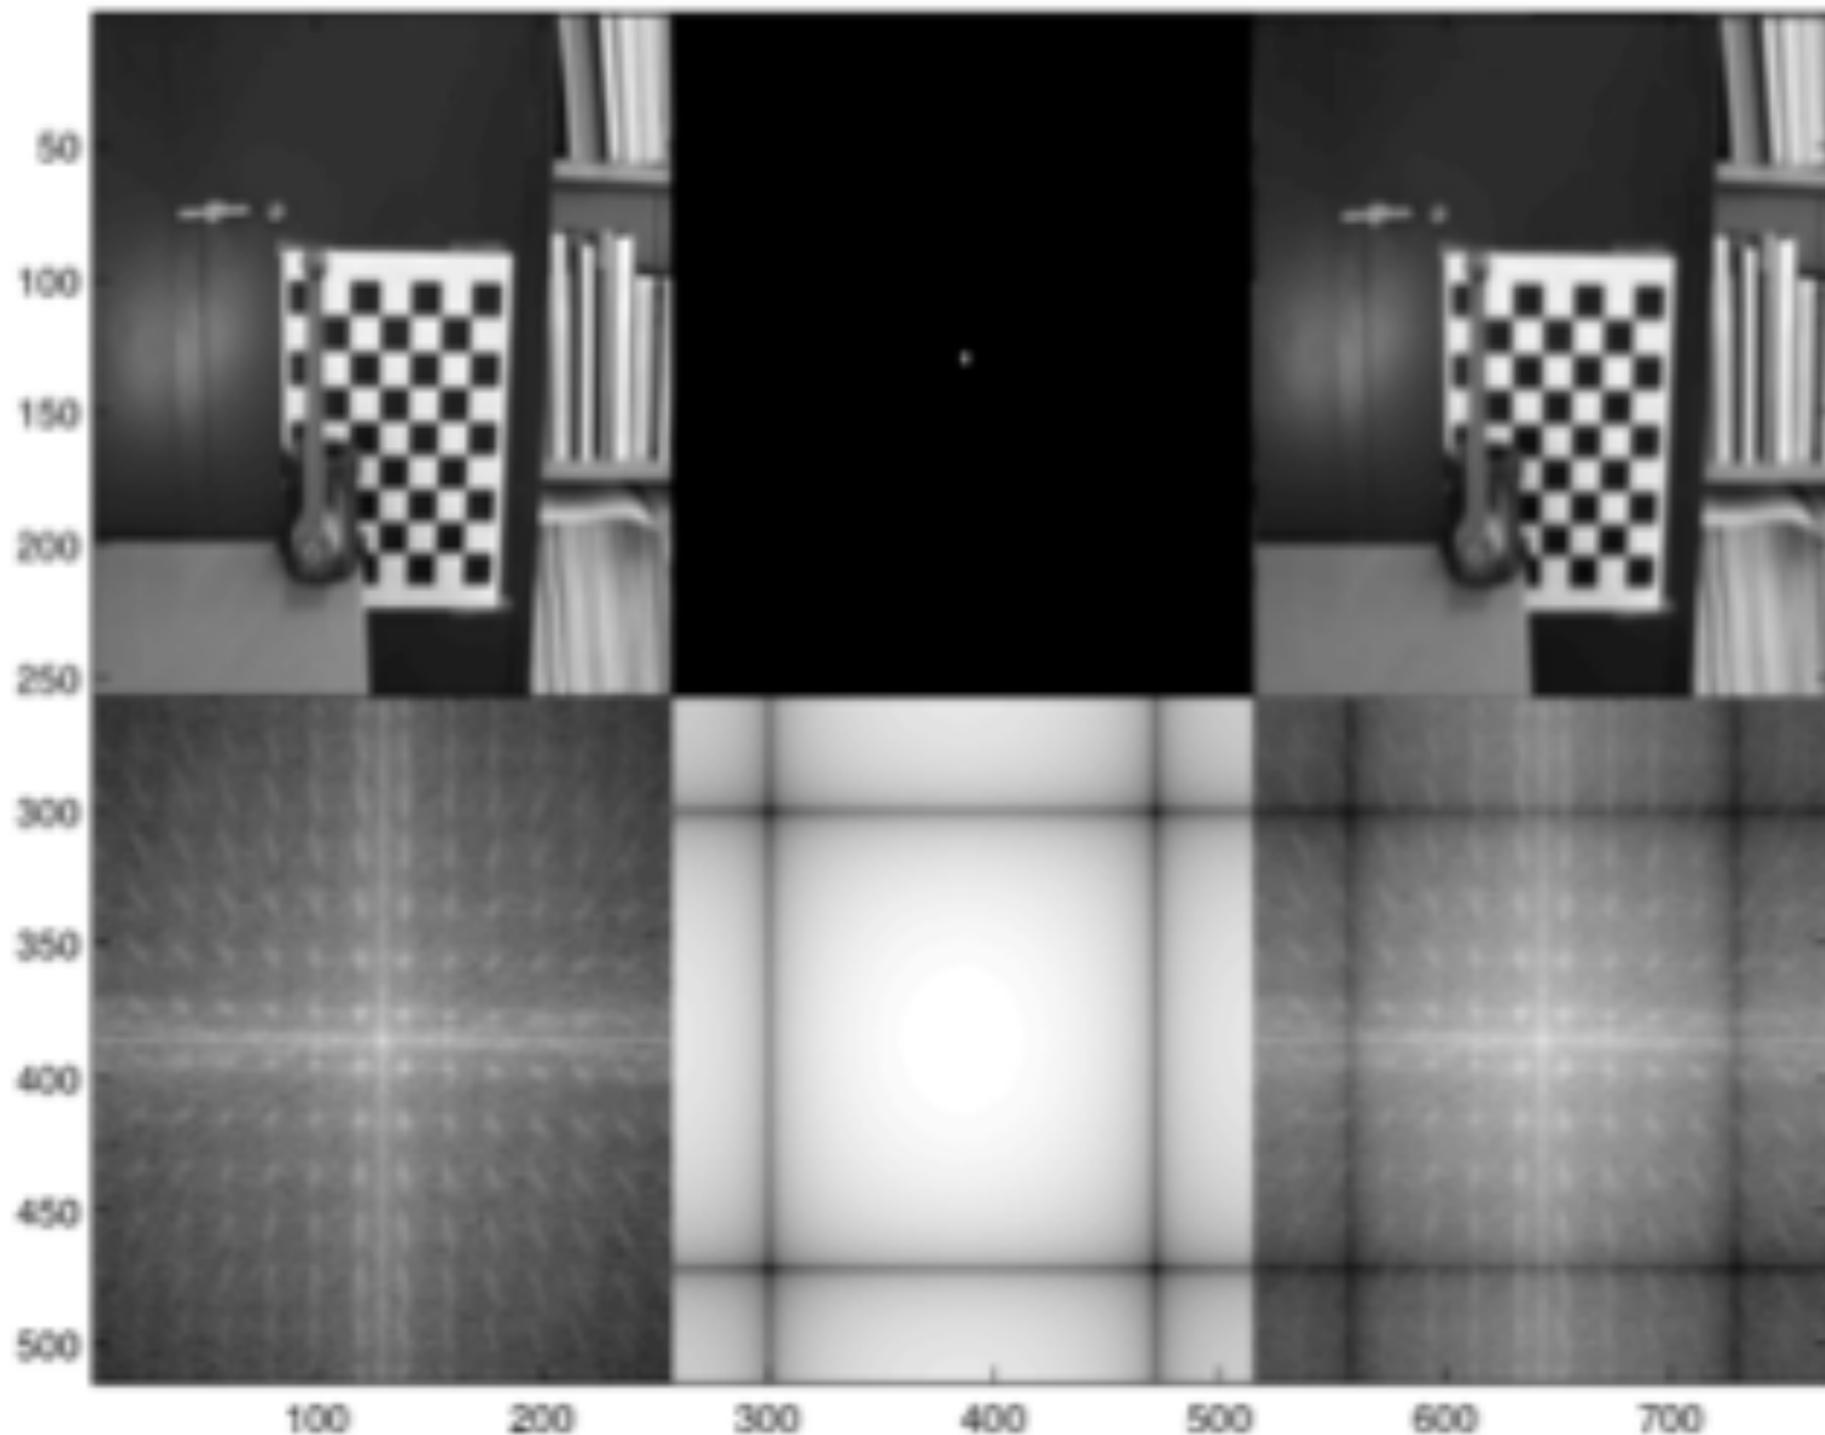
$$H(u) = 4 \operatorname{sinc}(2\pi u) \operatorname{sinc}(2\pi v) .$$

The scaling rule (page 148 in Forsythe-Ponce)

$$f(\lambda x) \rightarrow \frac{1}{\lambda} F\left(\frac{u}{\lambda}\right) .$$

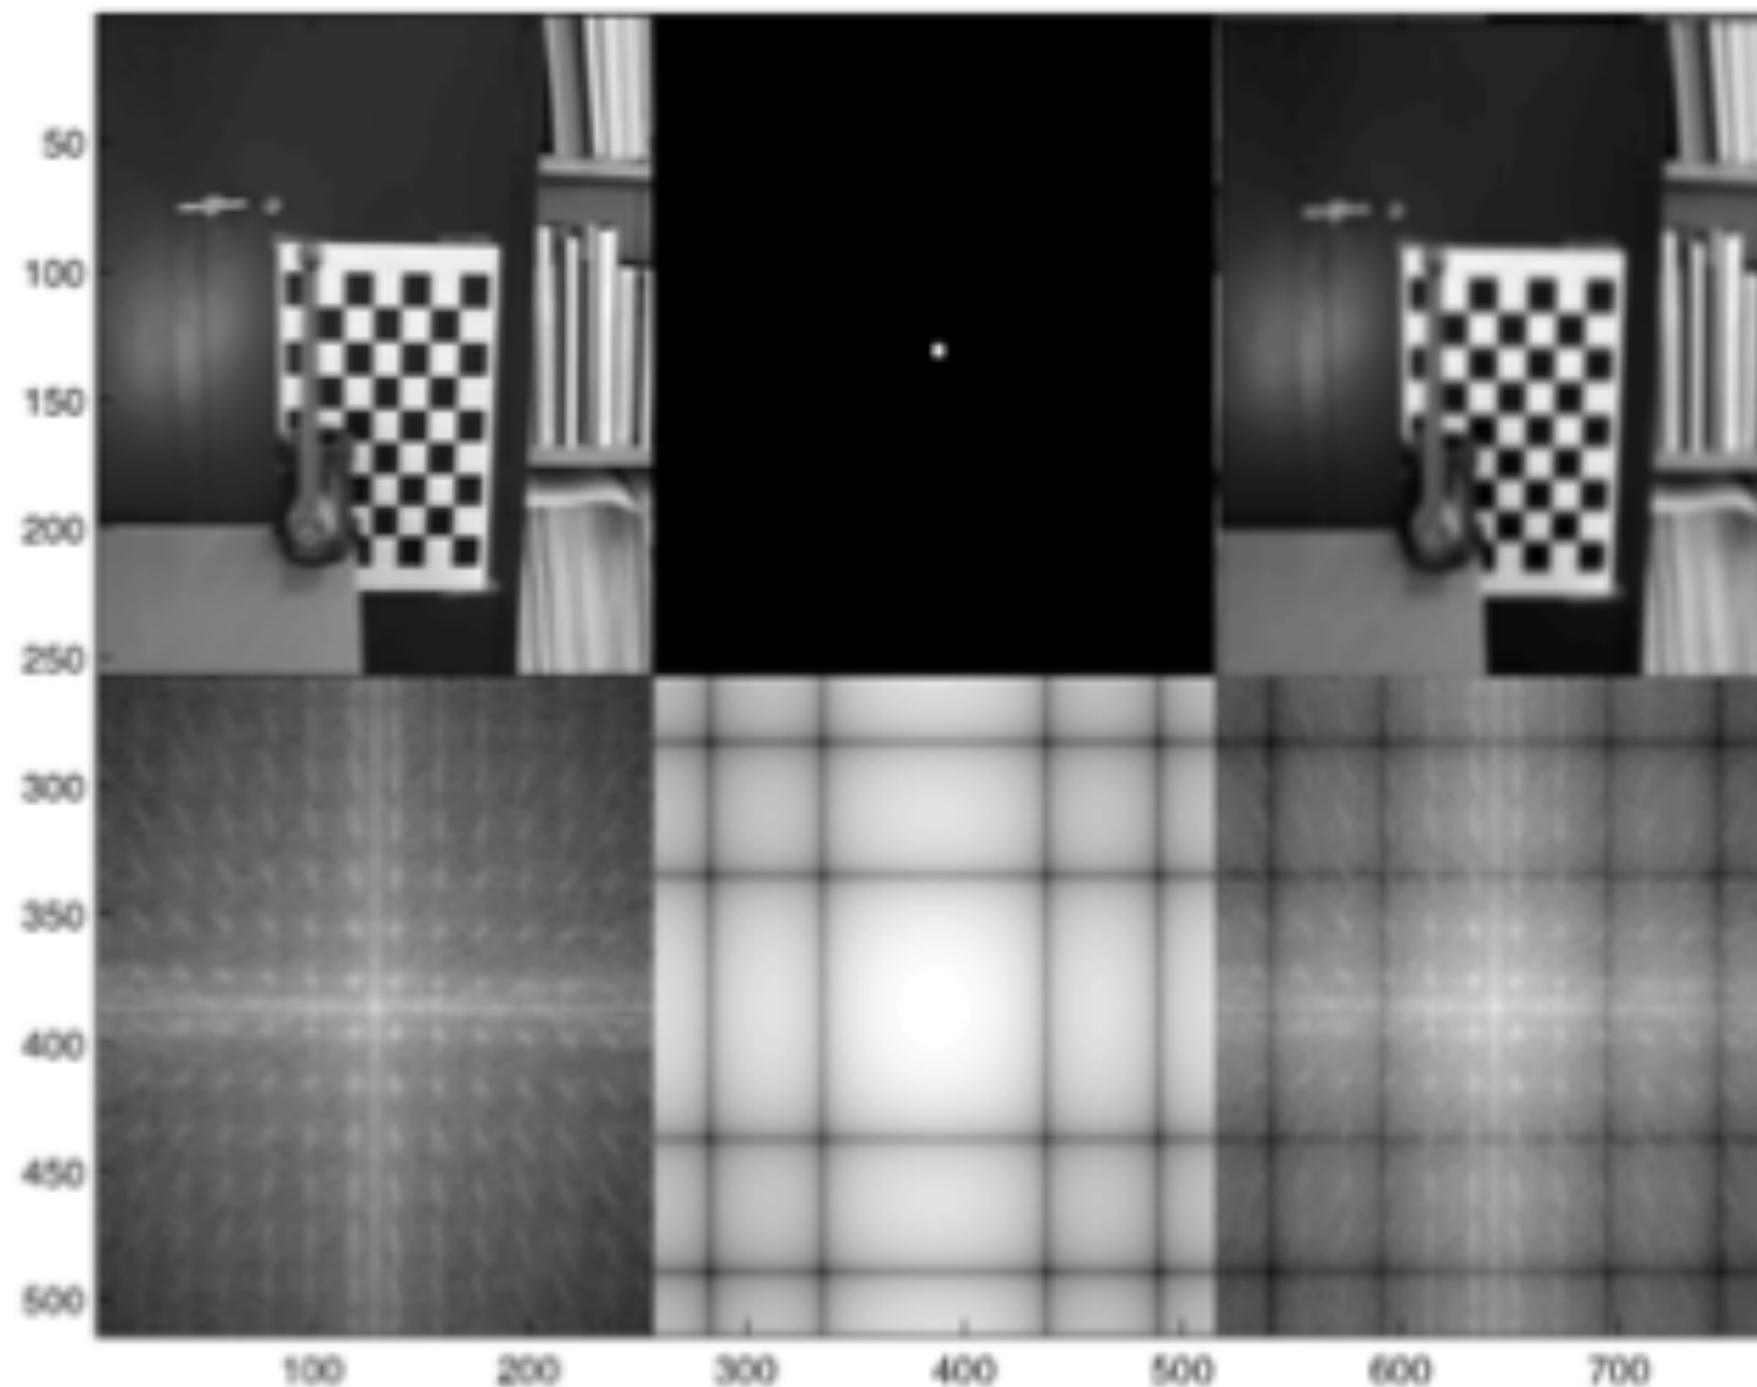


signal space: image, filter, result



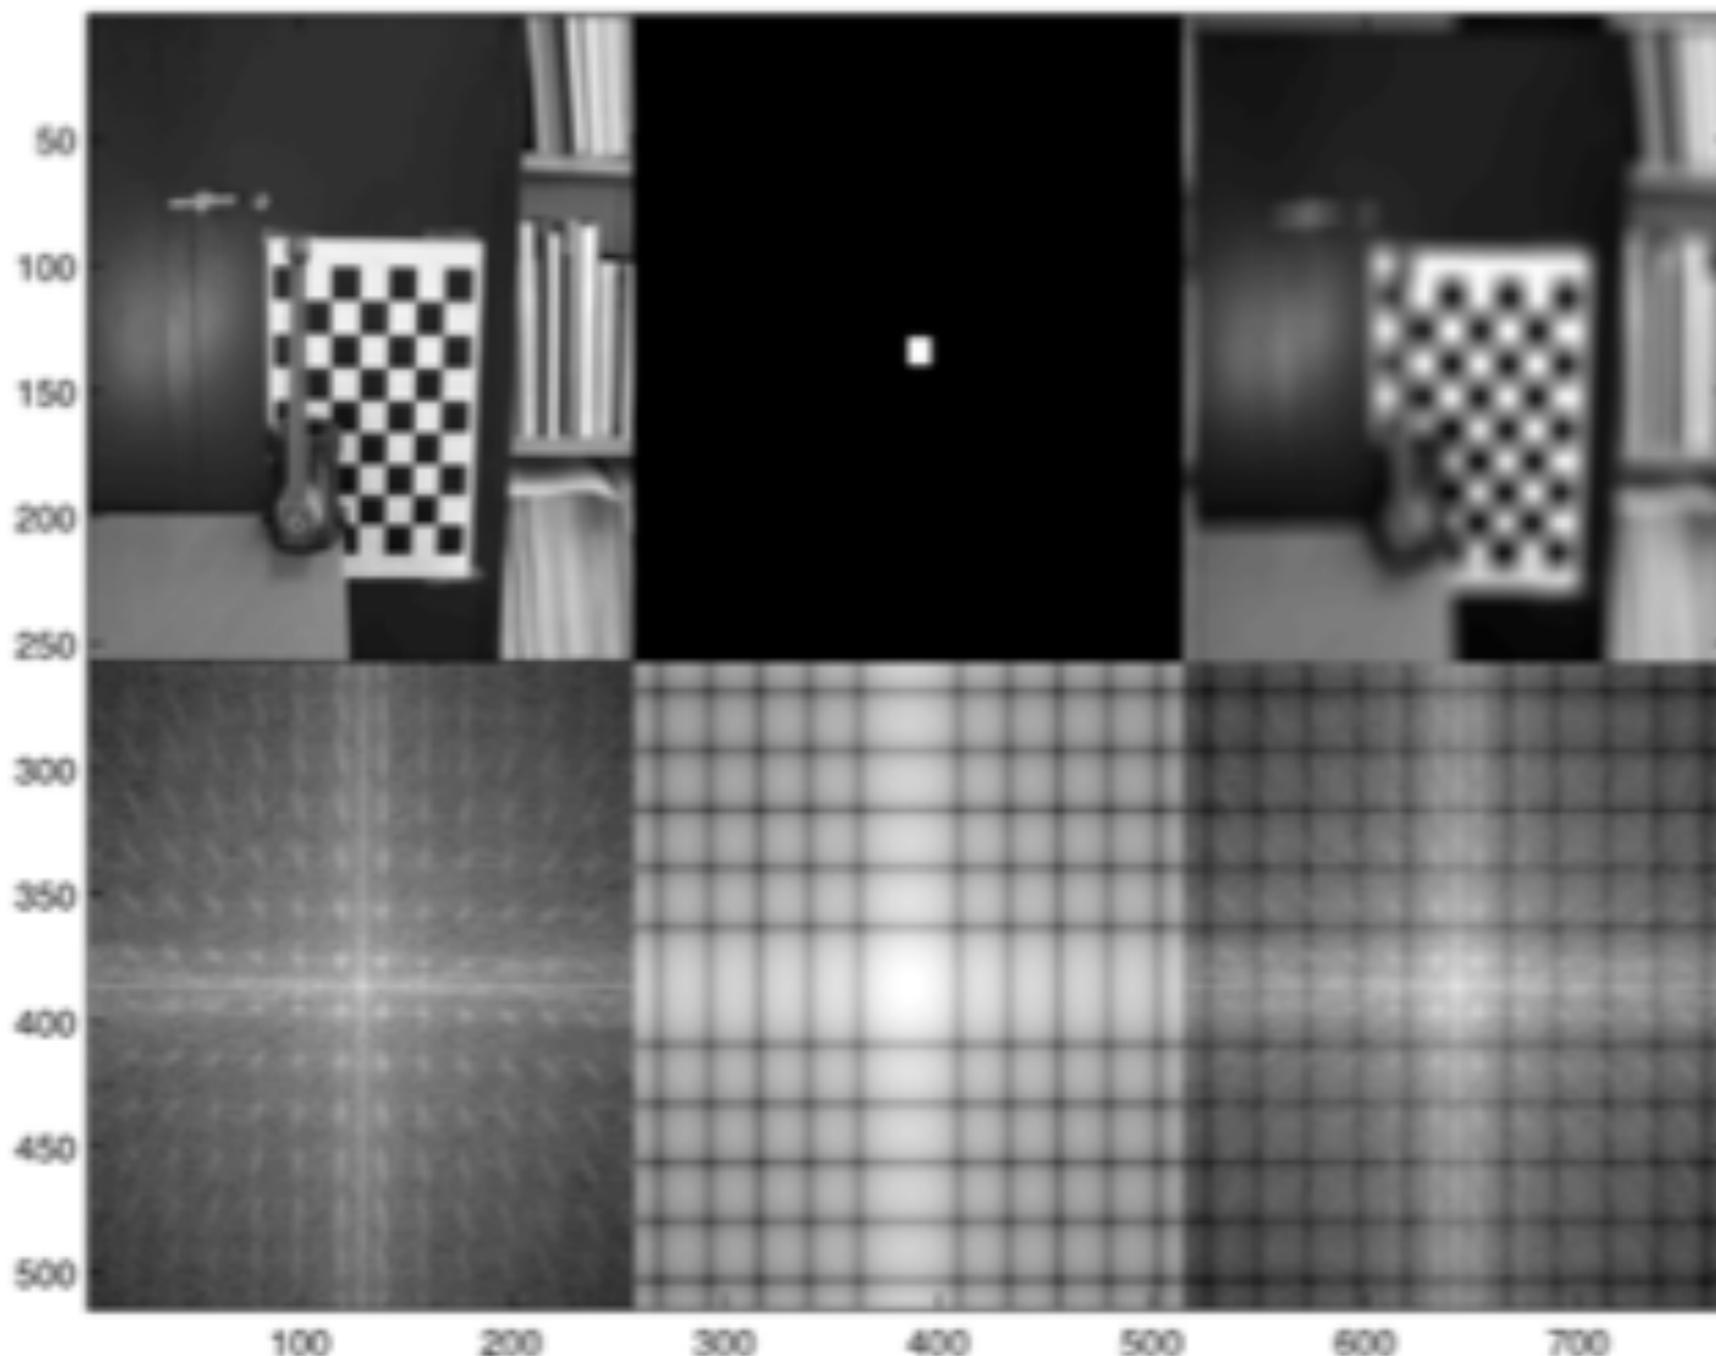
frequency space: image, filter result

signal space: image, filter, result



frequency space: image, filter result

signal space: image, filter, result

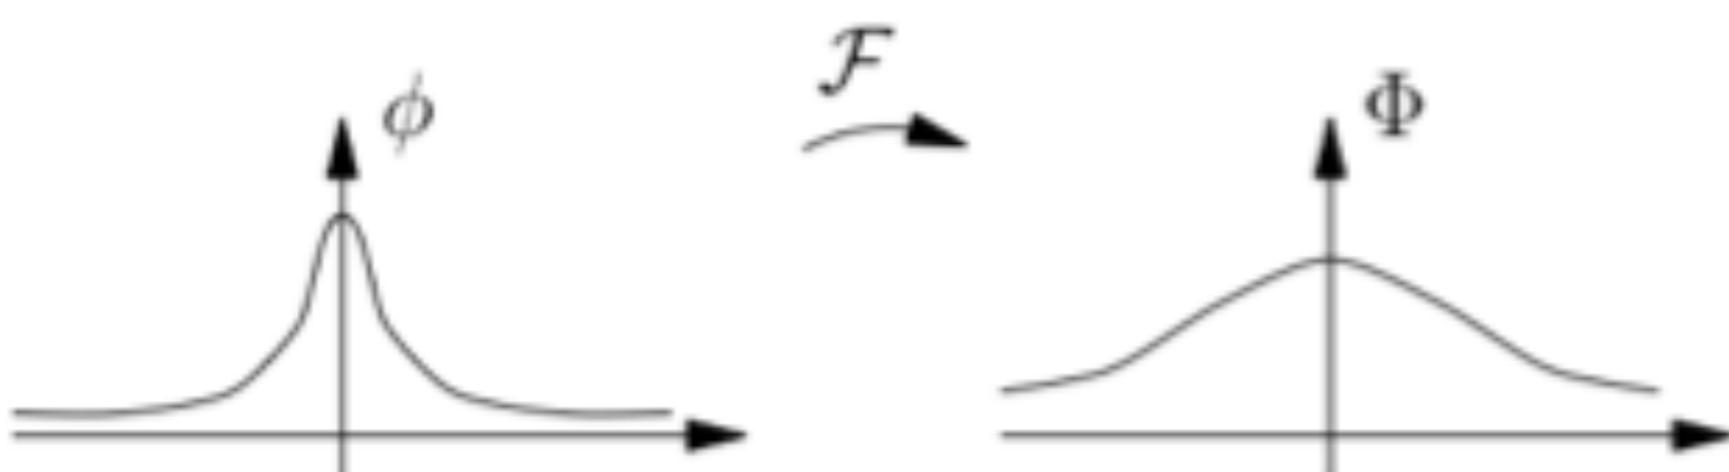


frequency space: image, filter result

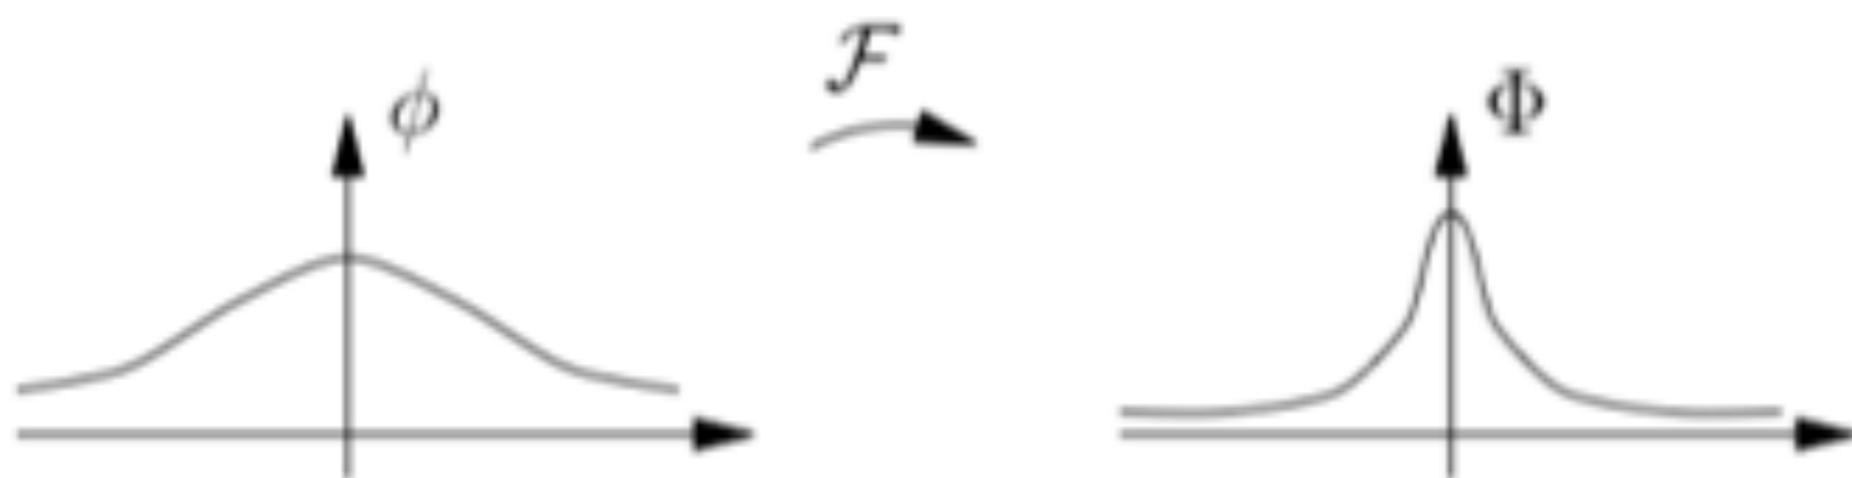
Example

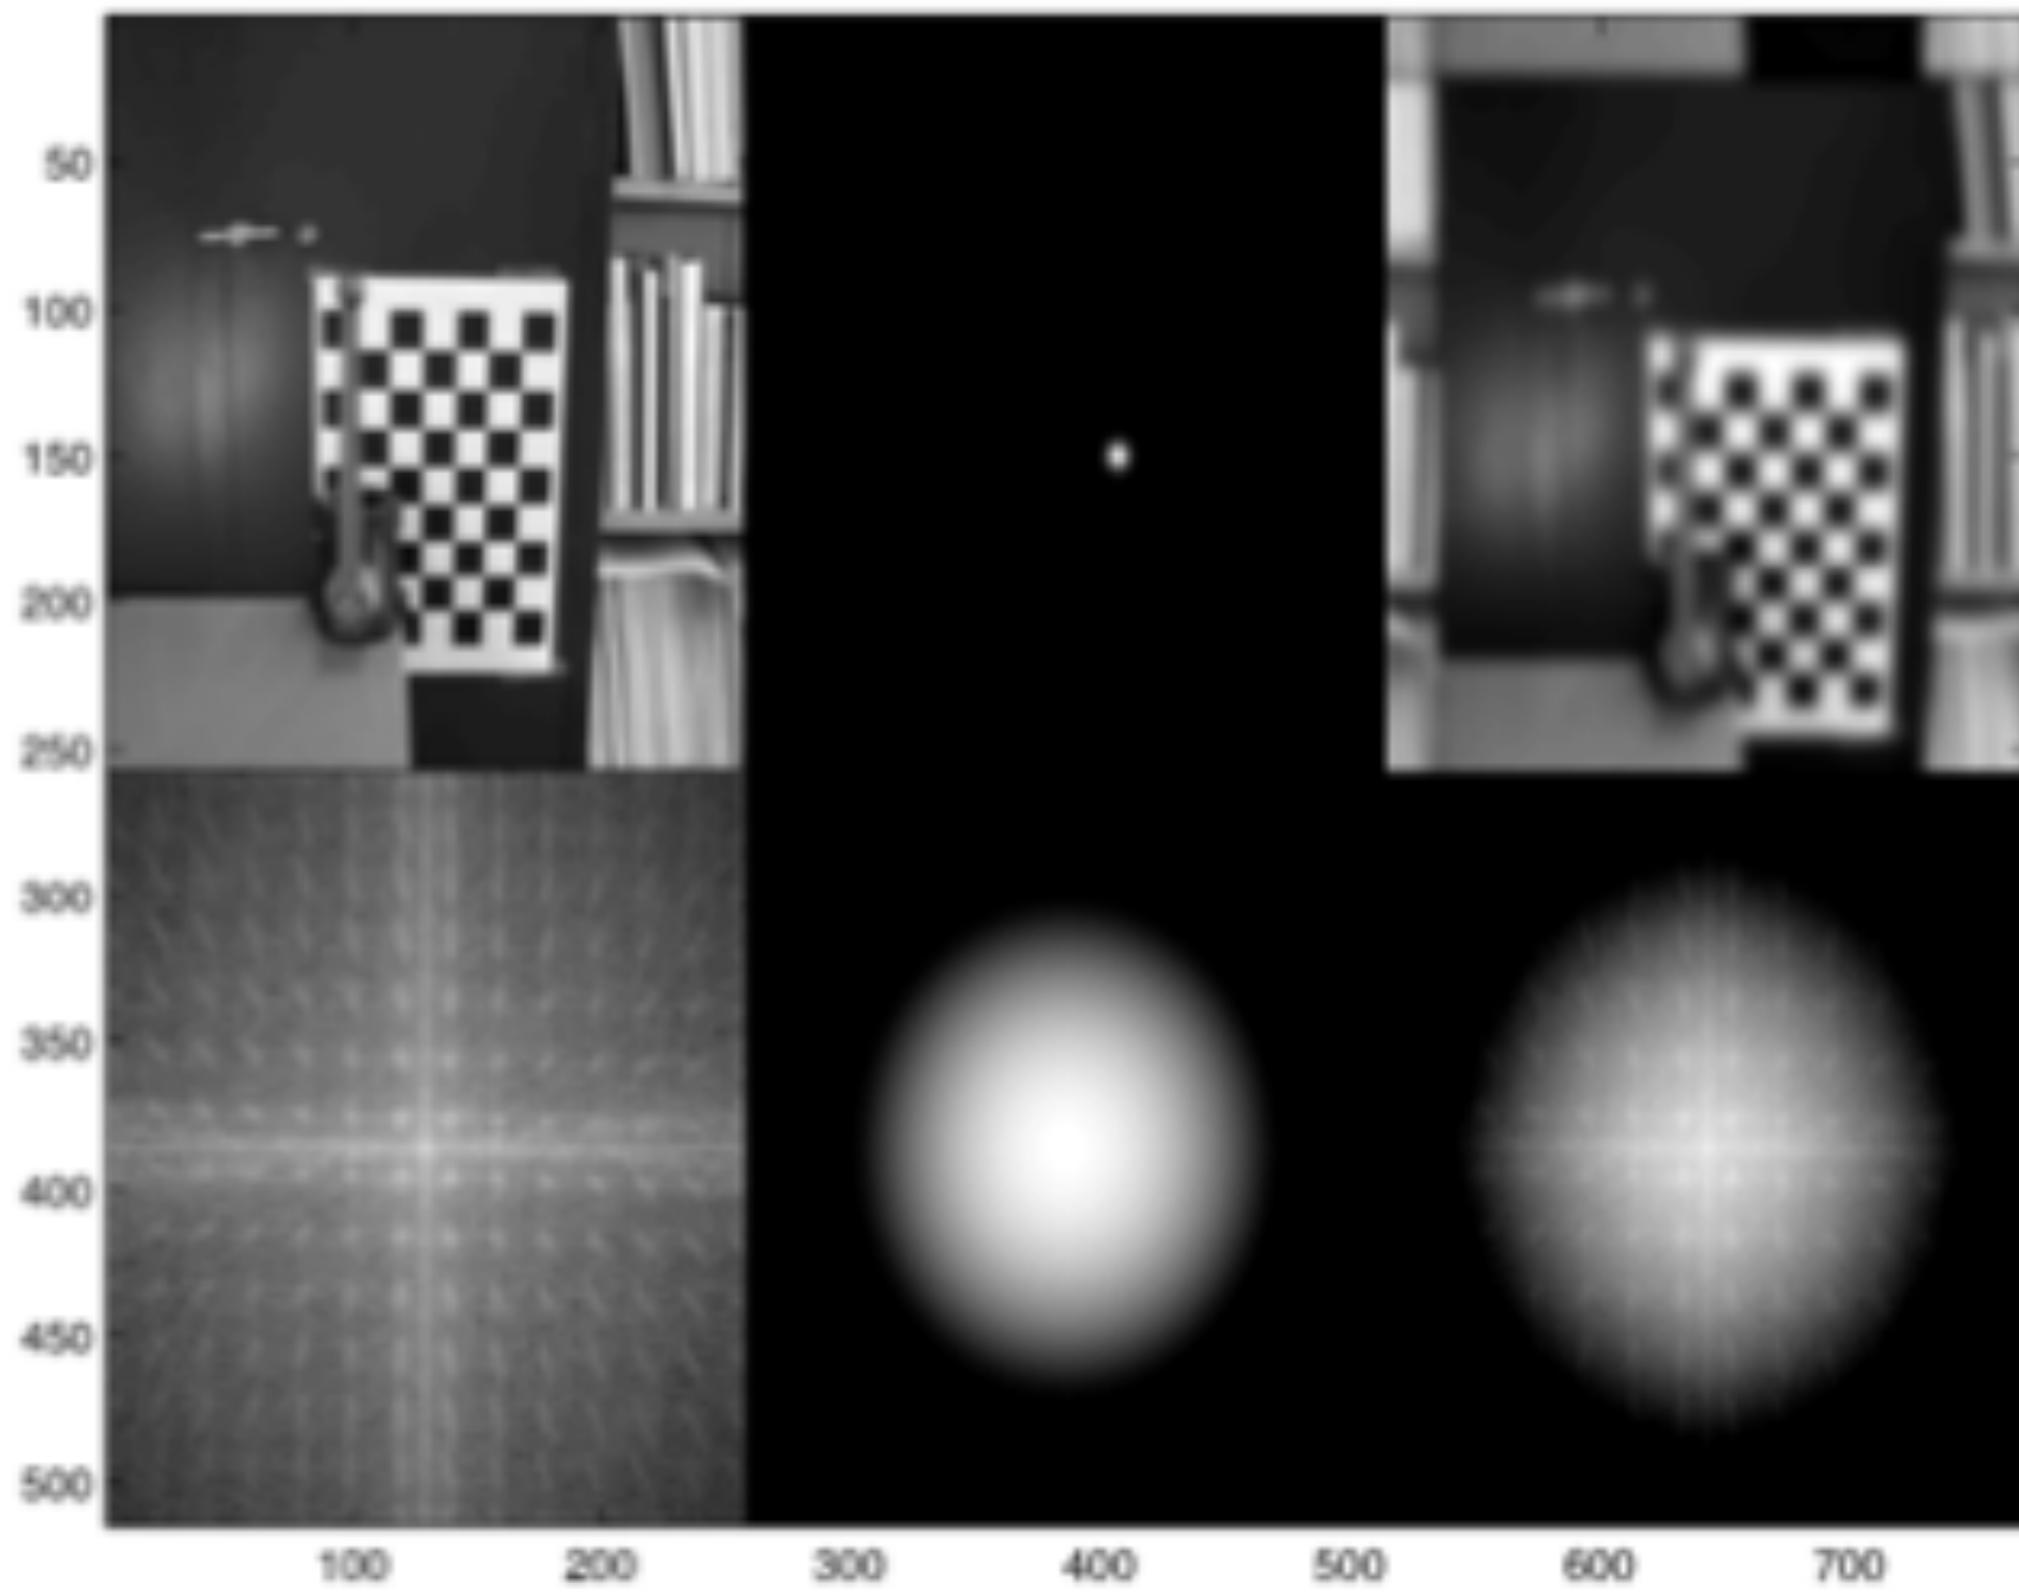
Notice that

$$\phi(x) = \frac{1}{\sqrt{2\sigma^2\pi}} e^{-x^2/(2\sigma^2)} \rightarrow \Phi(u) = e^{-2(\sigma\pi u)^2}.$$



Larger σ gives





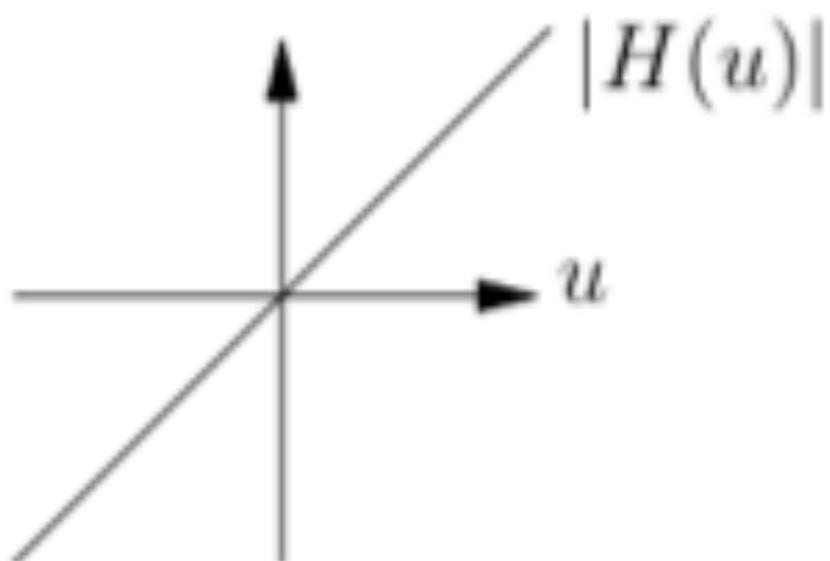
ILLINOIS
UNIVERSITY

Example

Differentiation

$$\frac{\partial f}{\partial x} \rightarrow 2\pi i u F(u)$$

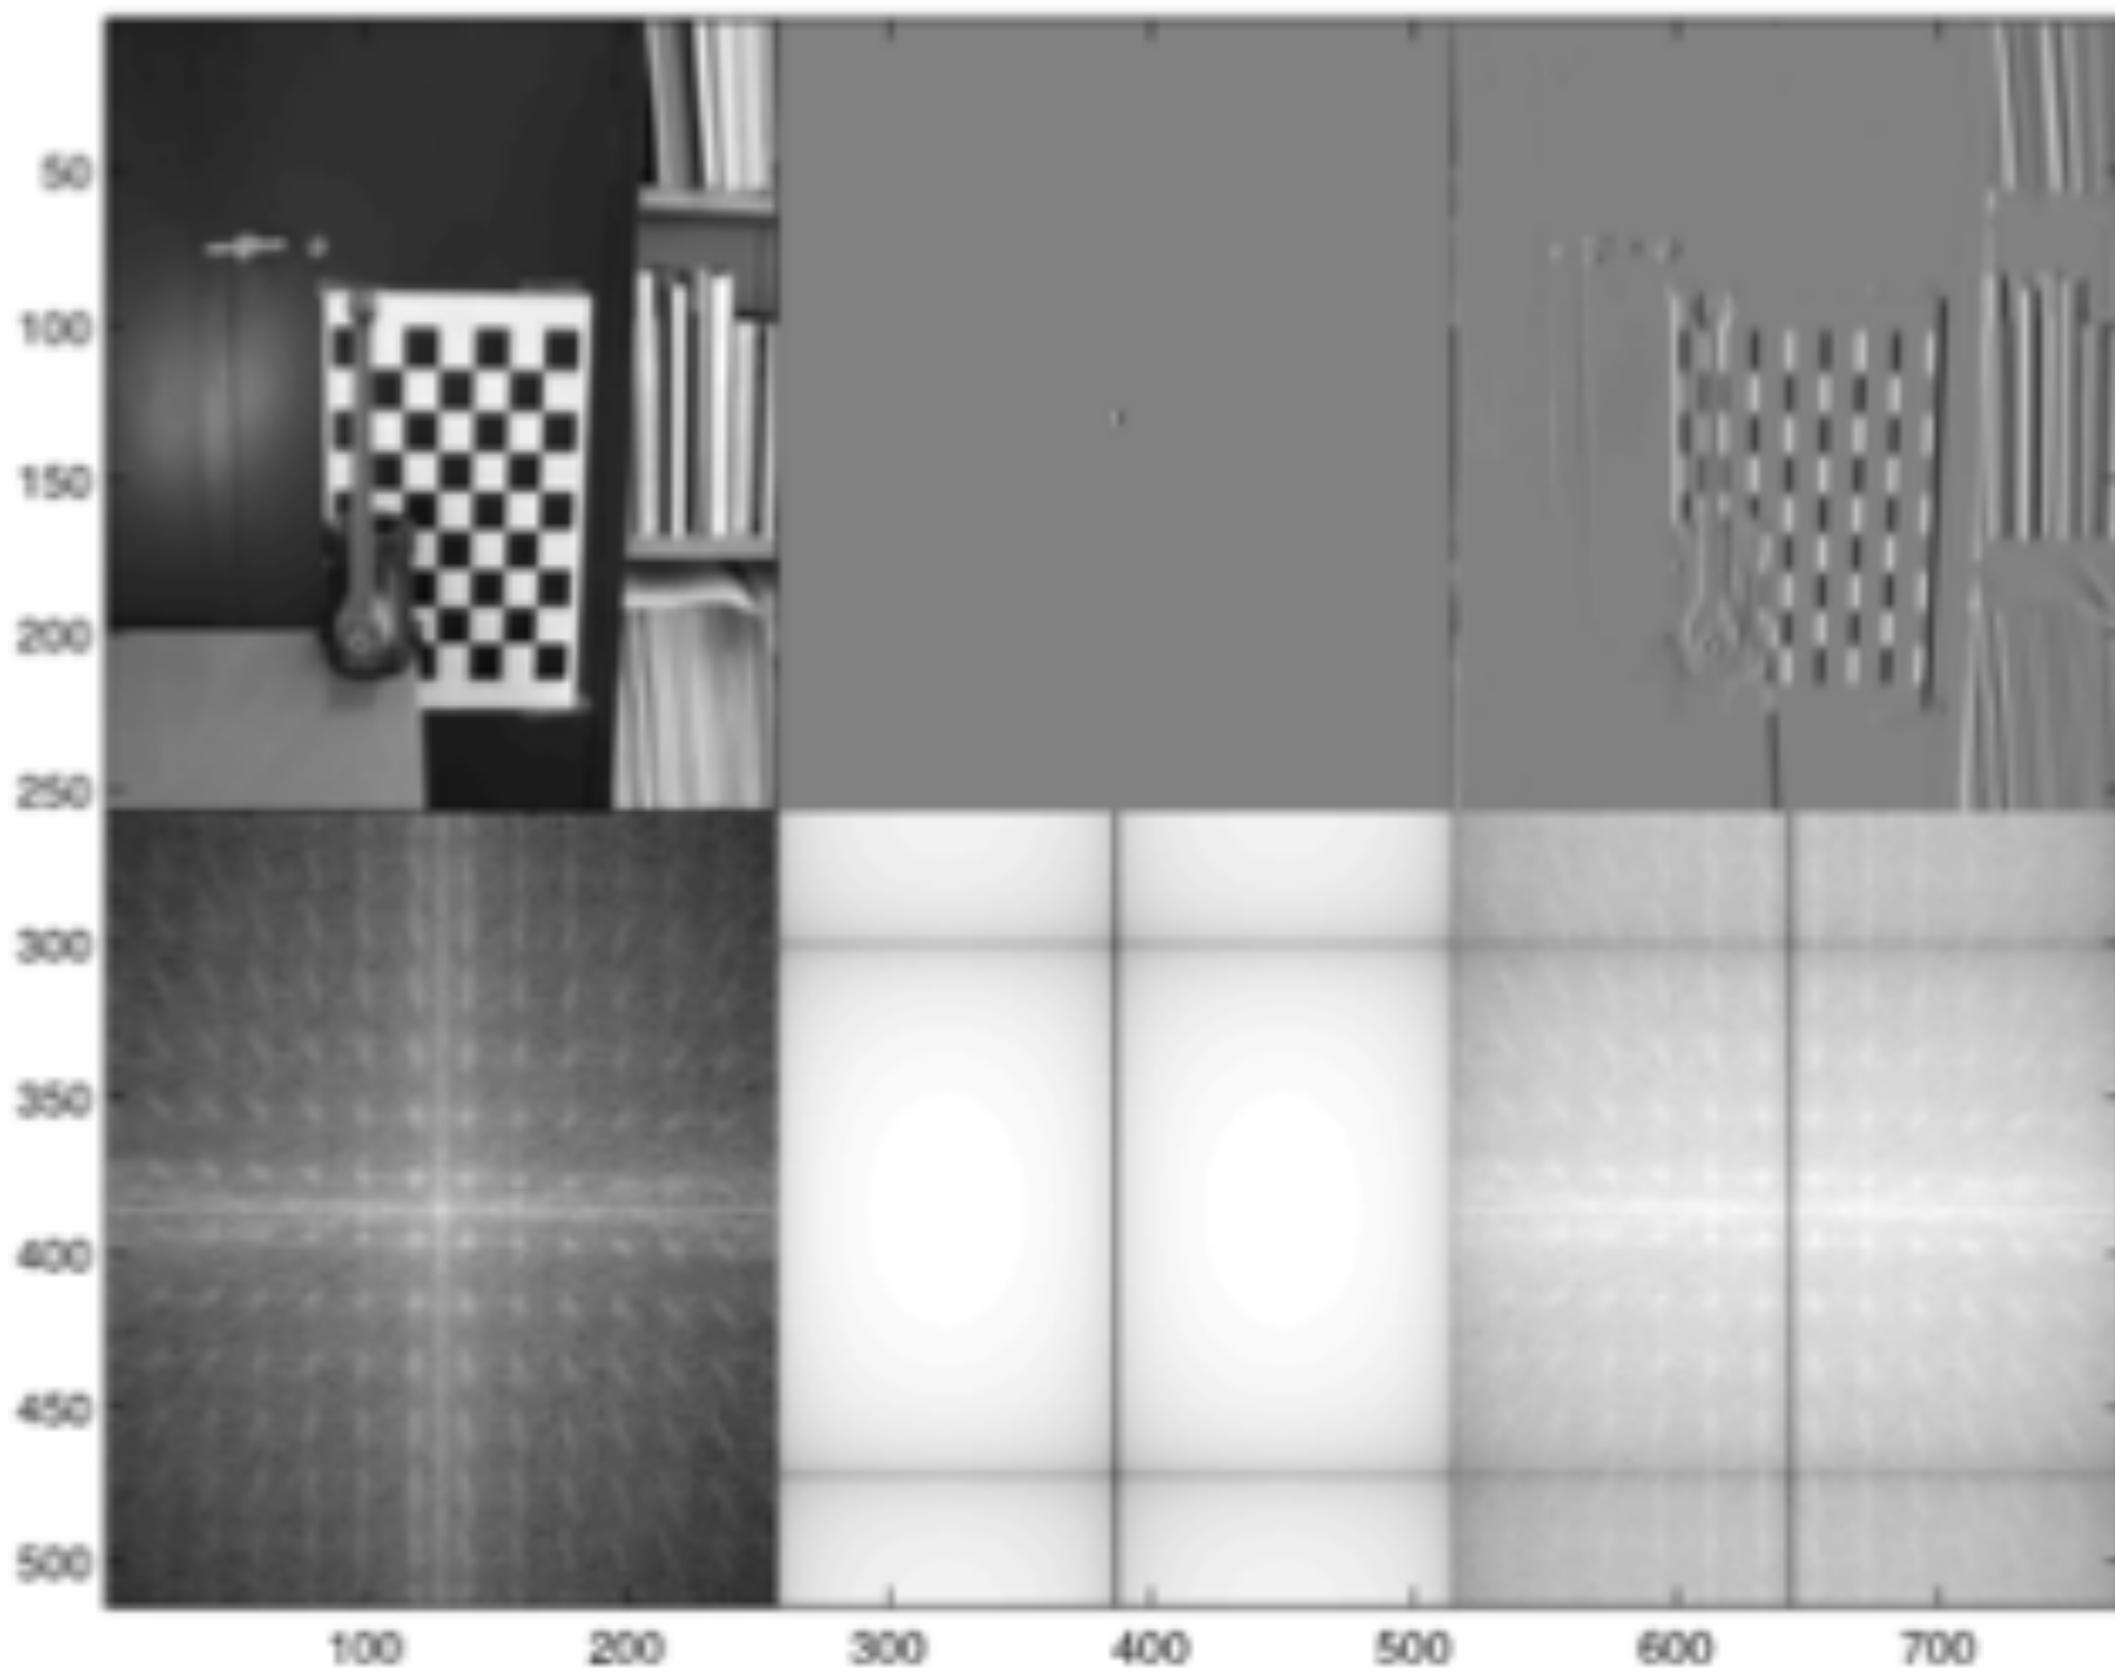
$$H(u) = 2\pi i u$$

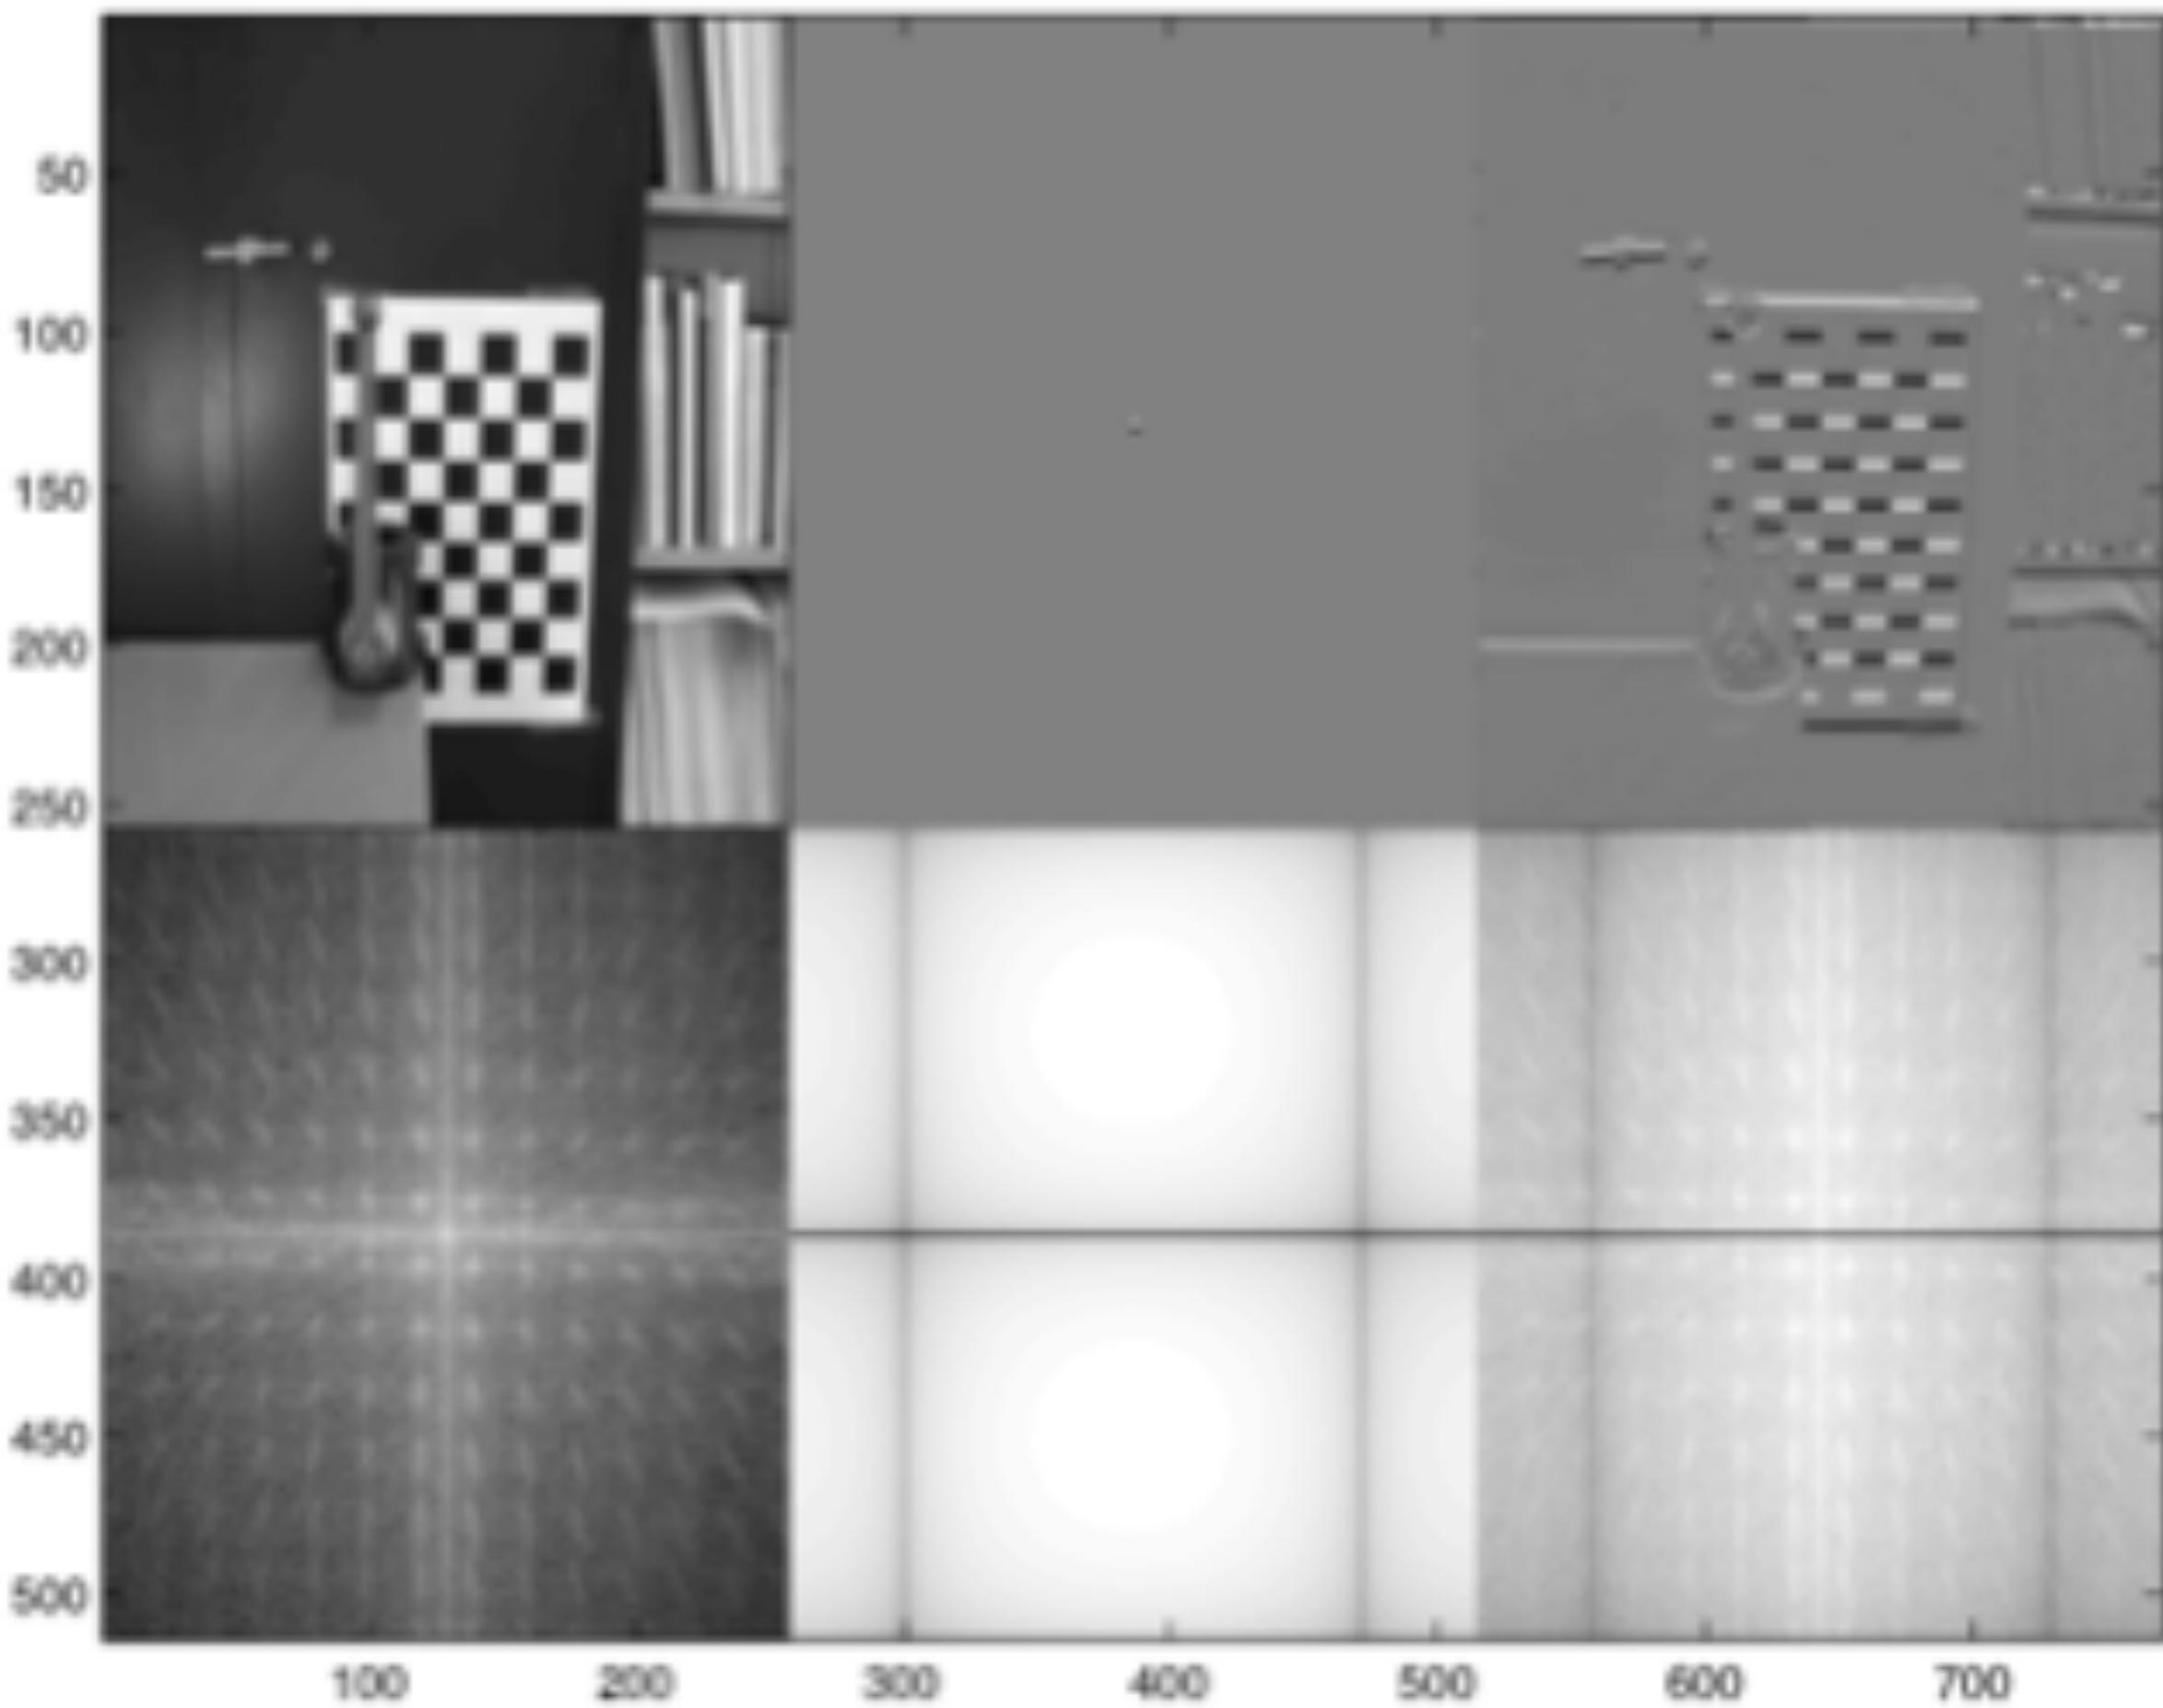


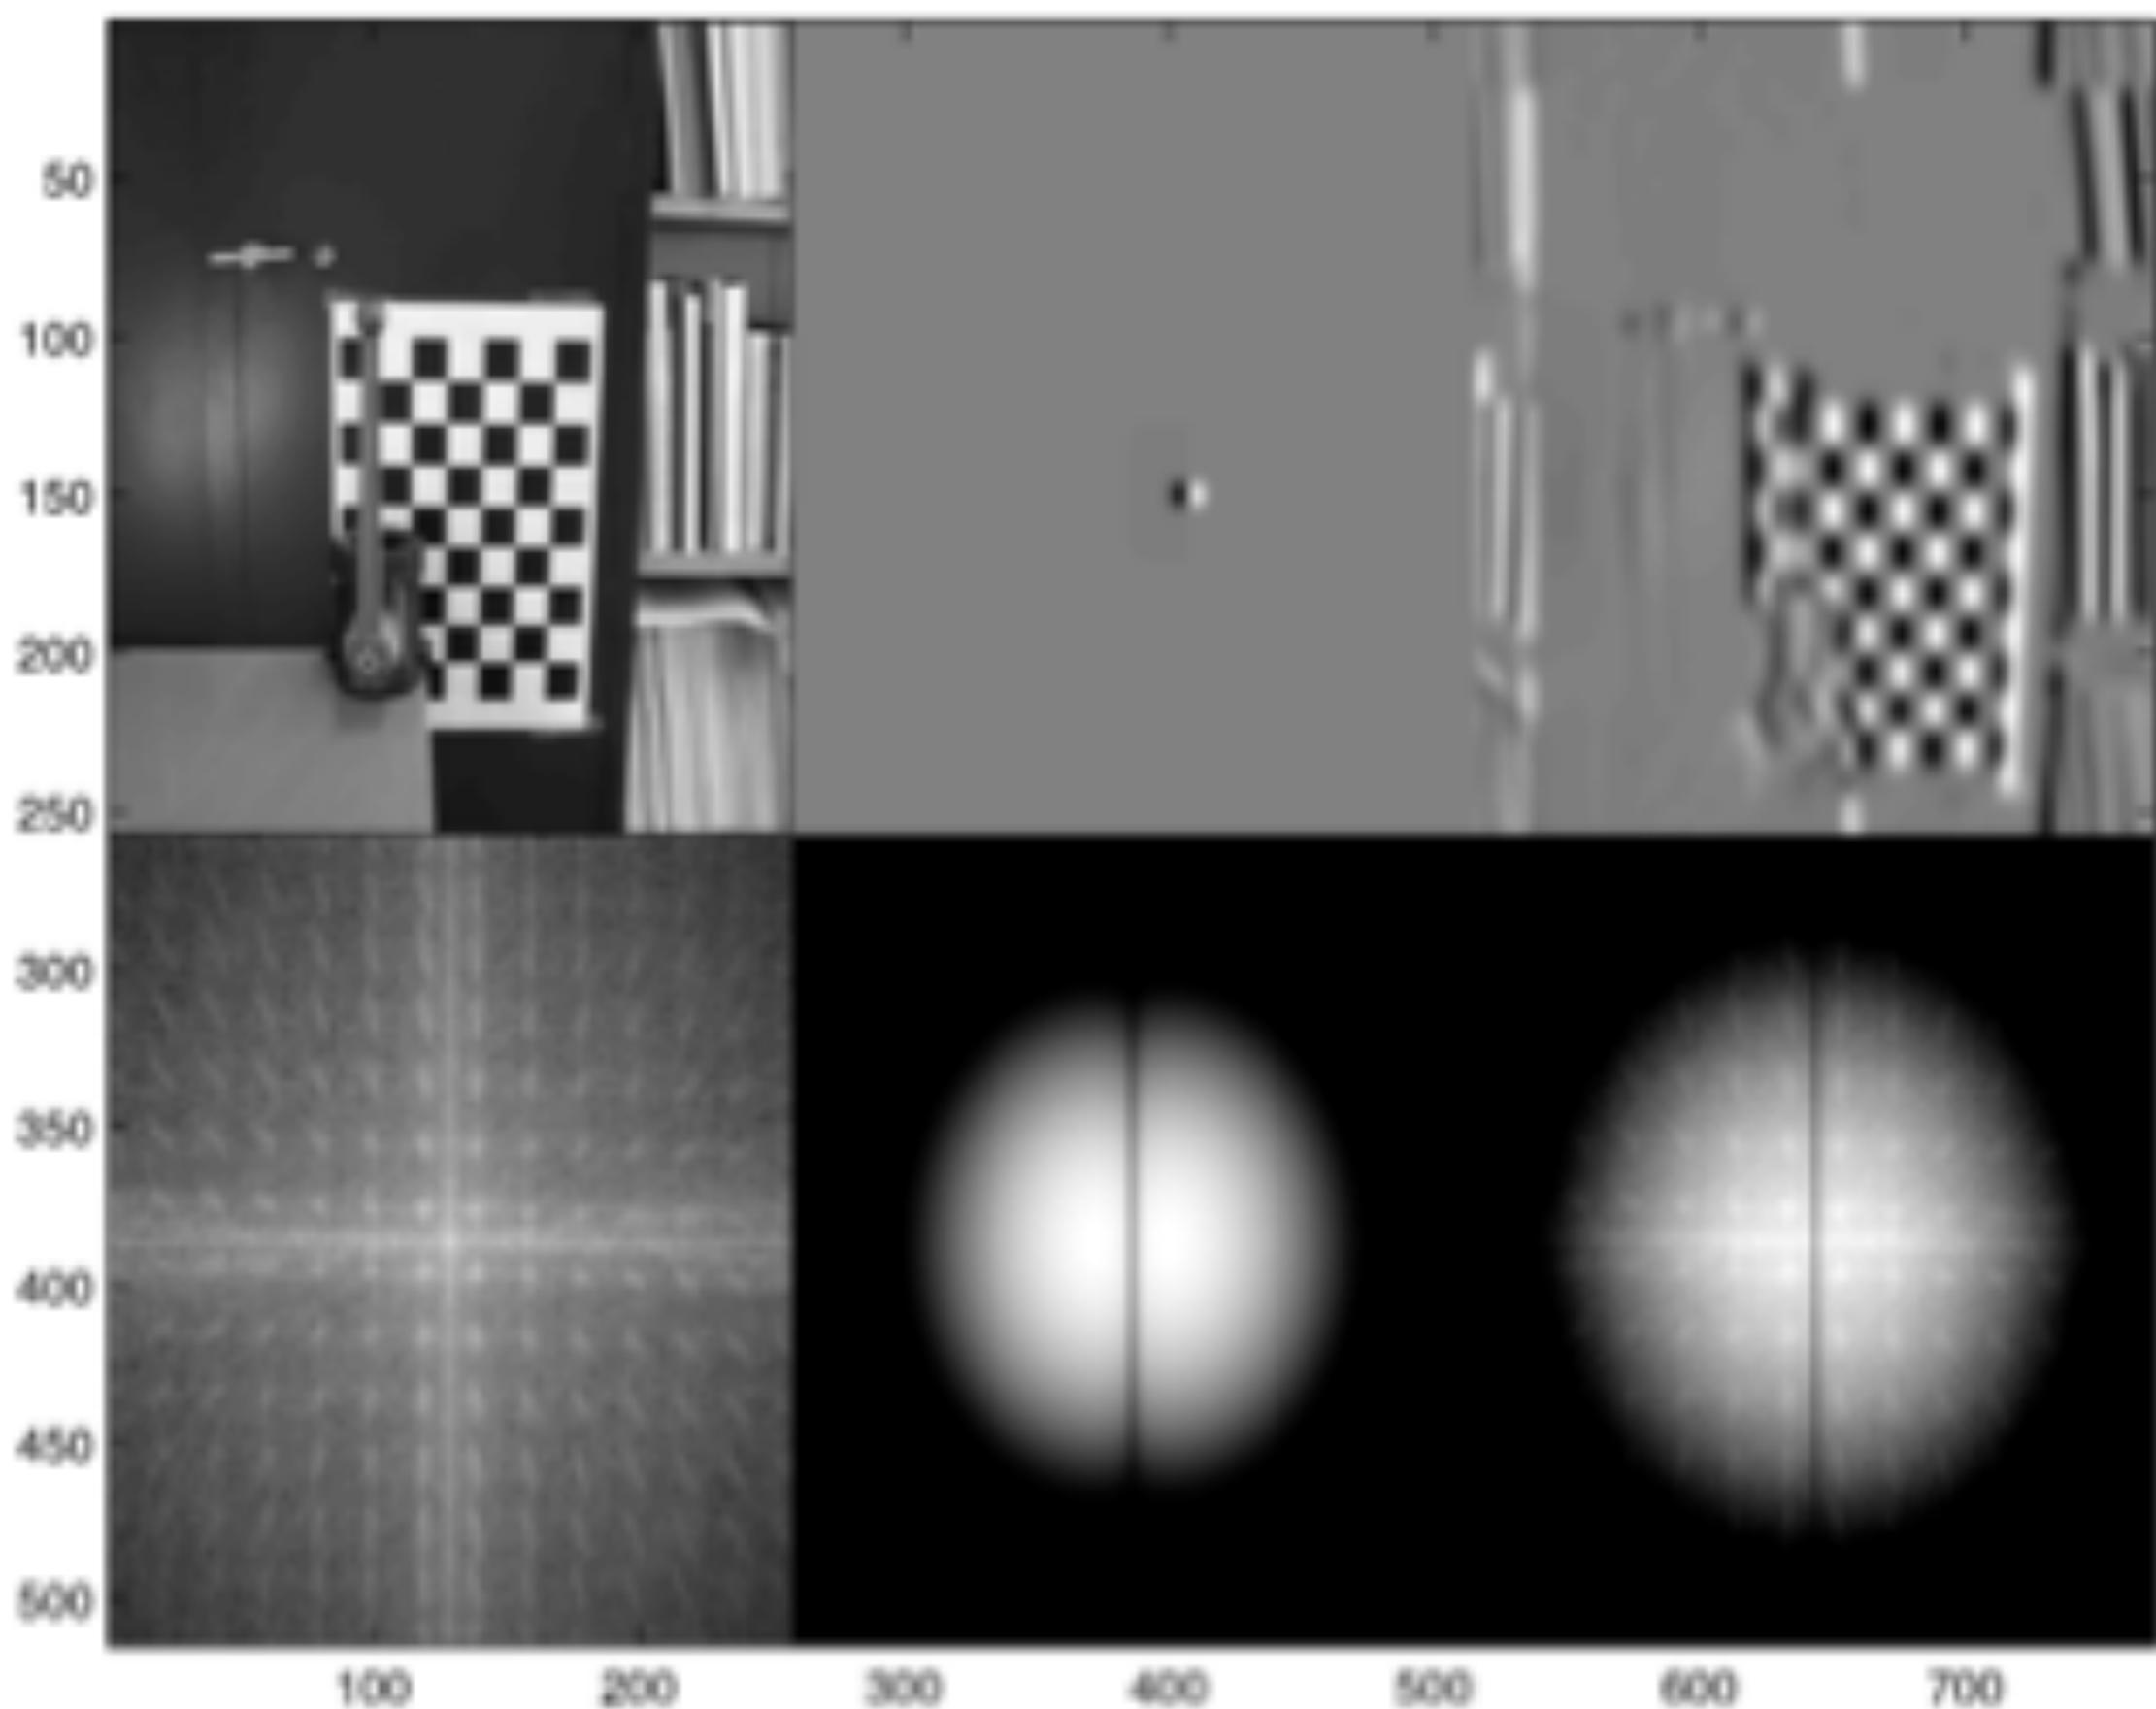
Sensitive to noise.
Combine with smoothing:

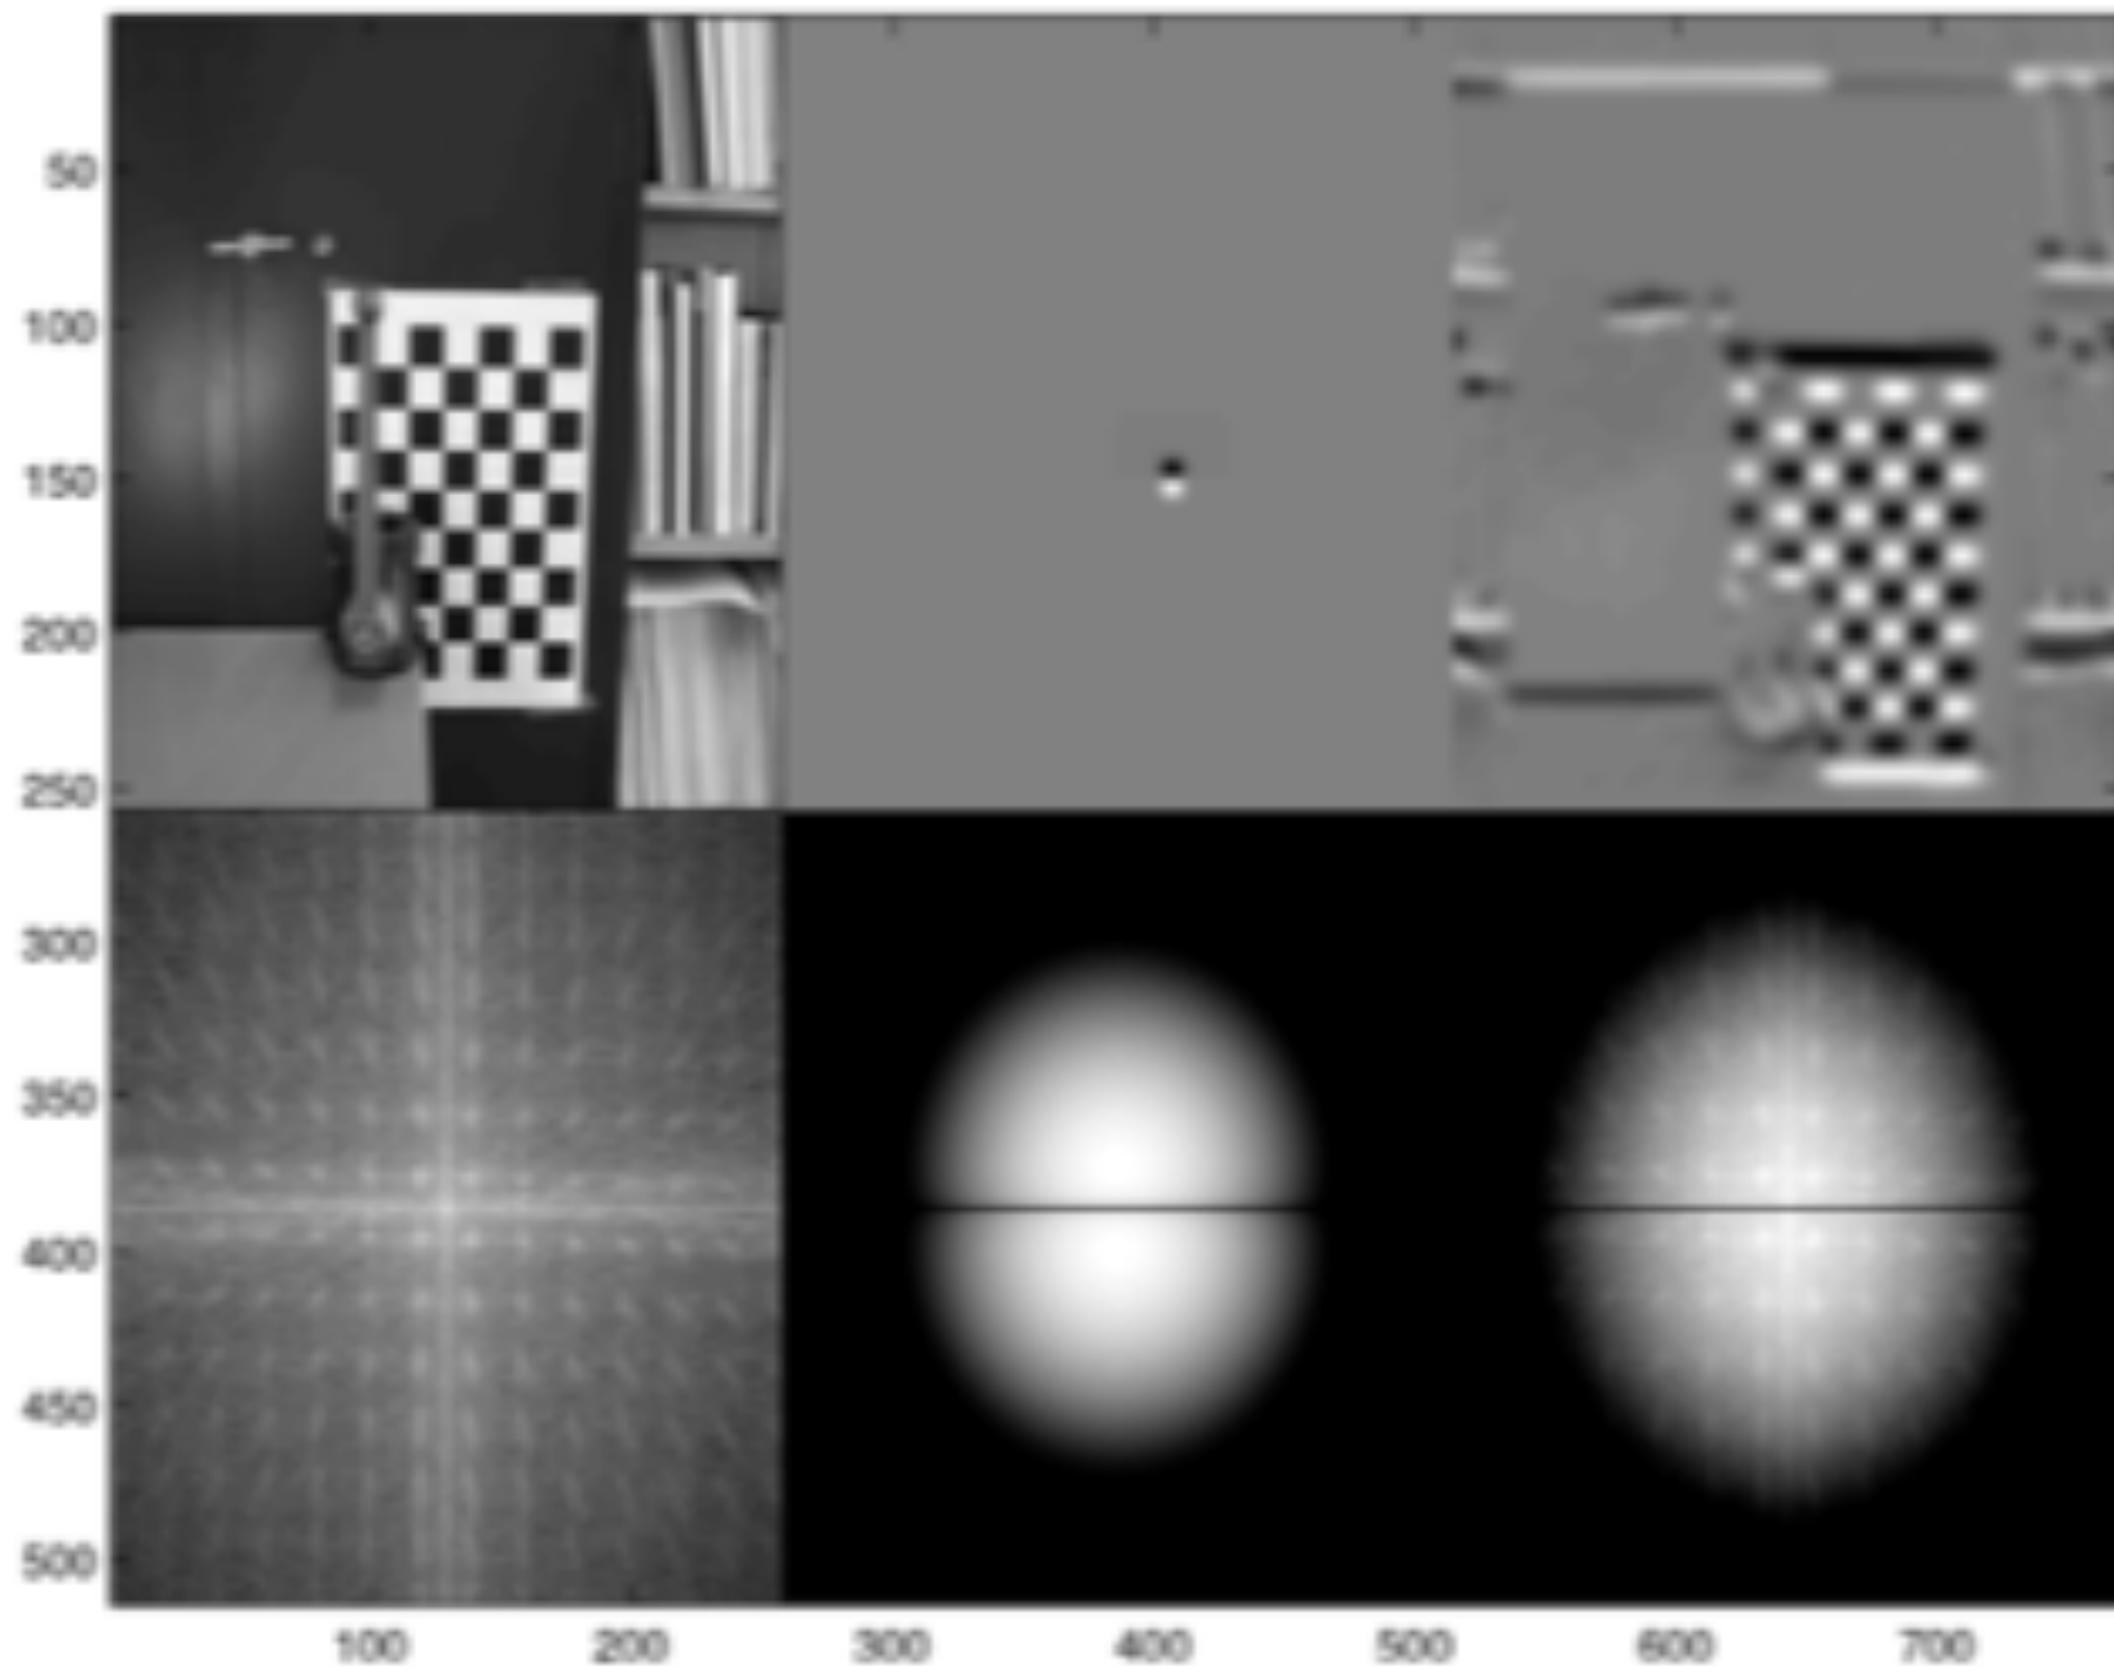
$$f \rightarrow \phi * f \rightarrow \frac{\partial}{\partial x} \phi * f$$

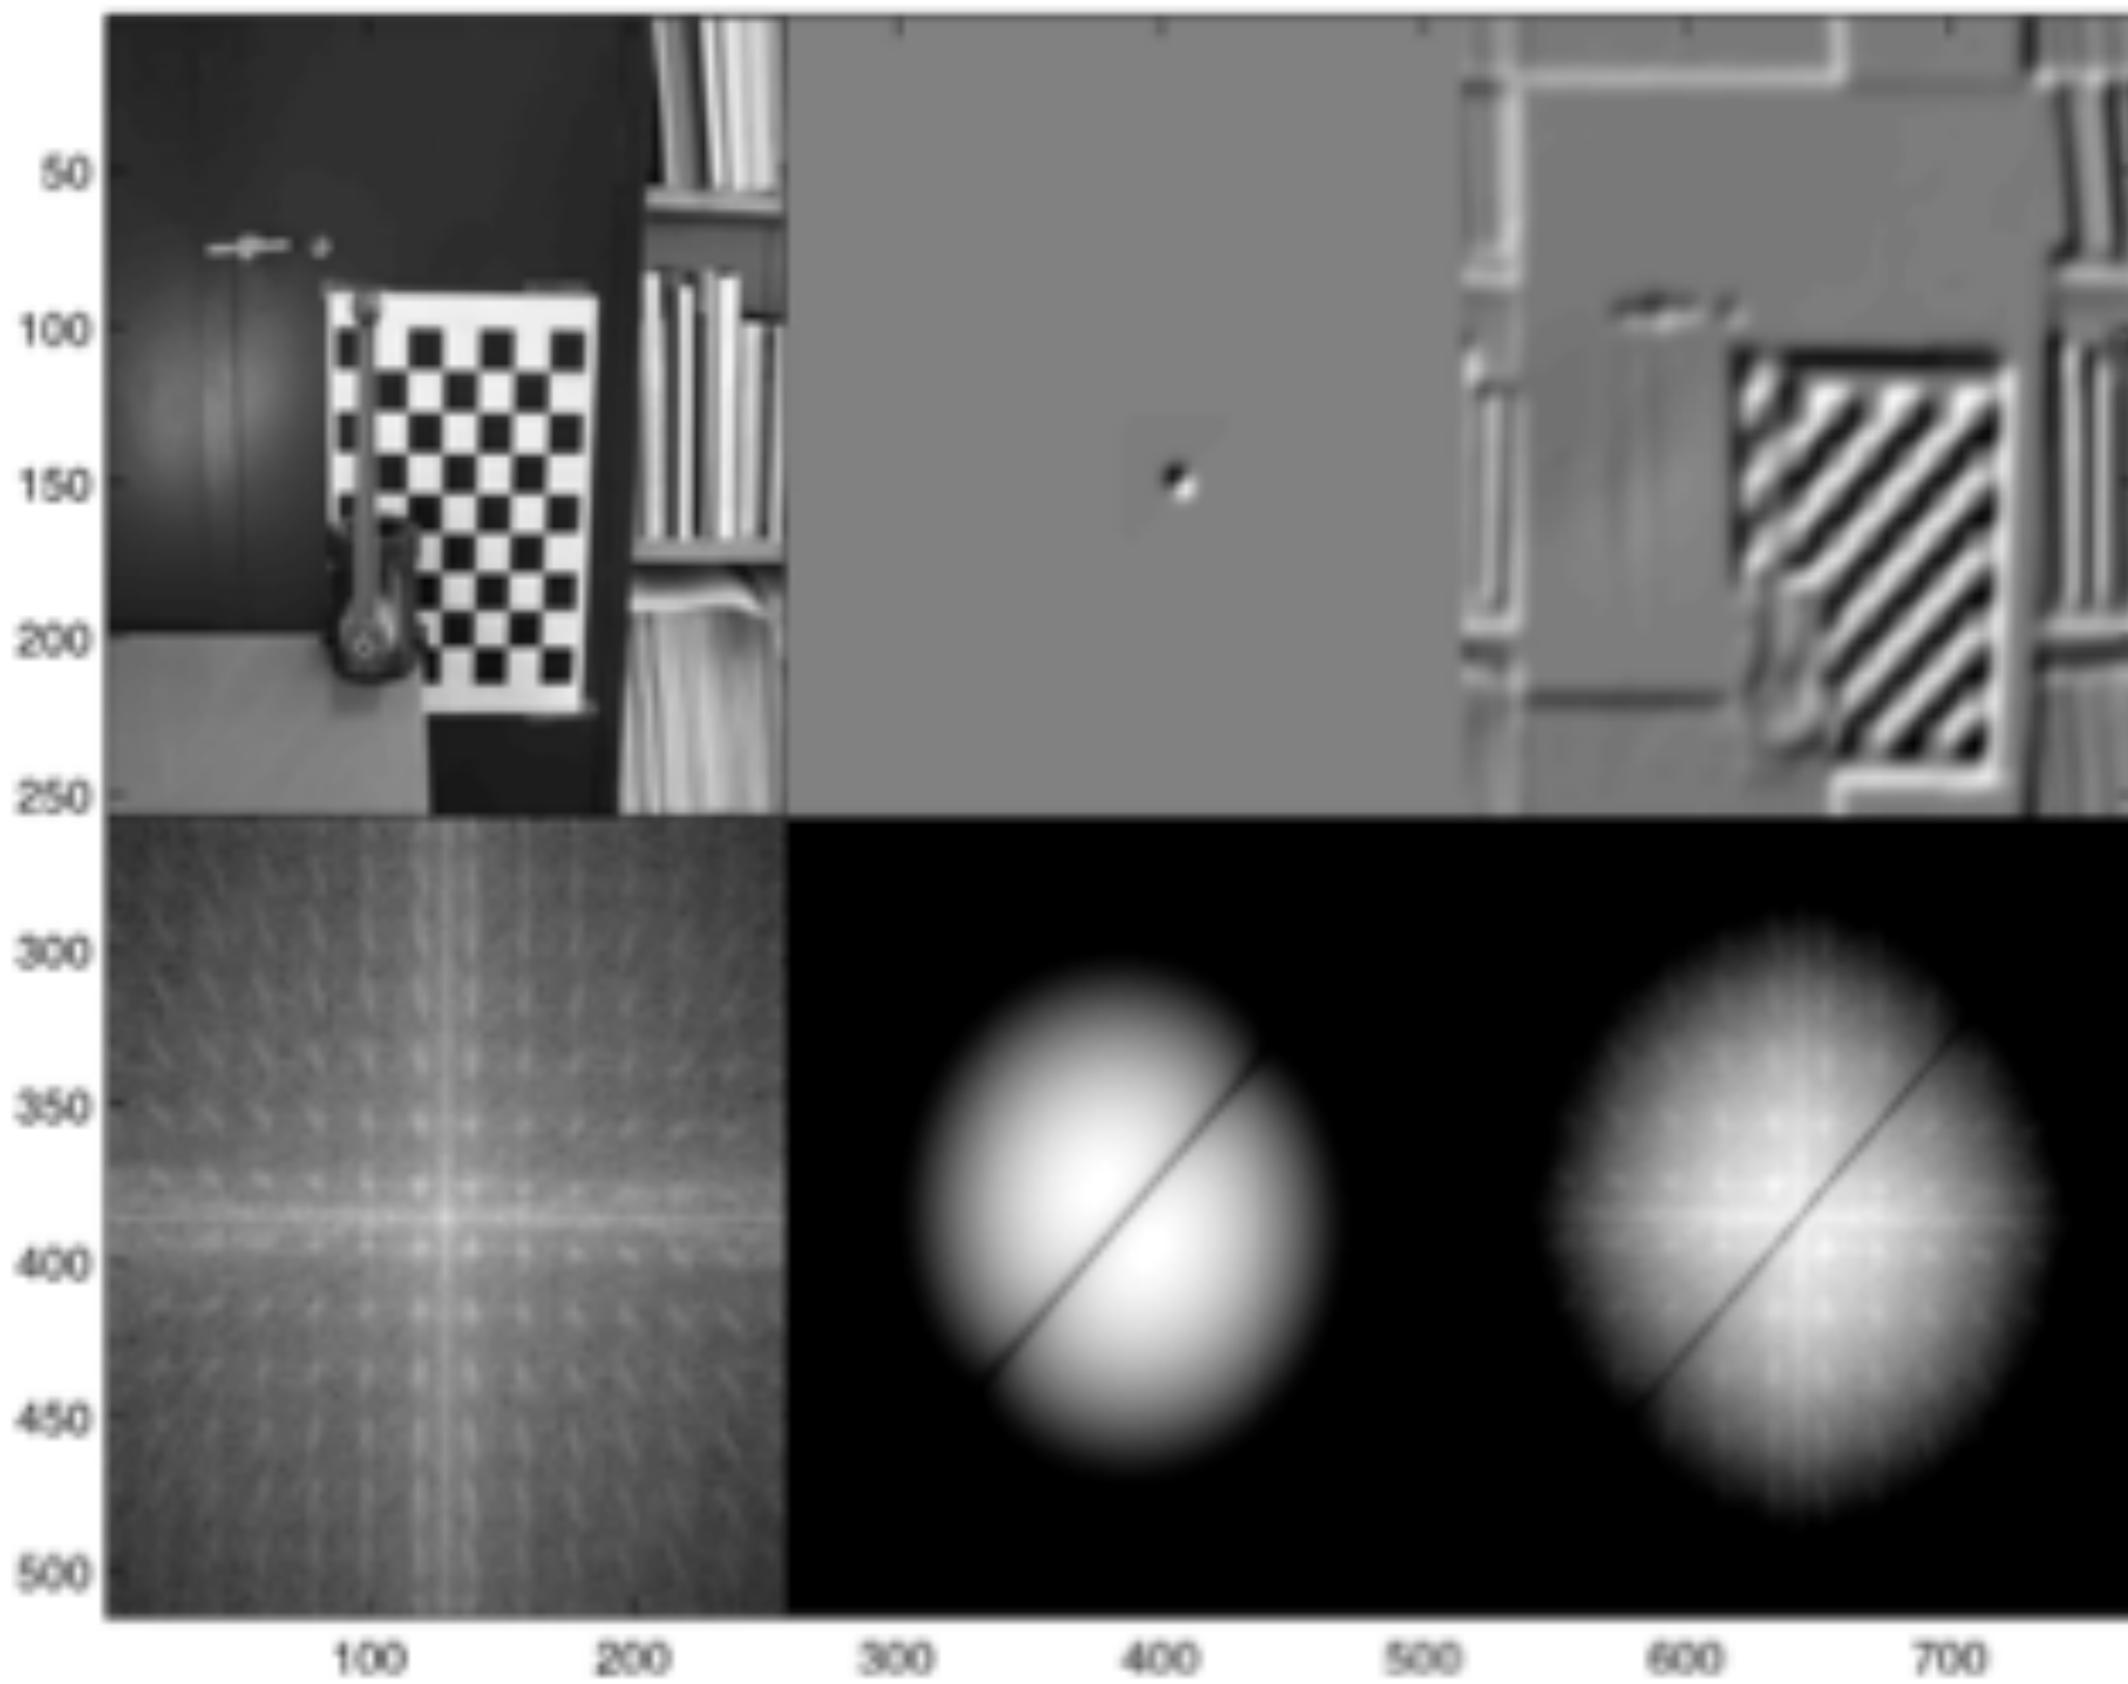
$$\frac{\partial}{\partial x} \phi = -\frac{x}{\sqrt{2\sigma^6\pi}} e^{-x^2/(2\sigma^2)}$$

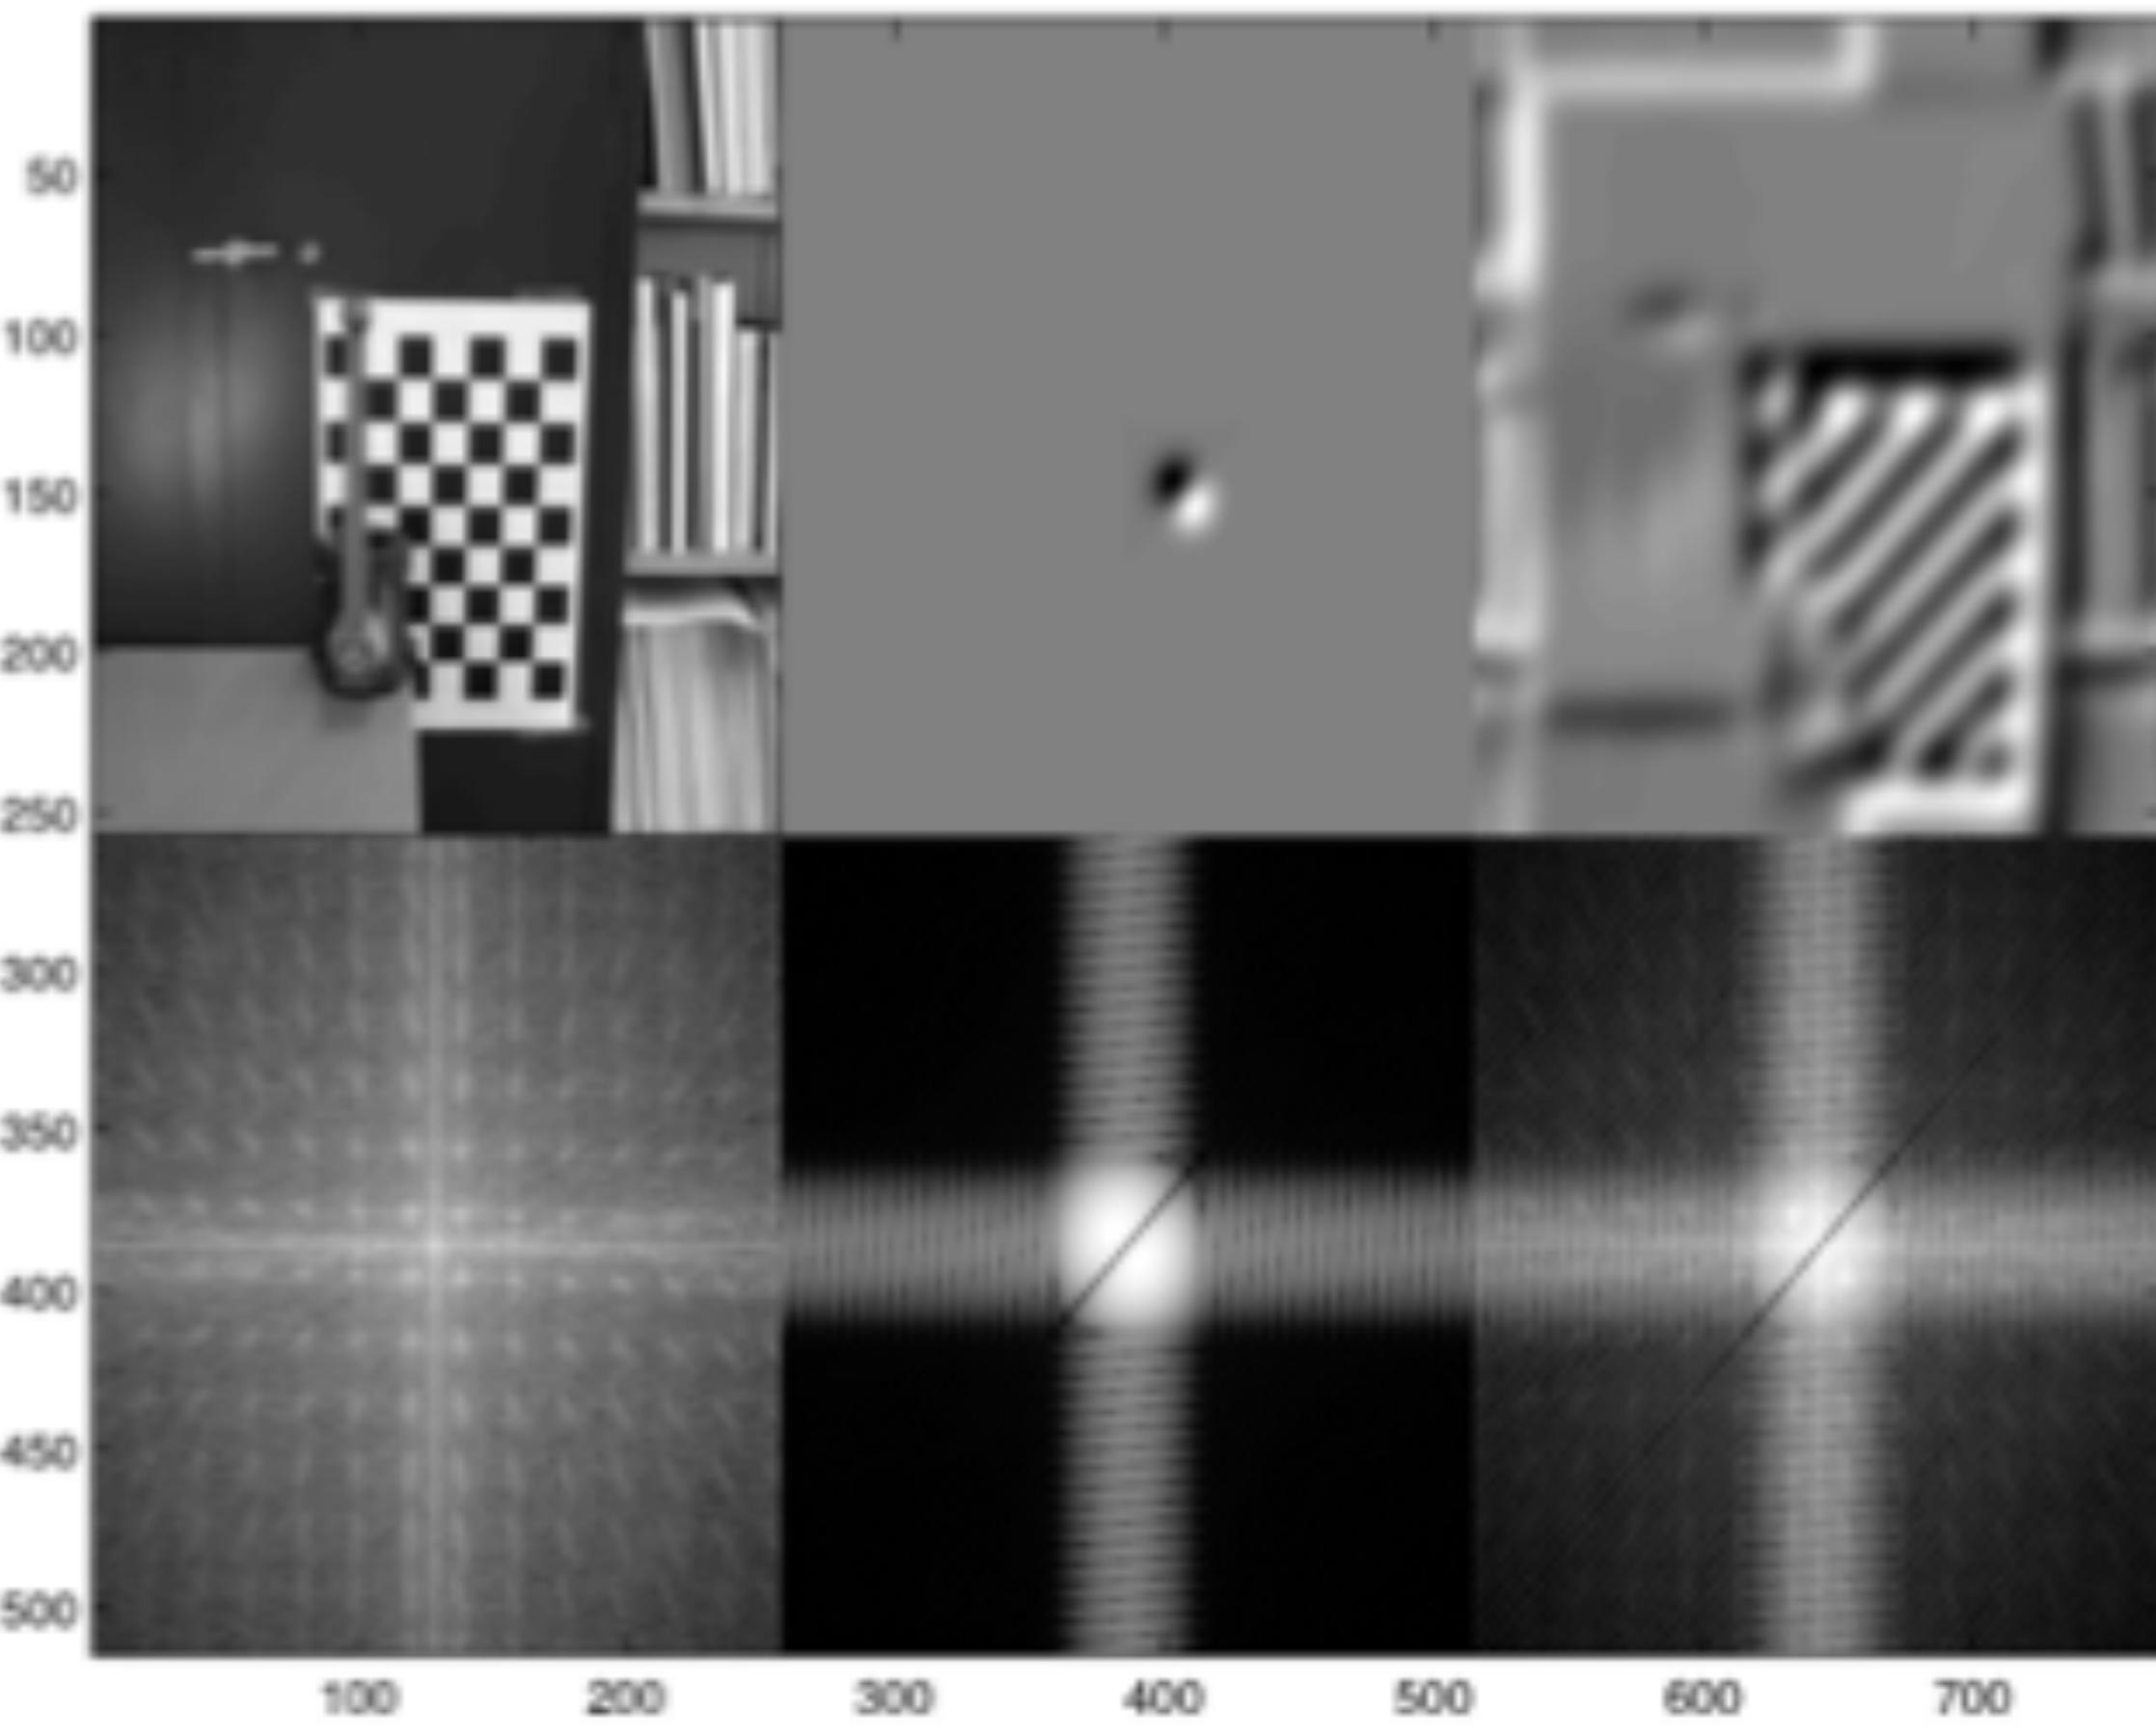












100

200

300

400

500

600

700

50

100

150

200

250

300

350

400

450

500



D

UNIVERSITY

Review

- Convolution (with flip) and cross-correlation (without flip)
 - Properties
 - Examples
- Convolution theorem
- Interpreting convolutions through the Fourier transform

- Read lecture notes
- Experiment with matlab demo scripts
- Finish assignment 1

Master's Thesis Suggestion of the day

- Visual SLAM and recognition for drones



LUND
UNIVERSITY

350