

Image Analysis (FMAN20)

Lecture 7, 2018

MAGNUS OSKARSSON

9 Decbel - München

COMPUTER

2.9

4

5.6

8

11

16

22



Image Analysis - Motivation

Aligning Books and Movies: Towards Story-like Visual Explanations by Watching Movies and Reading Books
Yukun Zhu*, Ryan Kiros*, Richard Zemel, Ruslan Salakhutdinov, Raquel Urtasun, Antonio Torralba, Sanja Fidler
Arxiv, June 2015



Image Analysis - Motivation

Aligning Books and Movies: Towards Story-like Visual Explanations by Watching Movies and Reading Books
Yukun Zhu*, Ryan Kiros*, Richard Zemel, Ruslan Salakhutdinov, Raquel Urtasun, Antonio Torralba, Sanja Fidler
Arxiv, June 2015



Overview – Deep Learning

1. History and Motivation
2. Components of Deep Learning
 1. Convolution
 2. Non-linear
 3. Max pooling
 4. Soft-Max
3. Network Design
4. Training
5. Examples

Computer vision

bridge the gap between pixels and meaning

0	3	2	5	4	7	6	9	8
3	0	1	2	3	4	5	6	7
2	1	0	3	2	5	4	7	6
5	2	3	0	1	2	3	4	5
4	3	2	1	0	3	2	5	4
7	4	5	2	3	0	1	2	3
6	5	4	3	2	1	0	3	2
9	6	7	4	5	2	3	0	1
8	7	6	5	4	3	2	1	0

**Images are
collections of
intensity
measurements (or
RGB, or ...)**


Computer vision

bridge the gap between pixels and meaning


0	3	2	5	4	7	6	9	8
3	0	1	2	3	4	5	6	7
2	1	0	3	2	5	4	7	6
5	2	3	0	1	2	3	4	5
4	3	2	1	0	3	2	5	4
7	4	5	2	3	0	1	2	3
6	5	4	3	2	1	0	3	2
9	6	7	4	5	2	3	0	1
8	7	6	5	4	3	2	1	0



Autonomous Precision Livestock Farming



Deep Learning



Deep learning

Convolutional Neural Networks

- Slides and material from
- <http://www.cs.nyu.edu/~yann/talks/lecun-ranzato-icml2013.pdf>
- MatConvNet
- <http://www.robots.ox.ac.uk/~vgg/practicals/cnn/>
- Gabrielle Flood's master's thesis
- Anna Gummesson's master's thesis

Components for deep learning


- One neuron
 - Example: Logistic regression
 - Classification model (x feature vector, (w,b) parameters, s smooth thresholding

$$x \in R^d, w \in R^d, b \in R, f(x) = s(w^T x + b)$$

- Logistic regression

$$s(z) = \frac{1}{1 + e^{-x}}$$

- ML estimate of parameters (w,b) is a convex optimization problem




Single Layer Neural Networks

One Neuron

- One neuron

$$x \in R^d, w \in R^d, b \in R, f(x) = s(w^T x + b)$$



Single Layer Neural Networks


Several Neurons

- Several parallel neurons

$$x \in R^d, y \in R^k, B \in R^d, W - k \times d \text{ matrix}$$

$$y = s(Wx + B)$$


- Elementwise smooth thresholding – s



Artificial Neural Networks

One hidden layer


- Multi-class classification
- One hidden layer
- Trained by back-propagation
- Popular since the 1990s




Deep Neural Networks

Many layers

- However
- Naively implemented would give too many parameters
- Example
- 1M pixel image
- 1M hidden layers
- 10^{12} parameters between each pair of layers




Convolutional neural network, CNN




CNN-Blocks - Convolutional layer

Convolution of an image as a filter-operation.




Convolution of an image represented as a sparsely connected ANN.



CNN-Blocks - Convolutional layer

- Input: Data block x of size $m \times n \times k_1$
- Output: Data block y of size $m \times n \times k_2$
- Filter: Filter kernel block w of size $m_w \times n_w \times k_1 \times k_2$
- Offsets: Vector w_o of length k_2




$$y(i, j, k) = w_o(k) + \sum_u \sum_v \sum_l x(i - u, j - v, l) w(u, v, l, k)$$

CNN-Blocks - Convolutional layer



$$y(i, j) = \sum_u \sum_v x(i - u, j - v) w(u, v)$$

$$y(i, j) = w_o + \sum_l \left(\sum_u \sum_v x(i - u, j - v, l) w(u, v, l) \right)$$



$$y(i, j, k) = w_o(k) + \sum_u \sum_v \sum_l x(i - u, j - v, l) w(u, v, l, k)$$



CNN-Blocks - Max-pooling



CNN-Blocks - RELU


$$f(x) = \max(0, x)$$

$$y(i, j, k) = \max(x(i, j, k), 0)$$




CNN-Blocks – Softmax
(convert from 'log probabilities' d_j to
'probabilities' that sum to 1)

$$p_j = \frac{e^{d_j}}{\sum_{k=1}^m e^{d_k}}$$



Gör om till 1x1 stav


Result, Network design



CNN-Blocks – input – output



$$y = f(x, w)$$



Input: image x of size $m \times n \times k$, typically $k=1$ (gray-scale) or $k=3$ (color)


Output: vector y of size $1 \times 1 \times N$, which we interpret as N probabilities y_j

The probability that the image x is of class j

Training data (x_i, c_i)



$$y = f(x, w)$$



Input: image x of size $m \times n \times k$, typically $k=1$ (gray-scale) or $k=3$ (colour)

Output: vector y of size $1 \times 1 \times N$, which we interpret as N probabilities y_j

The probability that the image x is of class j y_{c_i}

Training data $T = \{(x_1, c_1), \dots, (x_N, c_N)\}$

- Classification network $y = f(x, w)$
- Evaluate one example (x_k, c_k) (like adding another layer)

$$\sum_{k=1}^N -\log y(x_k, w)_{c_k}$$

- Evaluation function: $g(T, w) = \sum_{k=1}^N -\log y(x_k, w)_{c_k}$

- Solve

$$\min_w g(T, w)$$

Example: OCR, classify images as a-z, Network design

```
>> net
```

```
net =
```


```
layers: {1x7 cell}
imageMean: 0.9176775
```

```
>> vl_simplenn_display(net)
```

layer	1	2	3	4	5	6	7
type	cnv	mpool	cnv	mpool	cnv	relu	cnv
support	5x5	2x2	5x5	2x2	4x4	1x1	2x2
stride	1	2	1	2	1	1	1
pad	0	0	0	0	0	0	0
out dim	20	20	50	50	500	500	26
filt dim	1	n/a	20	n/a	50	n/a	500
rec. field	5	6	14	16	28	28	32
c/g net KB	4/0	0/0	196/0	0/0	3129/0	0/0	406/0

```
total network CPU/GPU memory: 3.6/0 MB
```

```
>>
```




Example: OCR, classify images as a-z

Training data


training chars for 'a'



Example: OCR, classify images as a-z, Training



$$g(T, w) = \sum_{k=1}^N -\log y(x_k, w)_{c_k}$$



$$\#\{c_k = \operatorname{argmax}_i y(x_k, w)_i\}$$

Tricks

- **Stochastic Gradient Descent**

- Computation of


$$g(T, w) = \sum_{k=1}^N -\log y(x_k, w)_{c_k}$$

- Requires going through all examples (all N).
- If N is large and/or if computing $y(x_k, w)$ is time-consuming, use stochastic gradient descent, i.e. update parameters using subsets of training data.


- **Jittering** - construct a larger training set by perturbing the examples, jittering, translating images, rotating images, warping, mirroring, adding noise, ...'

- **Dropout** – in each computation of $y(x_k, w)$ let a random subset of the neurons die, i.e. set the output to zero.

Generalisation, Expand data set



Generalisation, Dropout




Generalisation, Weight decay (Prior on small weights)

$$E(w) = - \sum_{n=1}^N \log \left(\frac{e^{y_{d(n)}(n)}}{\sum_{i=1}^k e^{y_i(n)}} \right) + \frac{\lambda}{2} \sum_l w_l^2$$

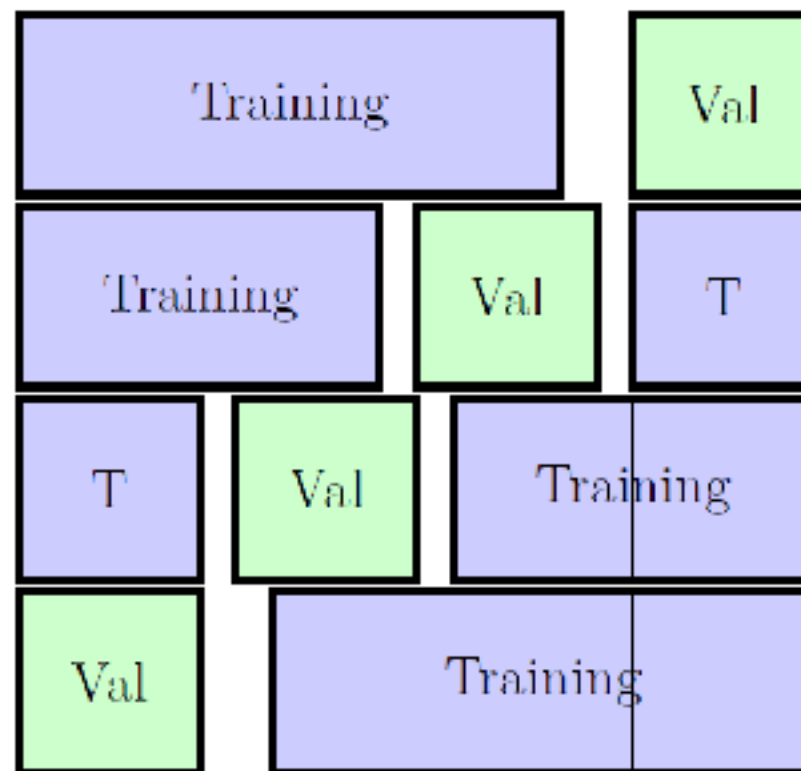
Thoughts

- Modelling. It takes time to
 - Figure out an appropriate network structure
 - Gather data and ground truth
- The optimization does not always work
 - Parameters explode
 - Nothing happens
- Visualization of the features is important for understanding.
- Feedback in networks

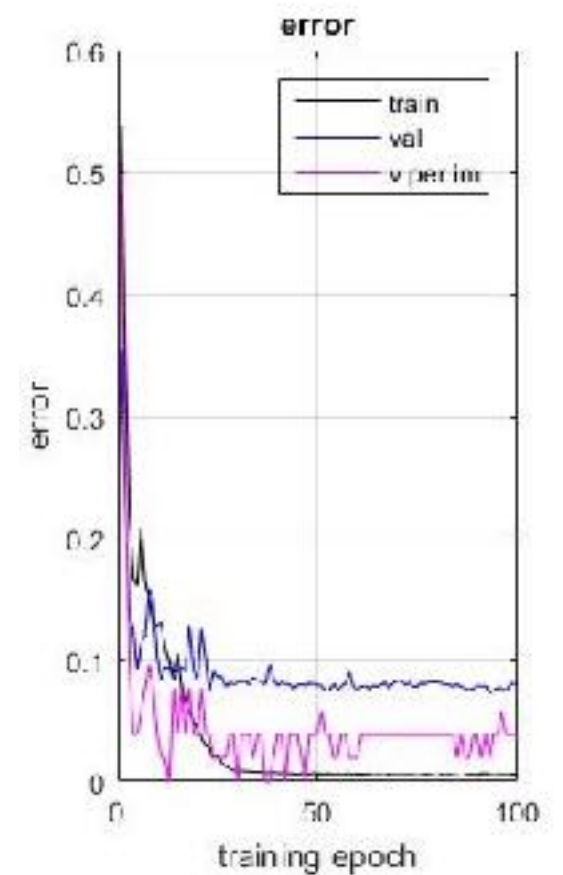


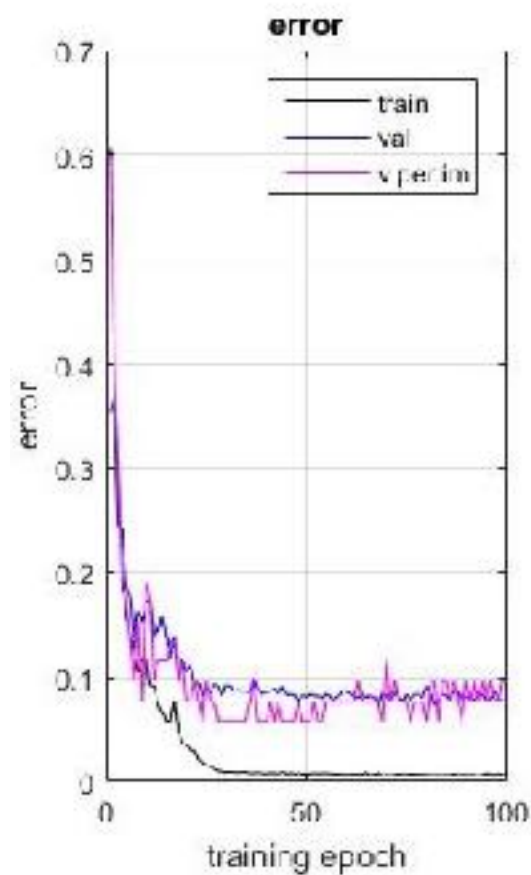
Example: Prostate cancer Data



Result, Cross-validation



Example: Prostate cancer Training

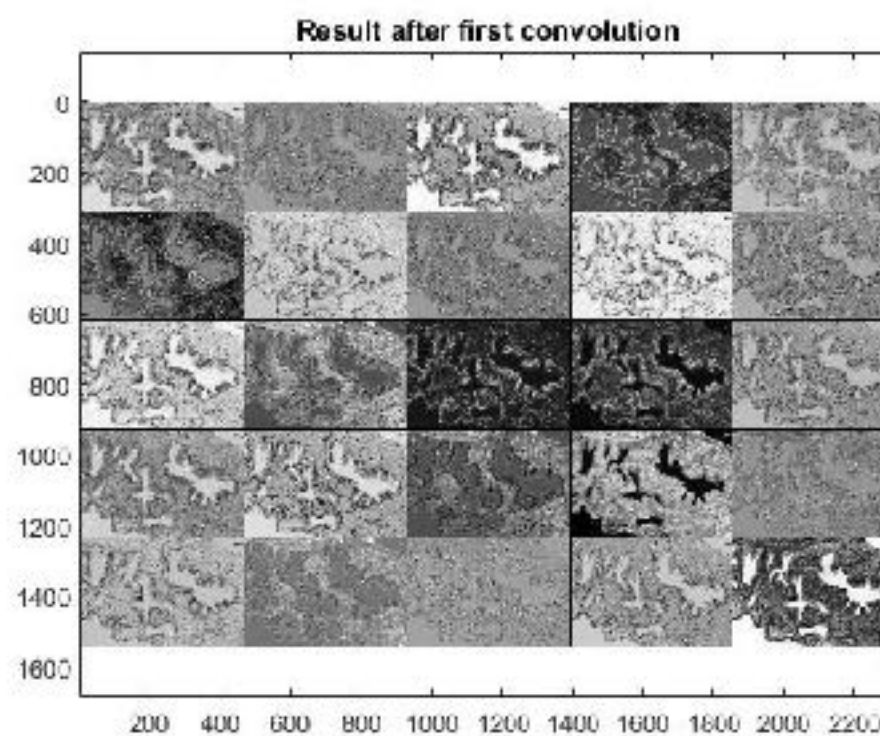



Example: Prostate cancer

Results: Confusion matrix

$$\begin{bmatrix} 51 & 0 & 0 & 0 \\ 3 & 46 & 3 & 1 \\ 0 & 6 & 43 & 0 \\ 0 & 3 & 0 & 52 \end{bmatrix}$$

Visualisation



Examples: Image net, Data

- ImageNet Large Scale Visual Recognition Challenge
- Yearly challenge since 2010
- 2011 - 25% error
- 2012 - 16% error (using CNN). This kicked off the deep learning hype
- 2015 – 4% error
- By 2015, researchers reported that current software exceeds human ability at the narrow ILSVCR tasks.
- However, as one of the challenge's organisers, Olga Russakovsky, pointed out in 2015, the programs only have to identify images as belonging to one of a thousand categories!

Training Deep Learning

- Data
 - Obtain data,
 - cut-outs of right size,
 - jittering,
 - Data expansion (translation, rotation, scaling, mirroring, adding noise, ...)
- Data
 - Obtain ground truth
 - How should the problem be coded

Training Deep Learning

- Hyperparameters
 - How many layers
 - Size of convolution kernels
 - Number of channels
 - Order of layers
- Training parameters
 - Initializing weights
 - Momentum
 - ...

Non-linear function

- Different choices of non-linear functions.

$$f(x) = \max(0, x)$$

- Faster learning

$$f(x) = \ln(1 + e^x)$$

- Rectifier ...

$$f(x) = \max(0, x)$$

$$f(x) = \begin{cases} x & \text{if } x > 0 \\ 0.01x & \text{otherwise} \end{cases}$$

- ... currently most popular, faster training

$$f(x) = (1 + e^{-x})^{-1}$$

- Arguments

$$f(x) = \tanh(x)$$

Training Deep Learning

- Once all of these are in place, there are several good systems for optimizing the parameters
 - MatConvNet, TensorFlow, Caffe, Torch7, Theano
 - Can train on single CPU
 - Faster if compiled for GPU
 - Even faster on cluster of computers with multiple GPU (e g LUNARC, <http://www.lunarc.lu.se>)
- More links on home page for PhD course
- <http://www.control.lth.se/Education/DoctorateProgram/deep-learning-study-circle.html>

Deep learning - summary

- What is deep learning
- Supervised vs unsupervised learning
- Goal function, energy function E
- Choice of non-linearity ReLU
- Optimization Back-propagation, SGD
- Tricks dropout
- Examples from speech and vision
- Software
- References
- To learn more – take course on Machine Learning FMAN45


Master's Thesis Suggestion

Masters thesis suggestion of the day



- How do we get training data for this?

Masters thesis suggestion of the day



- Turn it around!
- This is easy to get training data for

Masters thesis suggestion: Action-Based Learning in Brain Circuits and Deep Artificial Neural Networks

Research collaboration with
Christian Balkenius,
Per Petersson,
Jeanette Hellgren Kotaleski




Figure 1. The robot Epi grasping an object.

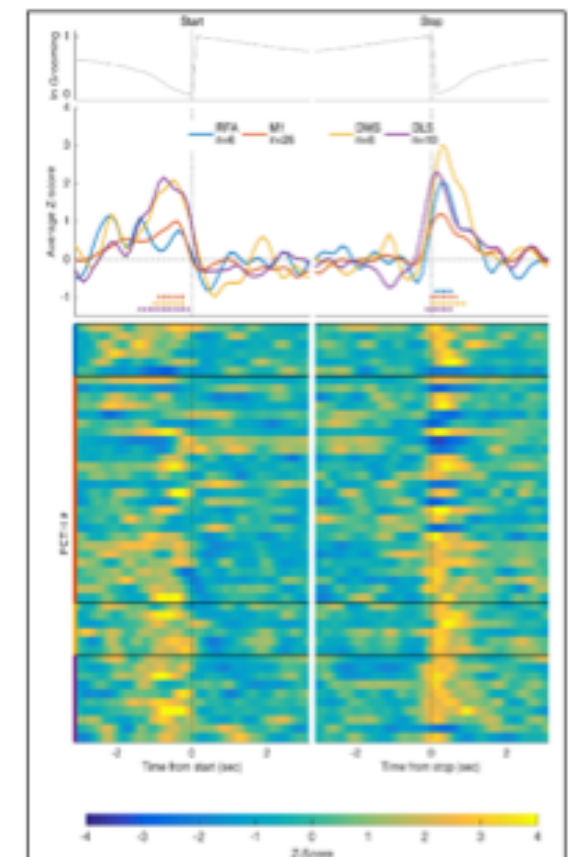



Figure 4. Neuronal firing rate modulation in cortico-basal ganglia circuits reveals distinct action bracketing in spontaneous grooming behavior (Tamte et al., unpublished).





LUND
UNIVERSITY

30